

Programming Modern Processors

Hung-Wei Tseng

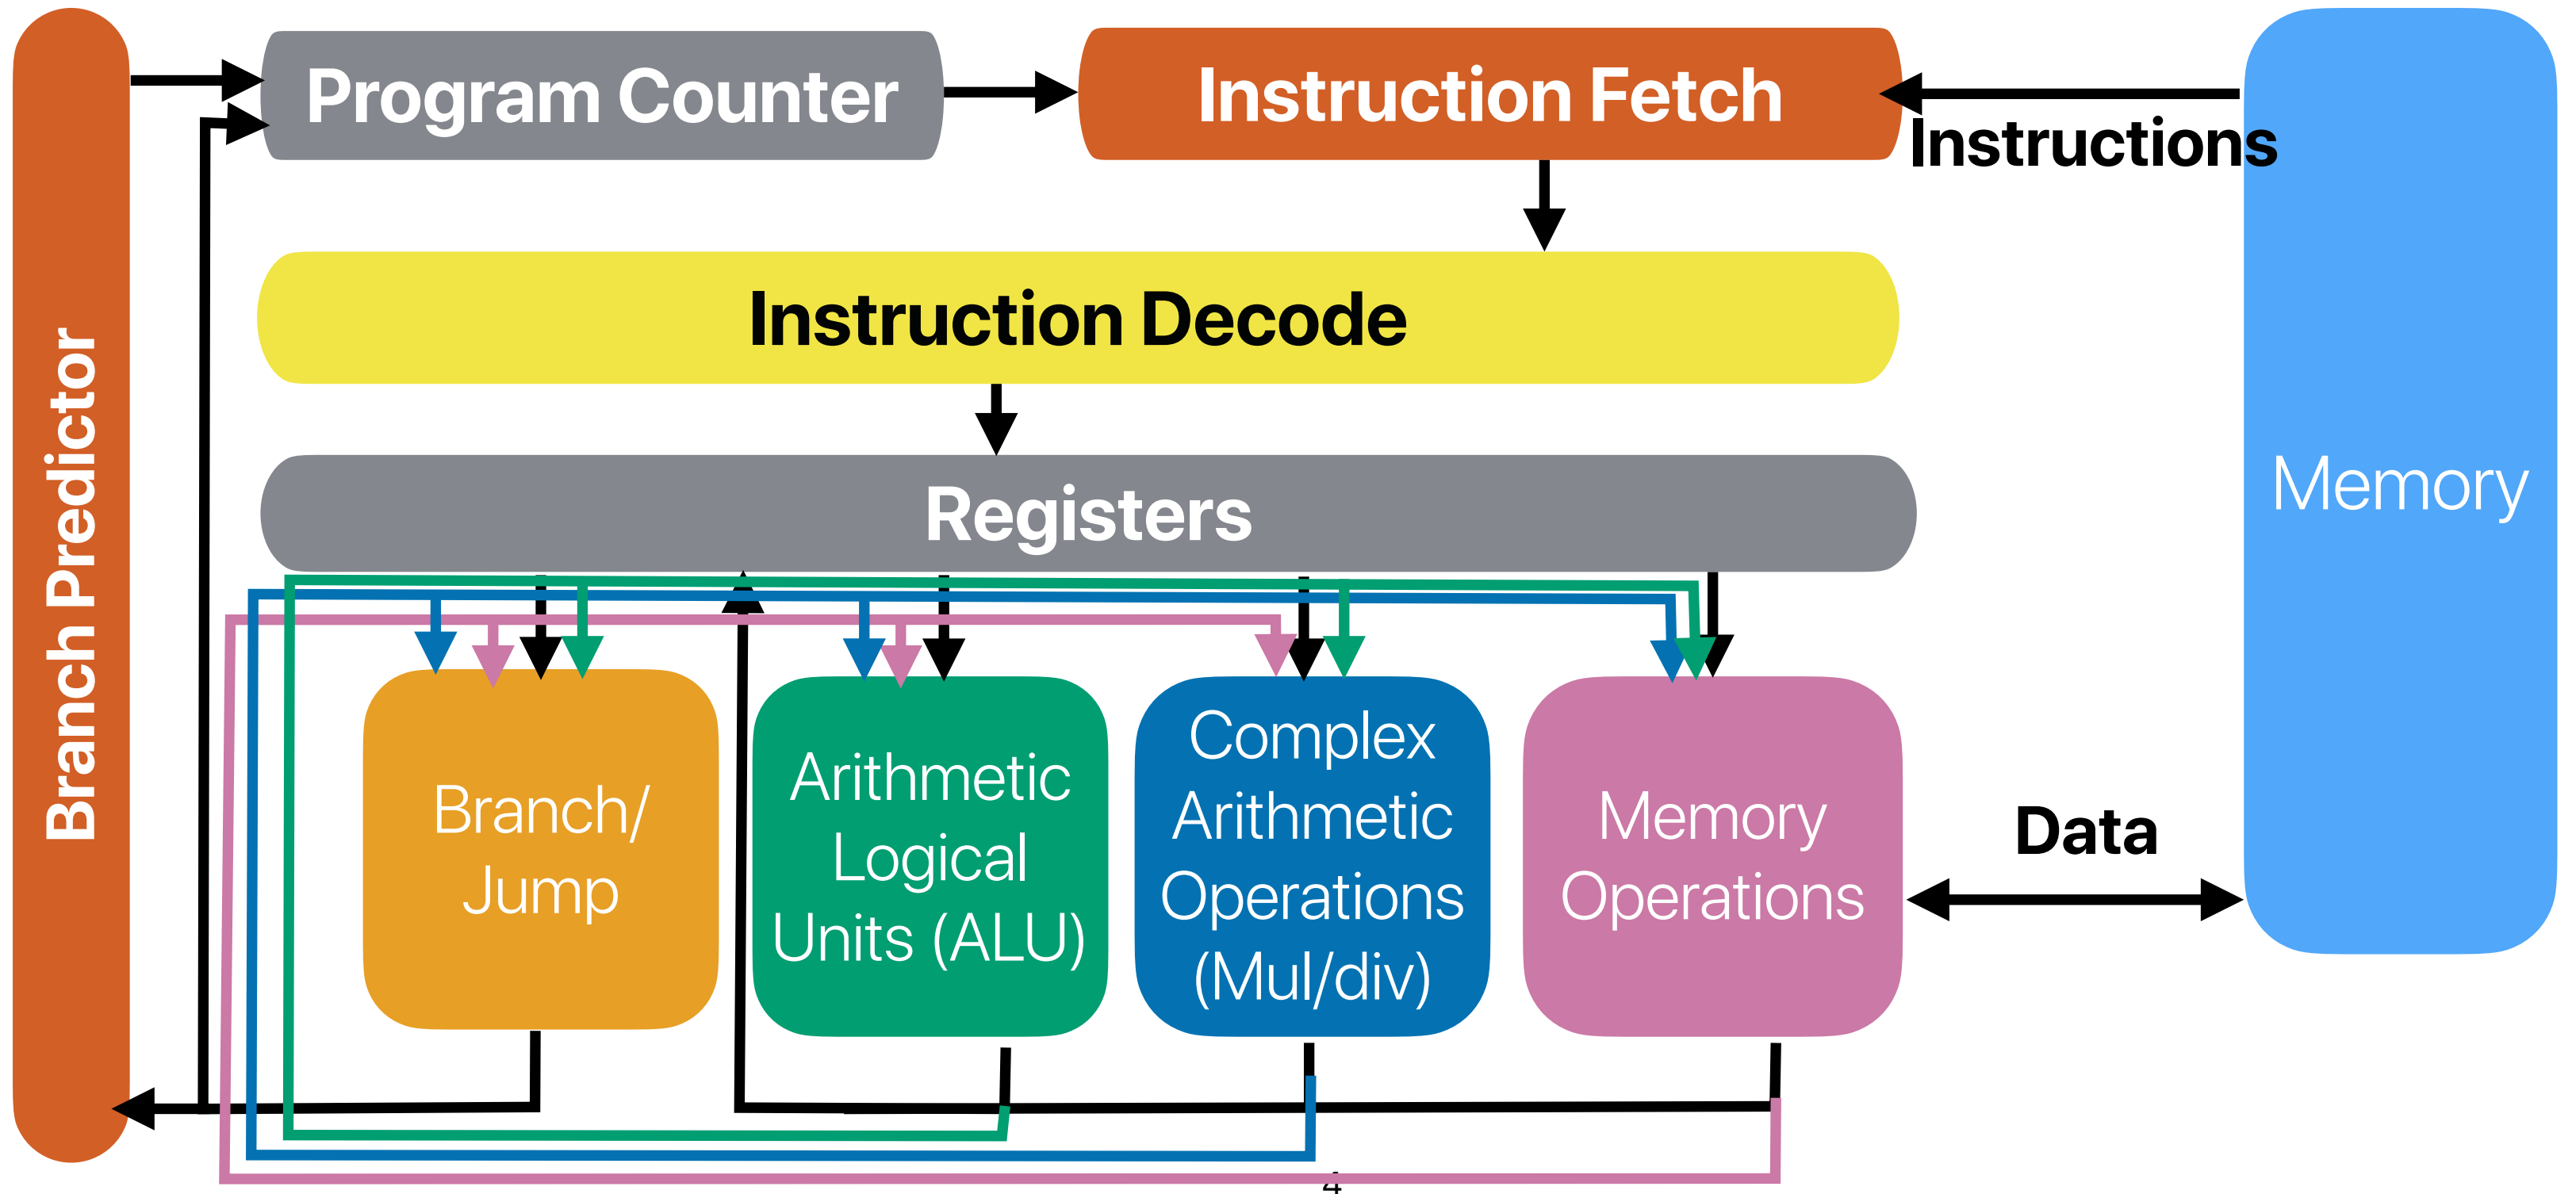
Recap: Three pipeline hazards

- Structural hazards — resource conflicts cannot support simultaneous execution of instructions in the pipeline
- Control hazards — the PC can be changed by an instruction in the pipeline
- Data hazards — an instruction depending on a the result that's not yet generated or propagated when the instruction needs that

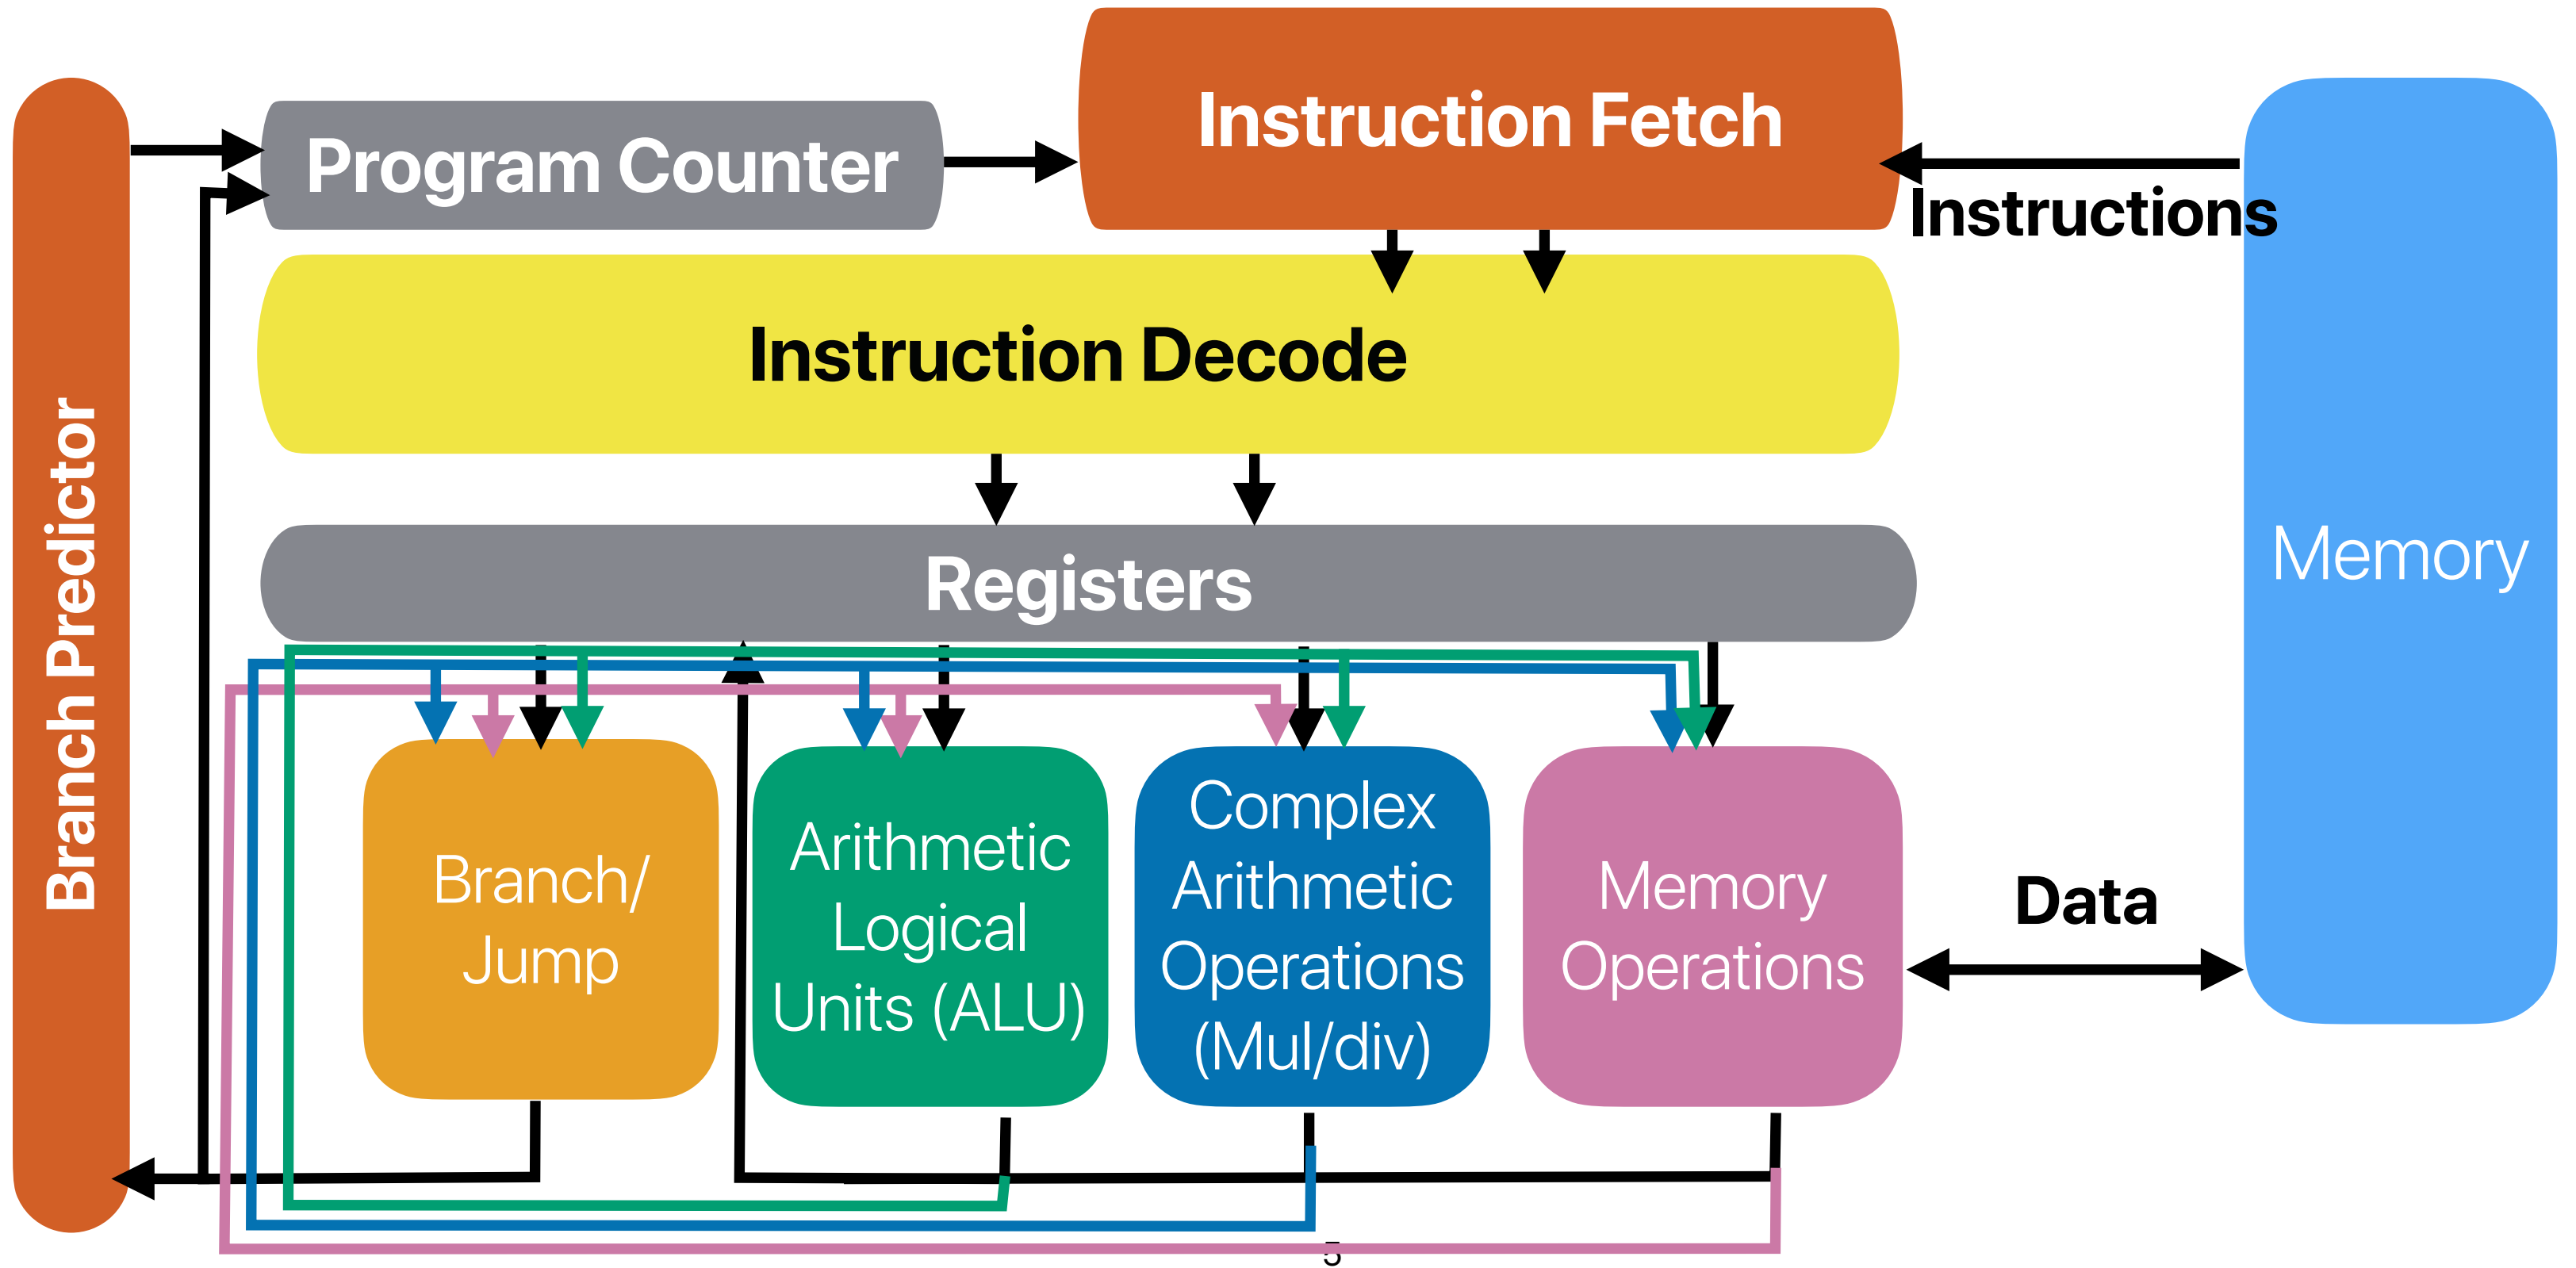
Recap: addressing hazards

- Structural hazards
 - Stall
 - Modify hardware design
- Control hazards
 - Stall
 - Static prediction
 - Dynamic prediction
- Data Hazards
 - Stall
 - Data forwarding
 - Dynamic instruction scheduling

Recap: Data "forwarding"



Super Scalar



Superscalar

- Since we have many functional units now, we should fetch/decode more instructions each cycle so that we can have more instructions to issue!
- Super-scalar: fetch/decode/issue more than one instruction each cycle
 - **Fetch width:** how many instructions can the processor fetch/decode each cycle
 - **Issue width:** how many instructions can the processor issue each cycle
- The theoretical CPI should now be


1

$\min(\text{issue width}, \text{fetch width}, \text{decode width})$

False dependencies


- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph
 - **WAR (Write After Read):** a later instruction overwrites the source of an earlier one
 - 2 and 1, 6 and 3, 7 and 4, 7 and 6
 - **WAW (Write After Write):** a later instruction overwrites the output of an earlier one
 - 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```



What if we can use more registers...

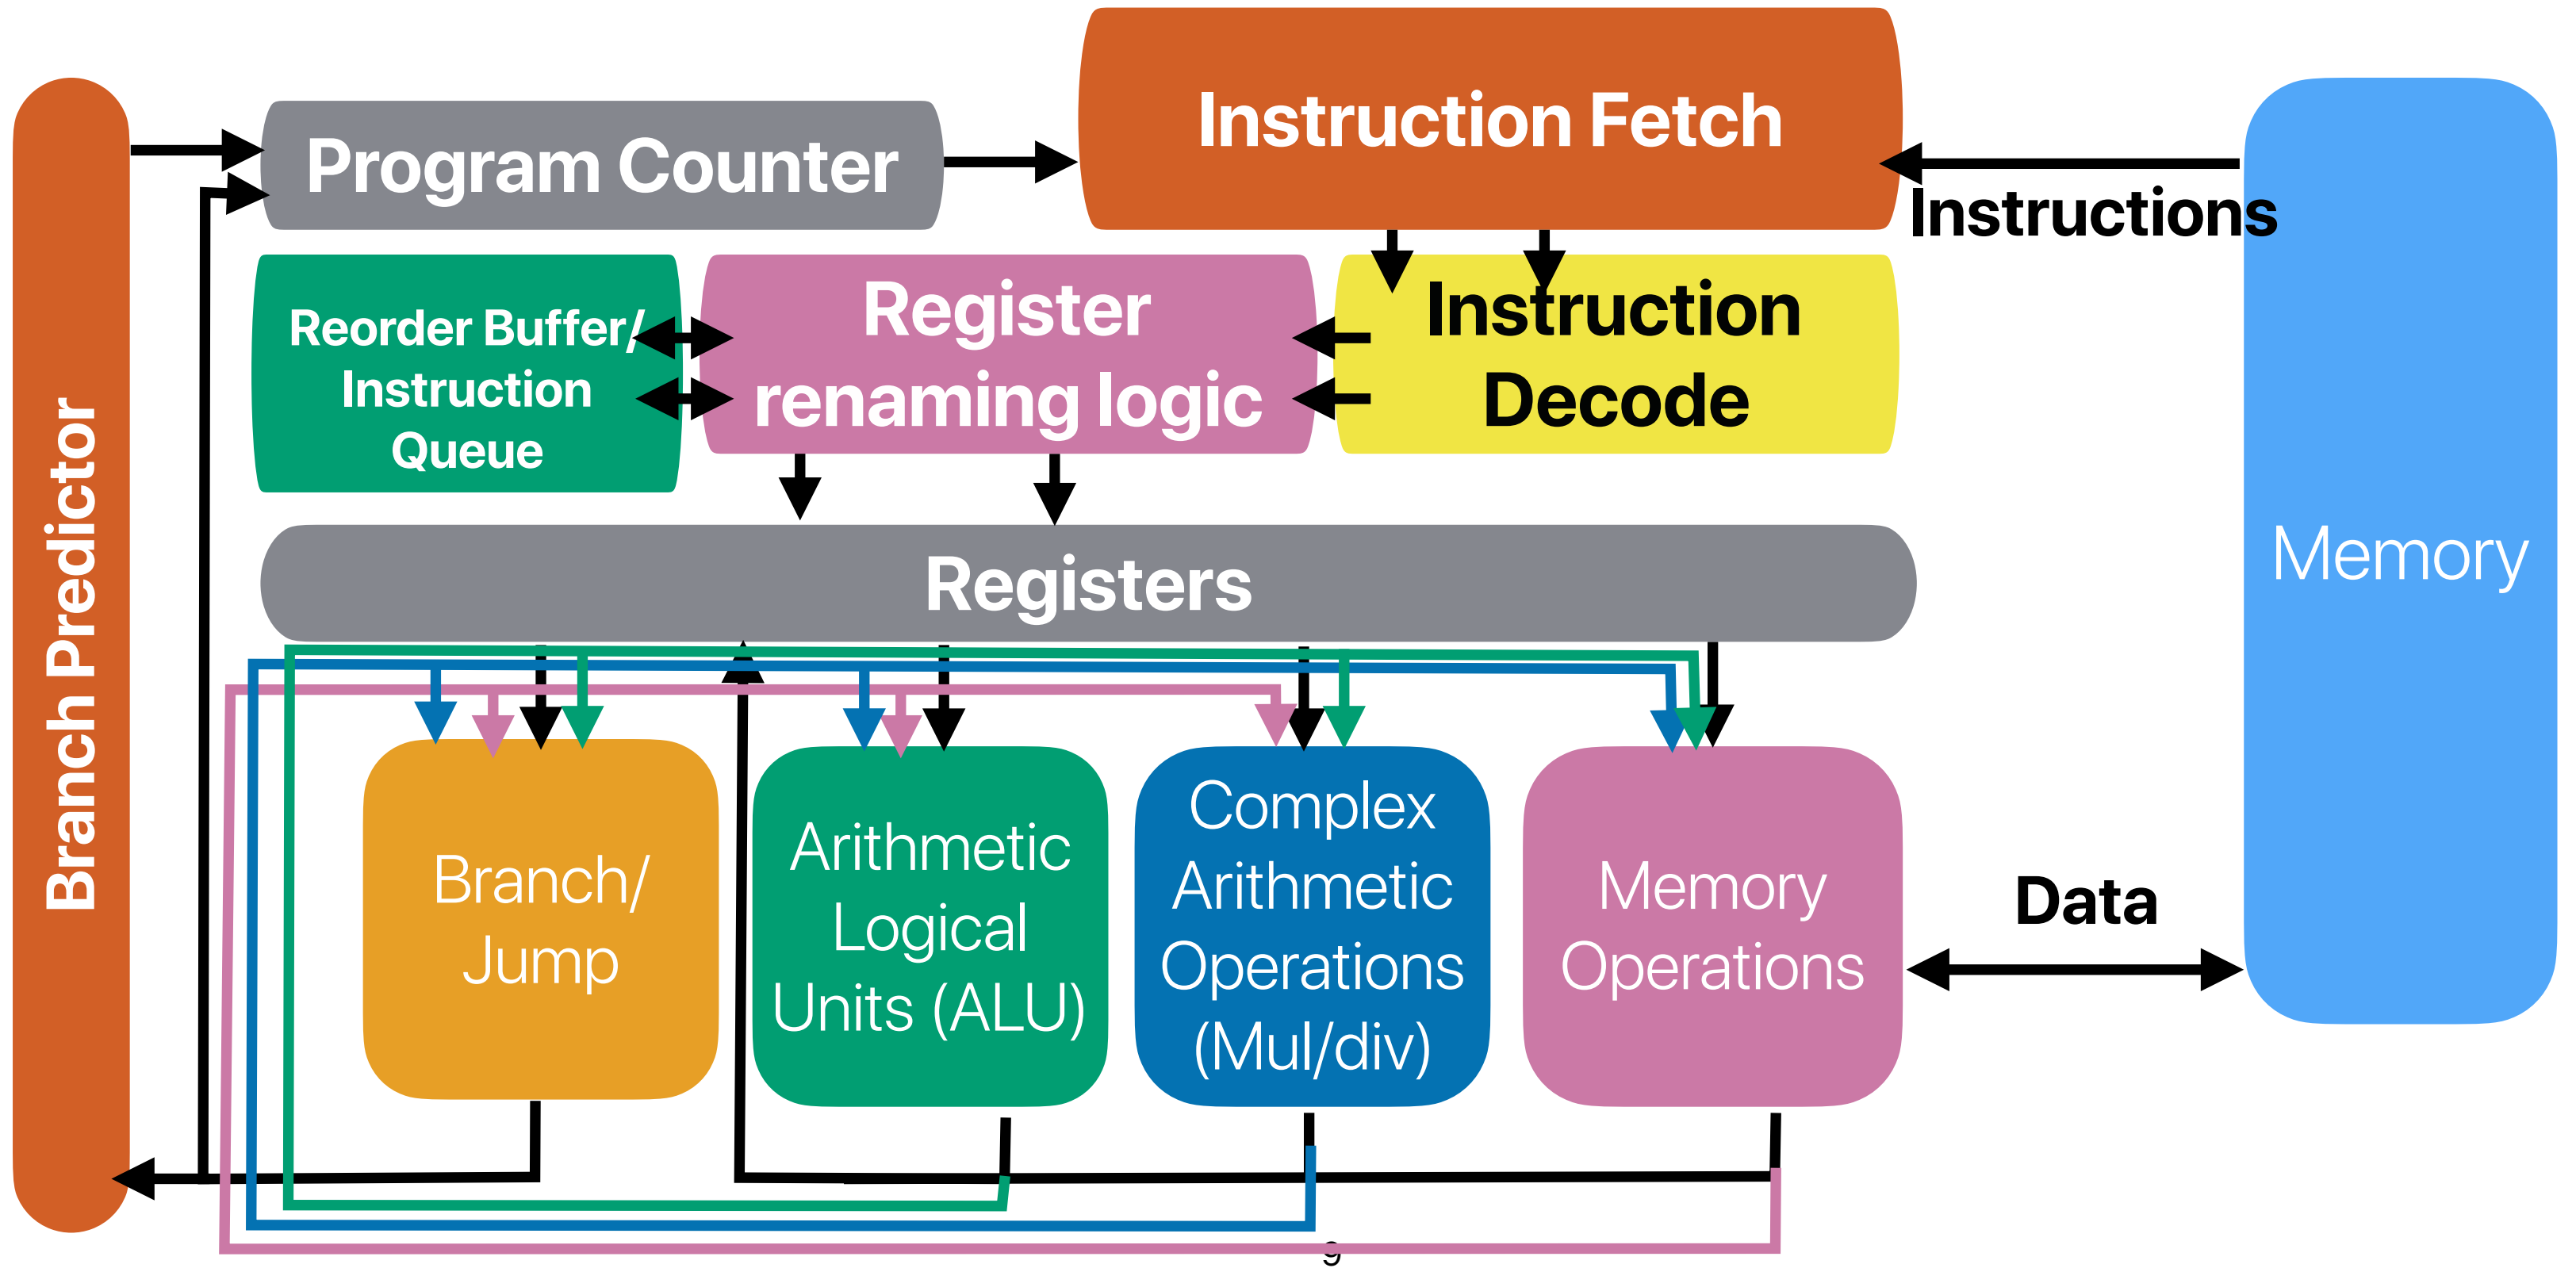
```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```



```
① movl    (%rdi), %ecx
② addq    $4, %rdi, %t0
③ addl    %ecx, %eax, %t1
④ cmpq    %rdx, %t0
⑤ jne     .L3
⑥ movl    (%t0), %t2
⑦ addq    $4, %t0, %t3
⑧ addl    %t1, %t2, %t4
⑨ cmpq    %rdx, %t3
⑩ jne     .L3
```

All false dependencies are gone!!!

Register renaming



Speculative Execution

- Exceptions (e.g. divided by 0, page fault) may occur anytime
 - A later instruction cannot write back its own result otherwise the architectural states won't be correct
- Hardware can schedule instruction across branch instructions with the help of branch prediction
 - Fetch instructions according to the branch prediction
 - However, branch predictor can never be perfect
- Execute instructions across branches
 - Speculative execution: execute an instruction before the processor know if we need to execute or not
 - Execute an instruction all operands are ready (the values of depending physical registers are generated)
 - Store results in **reorder buffer** before the processor knows if the instruction is going to be executed or not.

Register renaming in motion

① movl (%rdi), %ecx
② addq \$4, %rdi
③ addl %ecx, %eax
④ cmpq %rdx, %rdi
⑤ jne .L3
⑥ movl (%rdi), %ecx
⑦ addq \$4, %rdi
⑧ addl %ecx, %eax
⑨ cmpq %rdx, %rdi
⑩ jne .L3

(4) is now executing before (3)!

Assume issue width == 2, can
only put 2 instructions into
execution

IF	ID	REN	M1	M2	WB
IF	ID	REN	EX	-	WB
	IF	ID	REN	REN	EX
	IF	ID	REN	EX	-
		IF	ID	REN	BR
		IF	ID	REN	REN
			IF	ID	REN
			IF	ID	REN
				IF	ID
				IF	ID

ing before (3)!

depth == 2, can
actions into
on

CPI == 0.5!

Renamed instruction	
1	movl (%rdi), P1
2	addq \$4,%rdi, P2
3	addl P1, %eax, P3
4	cmpq %rdx, P2
5	jne .L3
6	movl (P2), P4
7	addq \$4, P2, P5
8	addl P4, P3, P6
9	cmpq %rdx, P5
10	jne .L3

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	1		1	P9			
P5	1		1	P10			

Outline

- Out-of-order, Dynamic instruction scheduling
- Programming on Modern Processor

Register renaming

2-issue: Only 2 of them can have instructions

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩ jne     .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3)(4)	(1) (2)								
3	(5)(6)	(3)(4)	(1) (2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5										
6										
7										
8										
9										
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions

① `movl (%rdi), %ecx`
② `addq $4, %rdi`
③ `addl %ecx, %eax`
④ `cmpq %rdx, %rdi`
⑤ `jne .L3`
⑥ `movl (%rdi), %ecx`
⑦ `addq $4, %rdi`
⑧ `addl %ecx, %eax`
⑨ `cmpq %rdx, %rdi`
⑩ `jne .L3`

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3) (4)	(1) (2)								
3	(5) (6)	(3) (4)	(1) (2)							
4	(7) (8)	(5) (6)	(3) (4)	(1)		(2)				
5					(1)	(4)				
6										
7										
8										
9										
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩ jne     .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3)(4)	(1) (2)								
3	(5)(6)	(3)(4)	(1) (2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6										
7										
8										
9										
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions

① `movl (%rdi), %ecx`
② `addq $4, %rdi`
③ `addl %ecx, %eax`
④ `cmpq %rdx, %rdi`
⑤ `jne .L3`
⑥ `movl (%rdi), %ecx`
⑦ `addq $4, %rdi`
⑧ `addl %ecx, %eax`
⑨ `cmpq %rdx, %rdi`
⑩ `jne .L3`

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3)(4)	(1) (2)								
3	(5)(6)	(3)(4)	(1) (2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7										
8										
9										
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3)(4)	(1) (2)								
3	(5)(6)	(3)(4)	(1) (2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7			(9)(10)	(6)		(7)	(3)		(5)	
8										
9										
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions

①	movl	(%rdi), %ecx									
②	addq	\$4, %rdi									
③	addl	%ecx, %eax									
④	cmpq	%rdx, %rdi									
⑤	jne	.L3									
⑥	movl	(%rdi), %ecx									
⑦	addq	\$4, %rdi									
⑧	addl	%ecx, %eax									
⑨	cmpq	%rdx, %rdi									
⑩	jne	.L3									

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3)(4)	(1) (2)								
3	(5)(6)	(3)(4)	(1) (2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7			(9)(10)	(6)		(7)	(3)		(5)	
8					(6)	(9)				(3)(4)
9										
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions

① `movl (%rdi), %ecx`
② `addq $4, %rdi`
③ `addl %ecx, %eax`
④ `cmpq %rdx, %rdi`
⑤ `jne .L3`
⑥ `movl (%rdi), %ecx`
⑦ `addq $4, %rdi`
⑧ `addl %ecx, %eax`
⑨ `cmpq %rdx, %rdi`
⑩ `jne .L3`

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1)(2)									
2	(3)(4)	(1)(2)								
3	(5)(6)	(3)(4)	(1)(2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7			(9)(10)	(6)		(7)	(3)		(5)	
8					(6)	(9)				(3)(4)
9						(8)		(10)		(5)(6)
10										
11										

Register renaming

2-issue: Only 2 of them can have instructions



```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1)(2)									
2	(3)(4)	(1)(2)								
3	(5)(6)	(3)(4)	(1)(2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7			(9)(10)	(6)		(7)	(3)		(5)	
8					(6)	(9)				(3)(4)
9						(8)		(10)		(5)(6)
10										(7)(8)
11										

Register renaming

2-issue: Only 2 of them can have instructions


```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1)(2)									
2	(3)(4)	(1)(2)								
3	(5)(6)	(3)(4)	(1)(2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7			(9)(10)	(6)		(7)	(3)		(5)	
8					(6)	(9)				(3)(4)
9						(8)		(10)		(5)(6)
10										(7)(8)
11										(9)(10)

CPI ==
0.5!

Through data flow graph analysis

```
① movl (%rdi), %ecx
② addq $4, %rdi
③ addl %ecx, %eax
④ cmpq %rdx, %rdi
⑤ jne .L3
⑥ movl (%rdi), %ecx
⑦ addq $4, %rdi
⑧ addl %ecx, %eax
⑨ cmpq %rdx, %rdi
⑩ jne .L3
⑪ movl (%rdi), %ecx
⑫ addq $4, %rdi
⑬ addl %ecx, %eax
⑭ cmpq %rdx, %rdi
⑮ jne .L3
⑯ movl (%rdi), %ecx
⑰ addq $4, %rdi
⑱ addl %ecx, %eax
⑲ cmpq %rdx, %rdi
⑳ jne .L3
㉑ movl (%rdi), %ecx
```

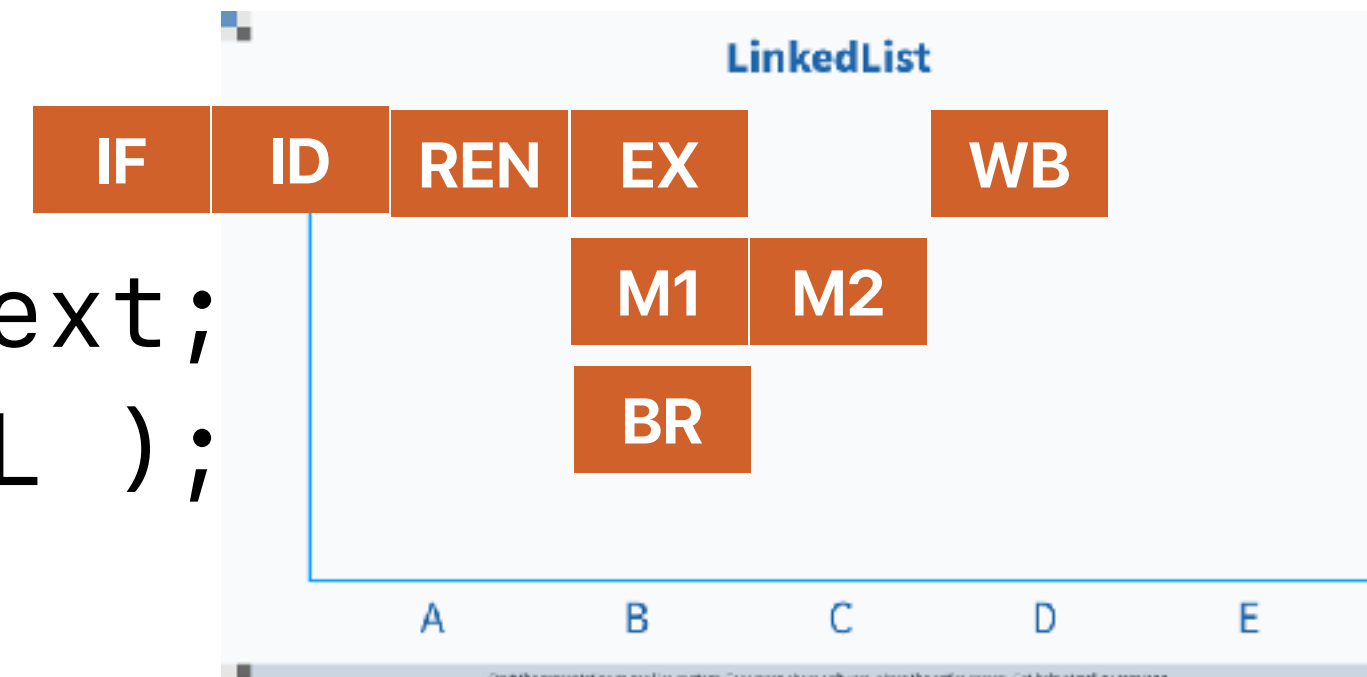


What about "linked list"

- Assume the current PC is already at instruction (1) and this linked list has only three nodes. This processor can fetch and issue 2 instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

Which of the following C state of the code snippet determines the performance?

- A. do {
- B. number_of_nodes++;
- C. current = current->next;
- D. } while (current != NULL);



Register renaming

2-issue: Only 2 of them can have instructions

```
① .L3:addq $8, %rdi
②      movq (%rdi), %rdi
③      addl $1, %eax
④      testq %rdi, %rdi
⑤      jne .L3
⑥ .L3:addq $8, %rdi
⑦      movq (%rdi), %rdi
⑧      addl $1, %eax
⑨      testq %rdi, %rdi
⑩     jne .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1) (2)									
2	(3)(4)	(1) (2)								
3	(5)(6)	(3)(4)	(1) (2)							
4	(7)(8)	(5)(6)	(3)(4)			(1)				
5	(9)(10)	(7)(8)	(5)(6)	(2)		(3)	(1)			
6		(9)(10)	(7)(8)		(2)		(3)			(1)
7			(9)(10)			(4)				(2)
8						(6)	(4)	(5)		(3)
9				(7)			(6)		(5)	(4)
10					(7)	(8)				(5)(6)
11						(9)	(8)			(7)
							(9)	(10)		(8)
									(10)	(9)(10)

Static instructions

- ①
- ②
- ③
- ④
- ⑤

Dynamic instructions

- ①
- ②
- ③
- ④
- ⑤
- ⑥
- ⑦
- ⑧
- ⑨
- ⑩
- ⑪
- ⑫
- ⑬
- ⑭
- ⑮



What about "linked list"

Performance determined by the critical path


4 cycles each iteration

5 instructions per iteration

$$CPI = \frac{4}{5} = 0.8$$

```
do {
    number_of_nodes++;
    current = current->next;
} while ( current != NULL );
```

```
① .L3:    addq    $8, %rdi
②        movq    (%rdi), %rdi
③        addl    $1, %eax
④        testq   %rdi, %rdi
⑤        jne     .L3
```



What about "linked list"

- For the following C code and it's translation in x86, **what's average CPI?** Assume the current PC is already at instruction (1) and this linked list has thousands of nodes. This processor can fetch and issue 2 instructions per cycle, with exactly the same register renaming hardware and pipeline as we showed previously.

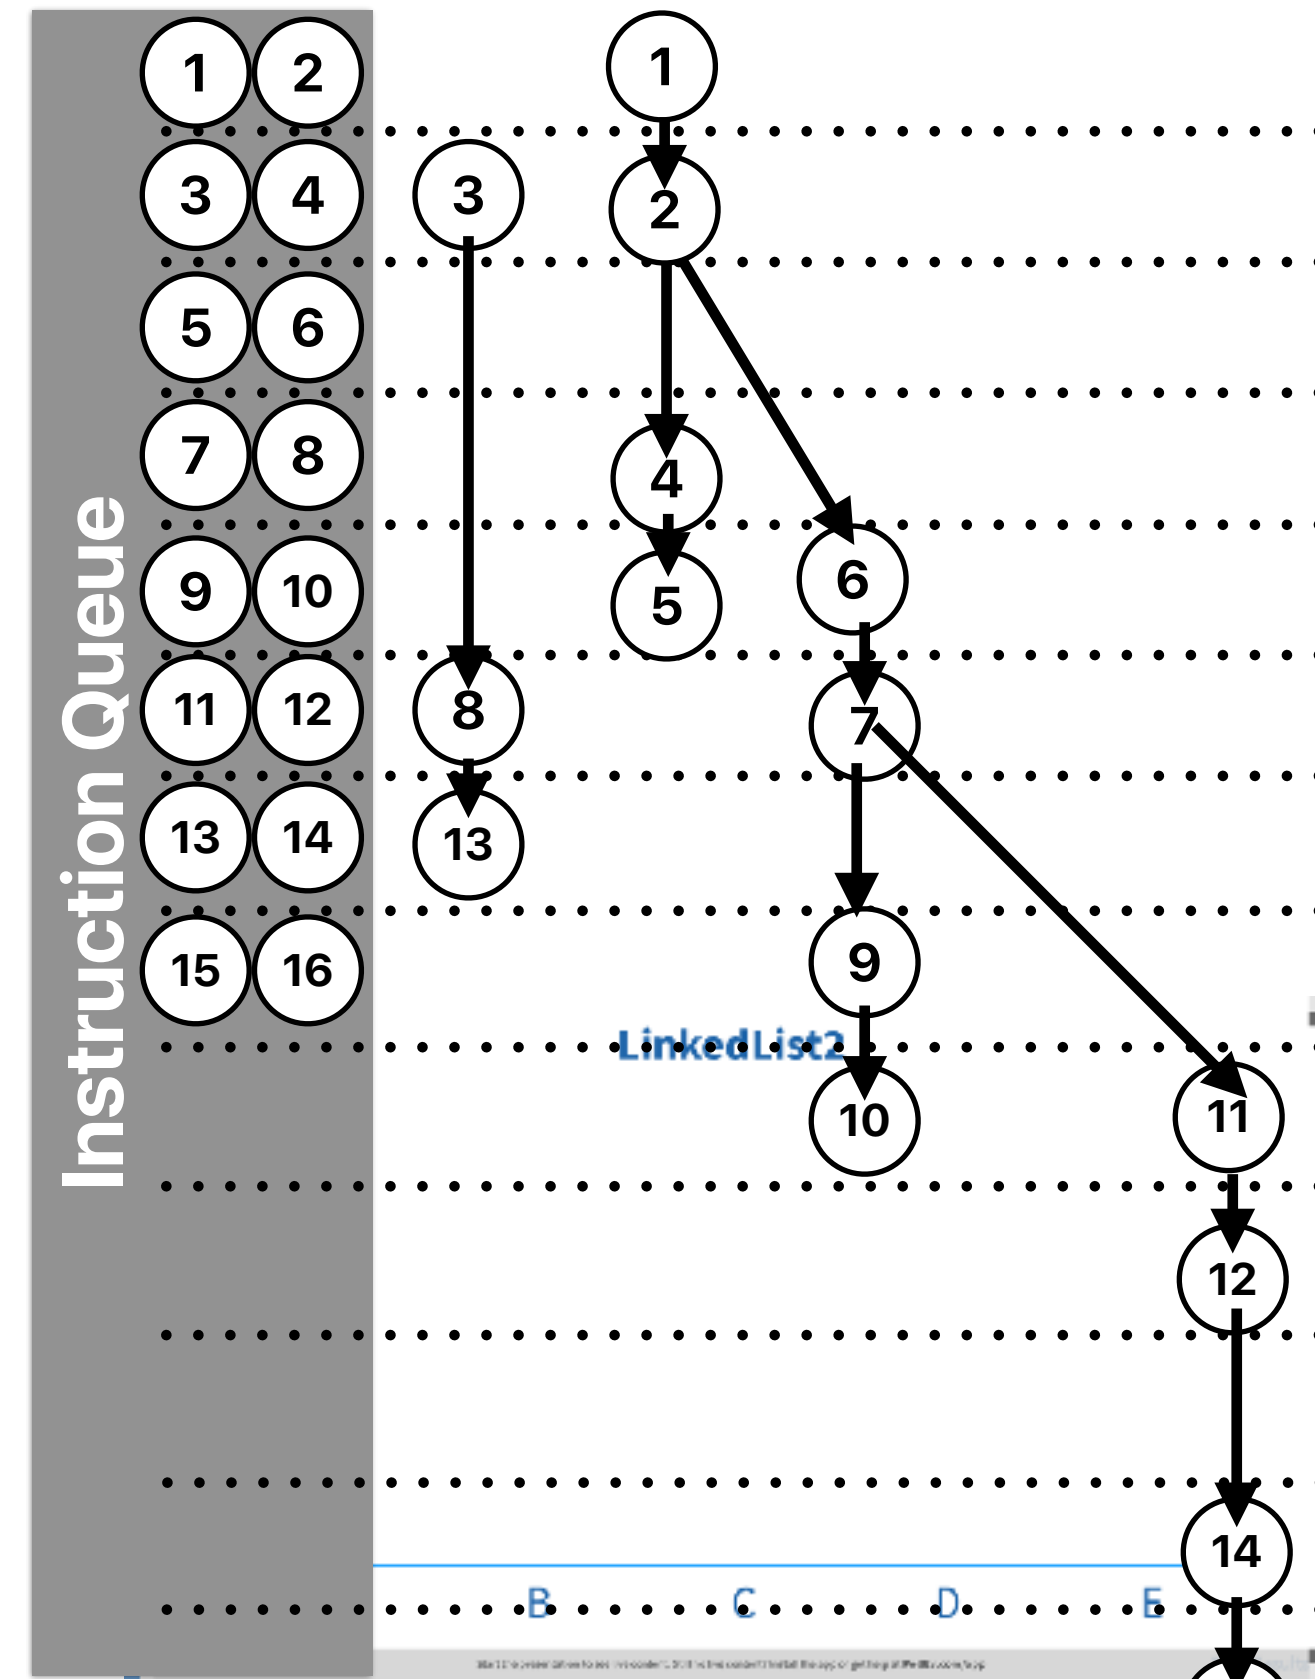
```
do {
```

```
    number_of_nodes++;
```

```
    current = current->next;
```

```
} while ( current != NULL )
```

A. 0.5	①	.L3:	addq	\$8, %rdi
B. 0.6	②		movq	(%rdi), %rdi
C. 0.7	③		addl	\$1, %eax
D. 0.8	④		testq	%rdi, %rdi
E. 0.9	⑤		jne	.L3



What about "linked list"

Performance determined by the critical path


4 cycles each iteration

5 instructions per iteration

$$CPI = \frac{4}{5} = 0.8$$

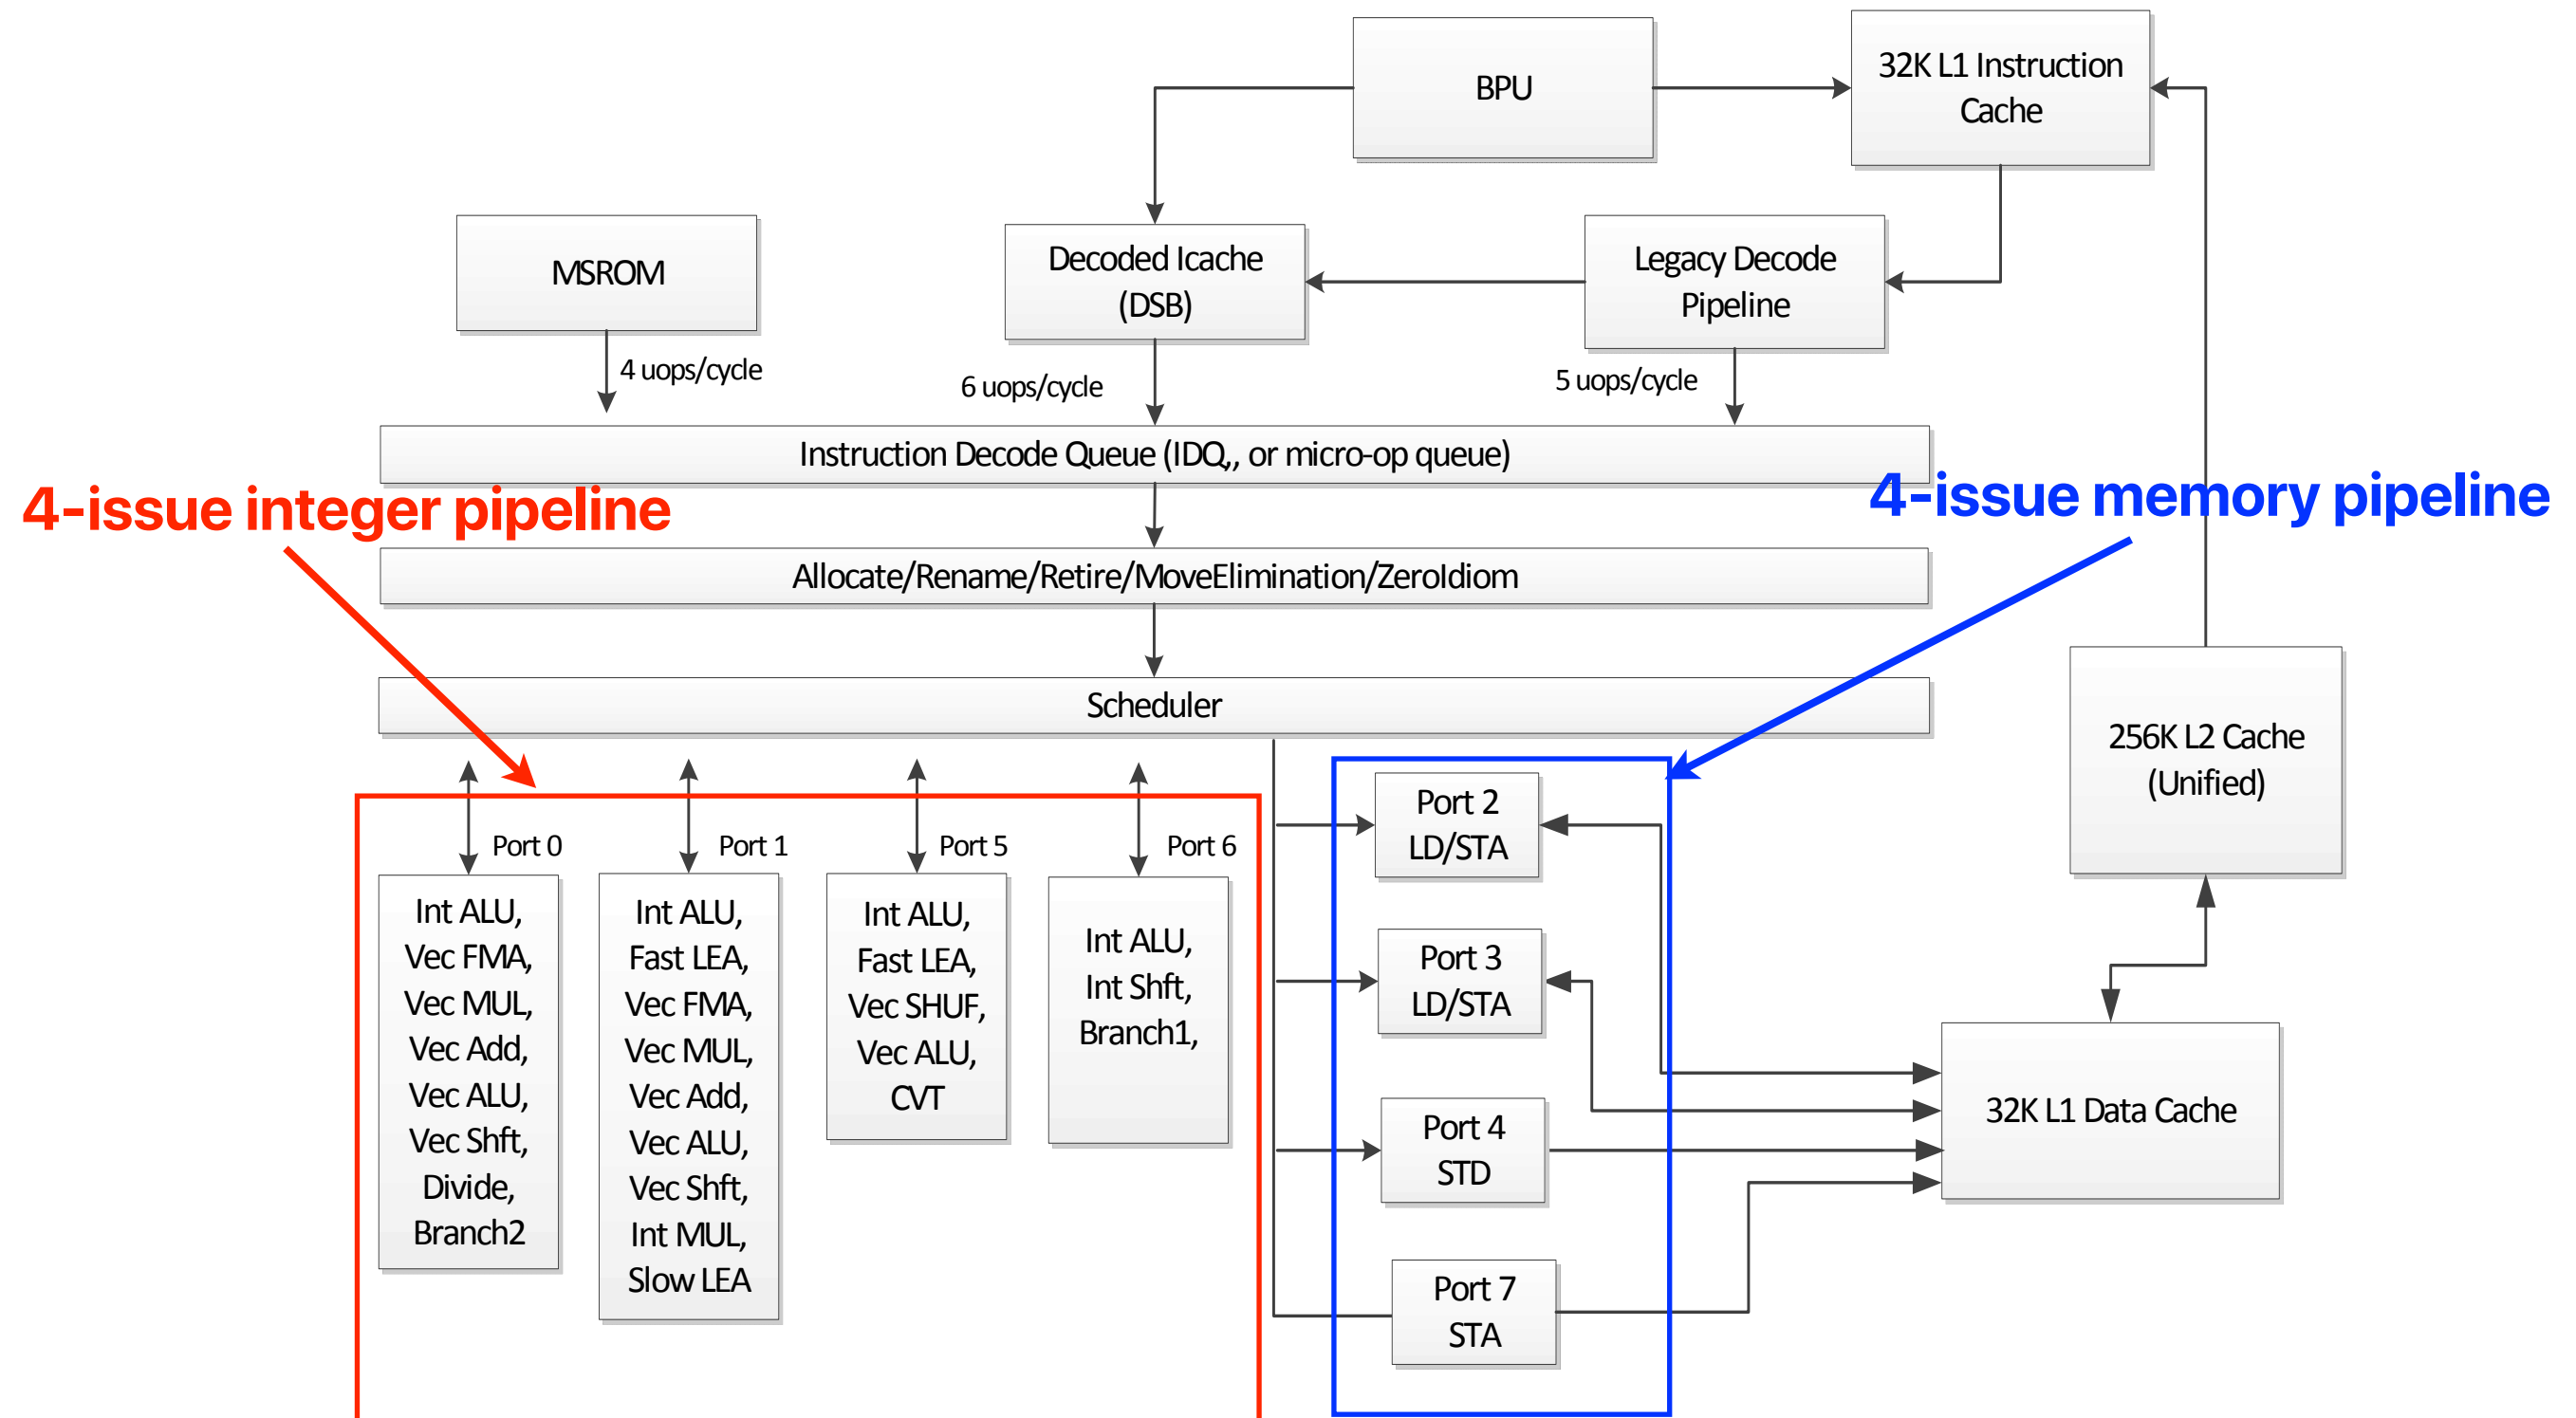
```
do {
    number_of_nodes++;
    current = current->next;
} while ( current != NULL );
```

```
① .L3:    addq    $8, %rdi
②        movq    (%rdi), %rdi
③        addl    $1, %eax
④        testq   %rdi, %rdi
⑤        jne     .L3
```



The pipelines of Modern Processors

Intel Skylake



AMD Zen 2 (RyZen 3000 Series)


3-issue memory pipeline

4-issue integer pipeline

$$MinCPI = \frac{1}{7}$$

$$MinINTInst . CPI = \frac{1}{4}$$

$$MinMEMInst . CPI = \frac{1}{3}$$



Recap: Intel Skylake


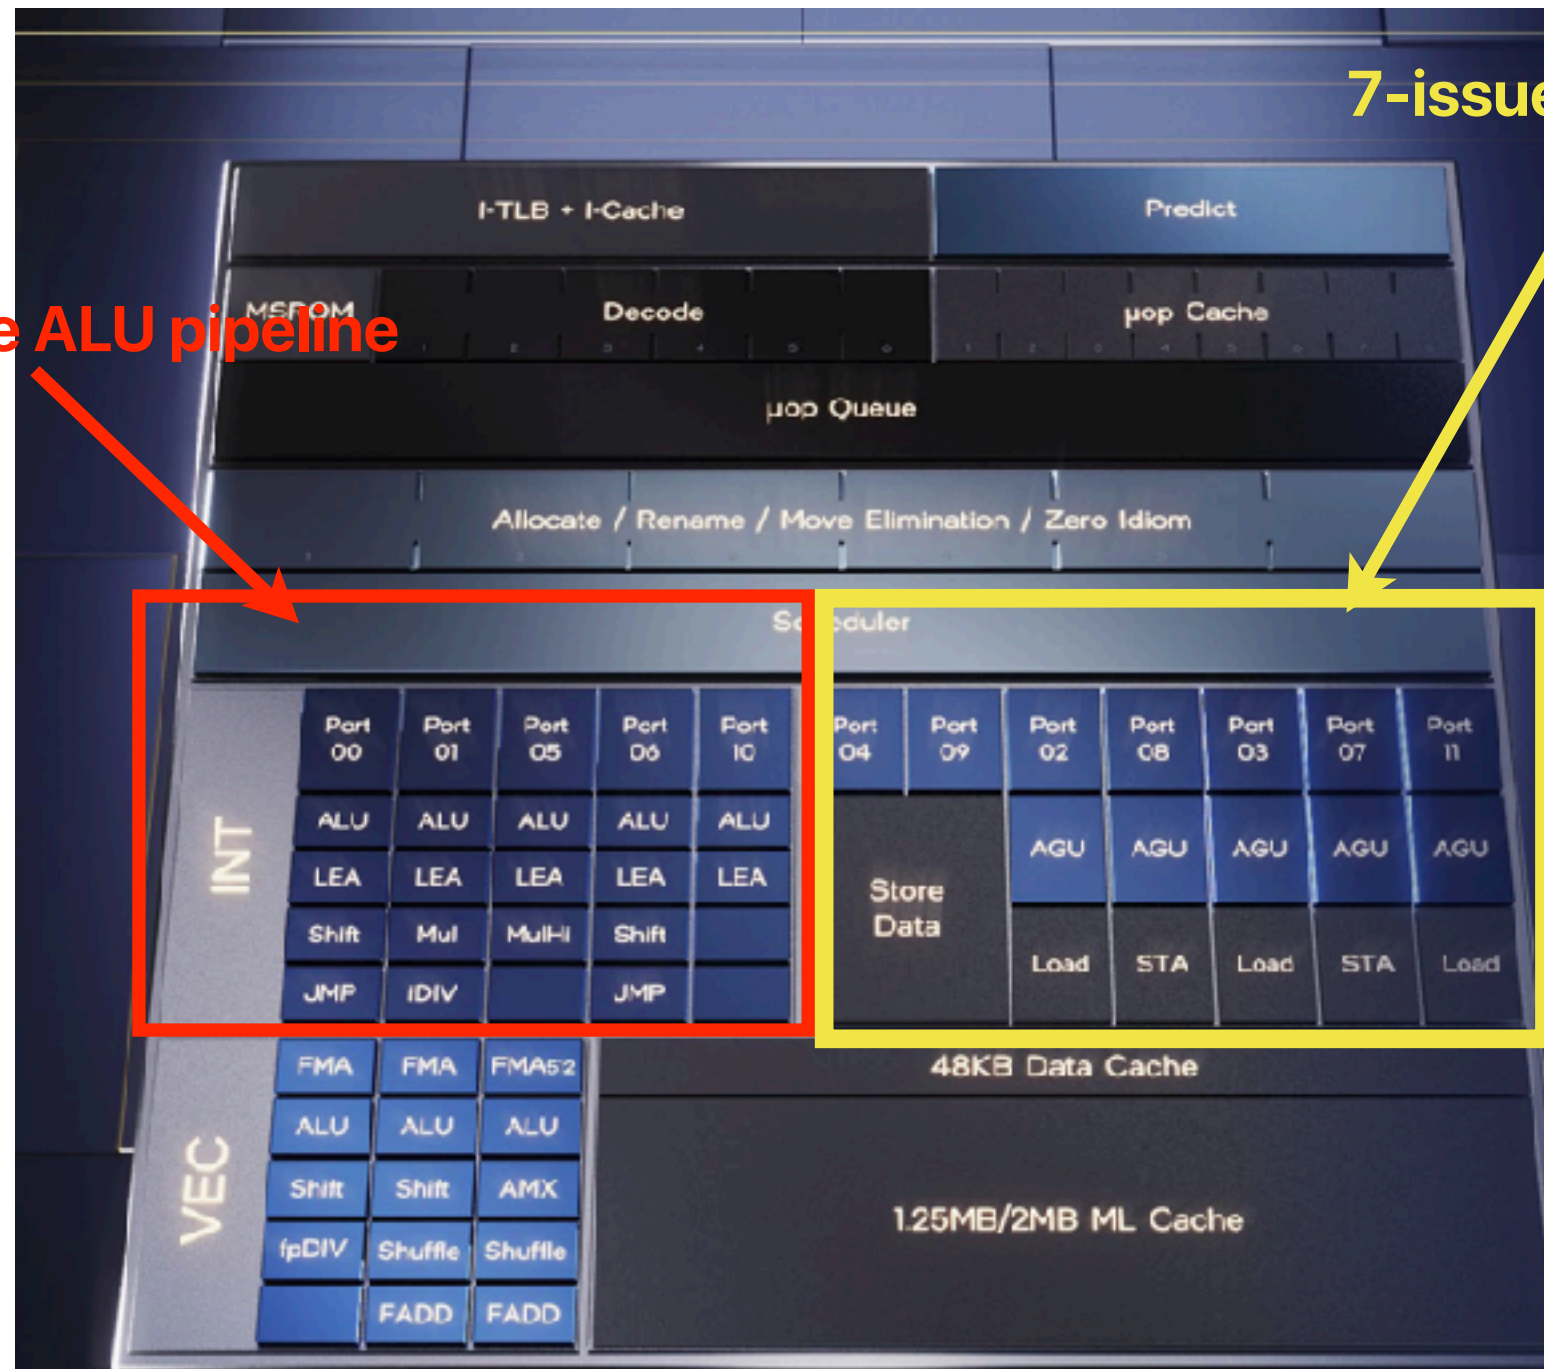


Figure 4. Skylake core block diagram.

Intel Alder Lake

5-issue ALU pipeline

7-issue memory pipeline



New

Performance

x86 Core

A Step Function in CPU Architecture
Performance For the Next Decade of
Compute

A significant IPC boost at high power efficiency

Wider

Deeper

Smarter

Better supports large data set and large code footprint applications

Enhanced power management improves frequency and power

Machine Learning Technology: Intel® AMX – Tile Multiplication

All in a tailored scalable architecture to serve the
full range of Laptops to Desktops to Data Centers

Summary: Characteristics of modern processor architectures

- Multiple-issue pipelines with multiple functional units available
 - Multiple ALUs
 - Multiple Load/store units
 - Dynamic OoO scheduling to reorder instructions whenever possible
- Cache
- Branch predictors

Performance Programming on Modern Processors

Demo: Popcount

- The population count (or popcount) of a specific value is the number of set bits (i.e., bits in 1s) in that value.
- Applications
 - Parity bits in error correction/detection code
 - Cryptography
 - Sparse matrix
 - Molecular Fingerprinting
 - Implementation of some succinct data structures like bit vectors and wavelet trees.

Demo: pop count

- Given a 64-bit integer number, find the number of 1s in its binary representation.

- Example 1:

Input: 9487

Output: 7

Explanation: 9487's binary representation is

0b10010100001111

```
int main(int argc, char *argv[]) {  
  
    uint64_t key = 0xdeadbeef;  
  
    int count = 1000000000;  
    uint64_t sum = 0;  
  
    for (int i=0; i < count; i++)  
    {  
        sum += popcount(RandLFSR(key));  
    }  
    printf("Result: %lu\n", sum);  
    return sum;  
}
```

Five implementations

- Which of the following implementations will perform the best on modern pipeline processors?

A

```
inline int popcount(uint64_t x){
    int c=0;
    while(x) {
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

B

```
inline int popcount(uint64_t x) {
    int c = 0;
    while(x) {
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

C

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1,
2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    while(x) {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

D

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1,
2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    for (uint64_t i = 0; i < 16; i++)
    {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

E

```
inline int popcount(uint64_t x) {
    int c = 0;
    for (uint64_t i = 0; i < 16; i++)
    {
        switch((x & 0xF))
        {
            case 1: c+=1; break;
            case 2: c+=1; break;
            case 3: c+=2; break;
            case 4: c+=1; break;
            case 5: c+=2; break;
            case 6: c+=2; break;
            case 7: c+=3; break;
            case 8: c+=1; break;
            case 9: c+=2; break;
            case 10: c+=2; break;
            case 11: c+=3; break;
            case 12: c+=2; break;
            case 13: c+=3; break;
            case 14: c+=3; break;
            case 15: c+=4; break;
            default: break;
        }
        x = x >> 4;
    }
    return c;
}
```

B

C

D

Five implementations

- Which of the following implementations will perform the best on modern pipeline processors?

A

```
inline int popcount(uint64_t x){
    int c=0;
    while(x) {
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

B

```
inline int popcount(uint64_t x) {
    int c = 0;
    while(x) {
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```

C

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1, 2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    while(x) {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

D

```
inline int popcount(uint64_t x) {
    int c = 0;
    int table[16] = {0, 1, 1, 2, 1, 2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4};
    for (uint64_t i = 0; i < 16; i++) {
        c += table[(x & 0xF)];
        x = x >> 4;
    }
    return c;
}
```

E

```
inline int popcount(uint64_t x) {
    int c = 0;
    for (uint64_t i = 0; i < 16; i++) {
        switch((x & 0xF)) {
            case 1: c+=1; break;
            case 2: c+=1; break;
            case 3: c+=2; break;
            case 4: c+=1; break;
            case 5: c+=2; break;
            case 6: c+=2; break;
            case 7: c+=3; break;
            case 8: c+=1; break;
            case 9: c+=2; break;
            case 10: c+=2; break;
            case 11: c+=3; break;
            case 12: c+=2; break;
            case 13: c+=3; break;
            case 14: c+=3; break;
            case 15: c+=4; break;
            default: break;
        }
        x = x >> 4;
    }
    return c;
}
```

Why is B better than A?

- How many of the following statements explains the reason why B outperforms A with compiler optimizations

- ① B has lower dynamic instruction count than A
- ② B has significantly lower branch mis-prediction rate than A
- ③ B has significantly fewer branch instructions than A
- ④ B can incur fewer data hazards

A. 0

B. 1

C. 2

D. 3

E. 4

A

```
inline int popcount(uint64_t x){  
    int c=0;  
    while(x) {  
        c += x & 1;  
        x = x >> 1;  
    }  
    return c;  
}
```

B

```
inline int popcount(uint64_t x) {  
    int c = 0;  
    while(x) {  
        c += x & 1;  
        x = x >> 1;  
        c += x & 1;  
        x = x >> 1;  
        c += x & 1;  
        x = x >> 1;  
    }  
    return c;  
}
```



Why is B better than A?

A

```
inline int popcount(uint64_t x){
    int c=0;
    while(x) {
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```


B

```
inline int popcount(uint64_t x) {
    int c = 0;
    while(x) {
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
        c += x & 1;
        x = x >> 1;
    }
    return c;
}
```



```
movl    %eax, %ecx
andl    $1, %ecx
addl    %ecx, %edx
shrq    %rax
jne     .L6
```

5*n instructions



15*(n/4) = 3.75*n instructions

compiler reorder
instructions and rename
registers to mitigate
hazards and exploit ILP

```
%ecx, %eax
andl    $1, %eax
addl    %edx, %eax
movq    %rcx, %rdx
shrq    %rdx
andl    $1, %edx
addl    %eax, %edx
movq    %rcx, %rax
shrq    $2, %rax
andl    $1, %eax
addl    %edx, %eax
movq    %rcx, %rdx
shrq    $3, %rdx
andl    $1, %edx
addl    %eax, %edx
shrq    $4, %rcx
jne     .L6
```

Only one branch for four iterations in A

Why is B better than A?

- How many of the following statements explains the reason why B outperforms A with compiler optimizations

- ① ✓ B has lower dynamic instruction count than A
- ② B has significantly lower branch mis-prediction rate than A
- ③ ✓ B has significantly fewer branch instructions than A
- ④ ✓ B can incur fewer data hazards

A. 0

B. 1

C. 2

D. 3

E. 4

A

```
inline int popcount(uint64_t x){  
    int c=0;  
    while(x) {  
        c += x & 1;  
        x = x >> 1;  
    }  
    return c;  
}
```

B

```
inline int popcount(uint64_t x) {  
    int c = 0;  
    while(x) {  
        c += x & 1;  
        x = x >> 1;  
        c += x & 1;  
        x = x >> 1;  
        c += x & 1;  
        x = x >> 1;  
        c += x & 1;  
        x = x >> 1;  
    }  
    return c;  
}
```

Announcements

- Assignment #3 due **Friday**
- Reading Quiz due next Monday

Computer Science & Engineering

203

つづく

