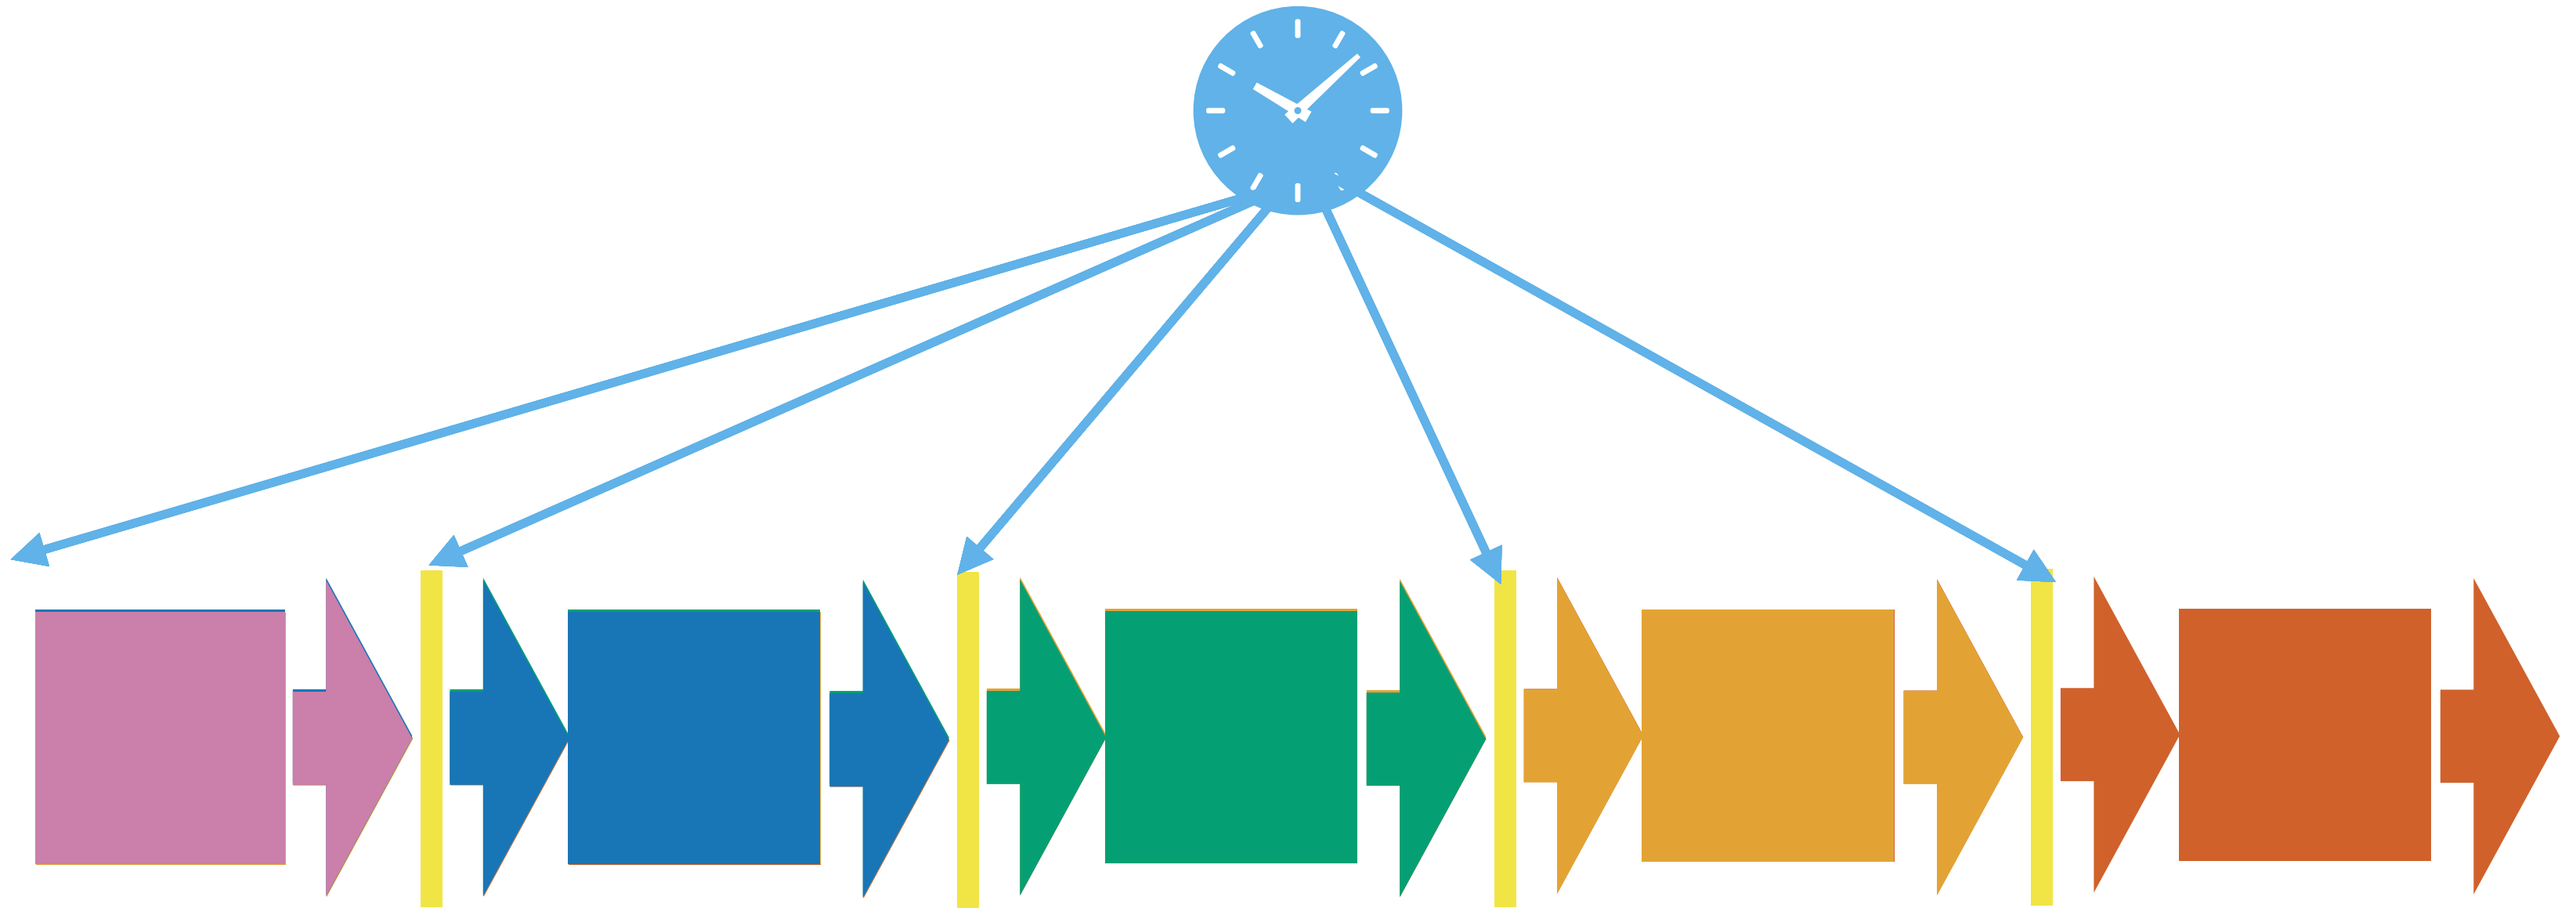


# Data Hazards & Dynamic Instruction Scheduling (I)

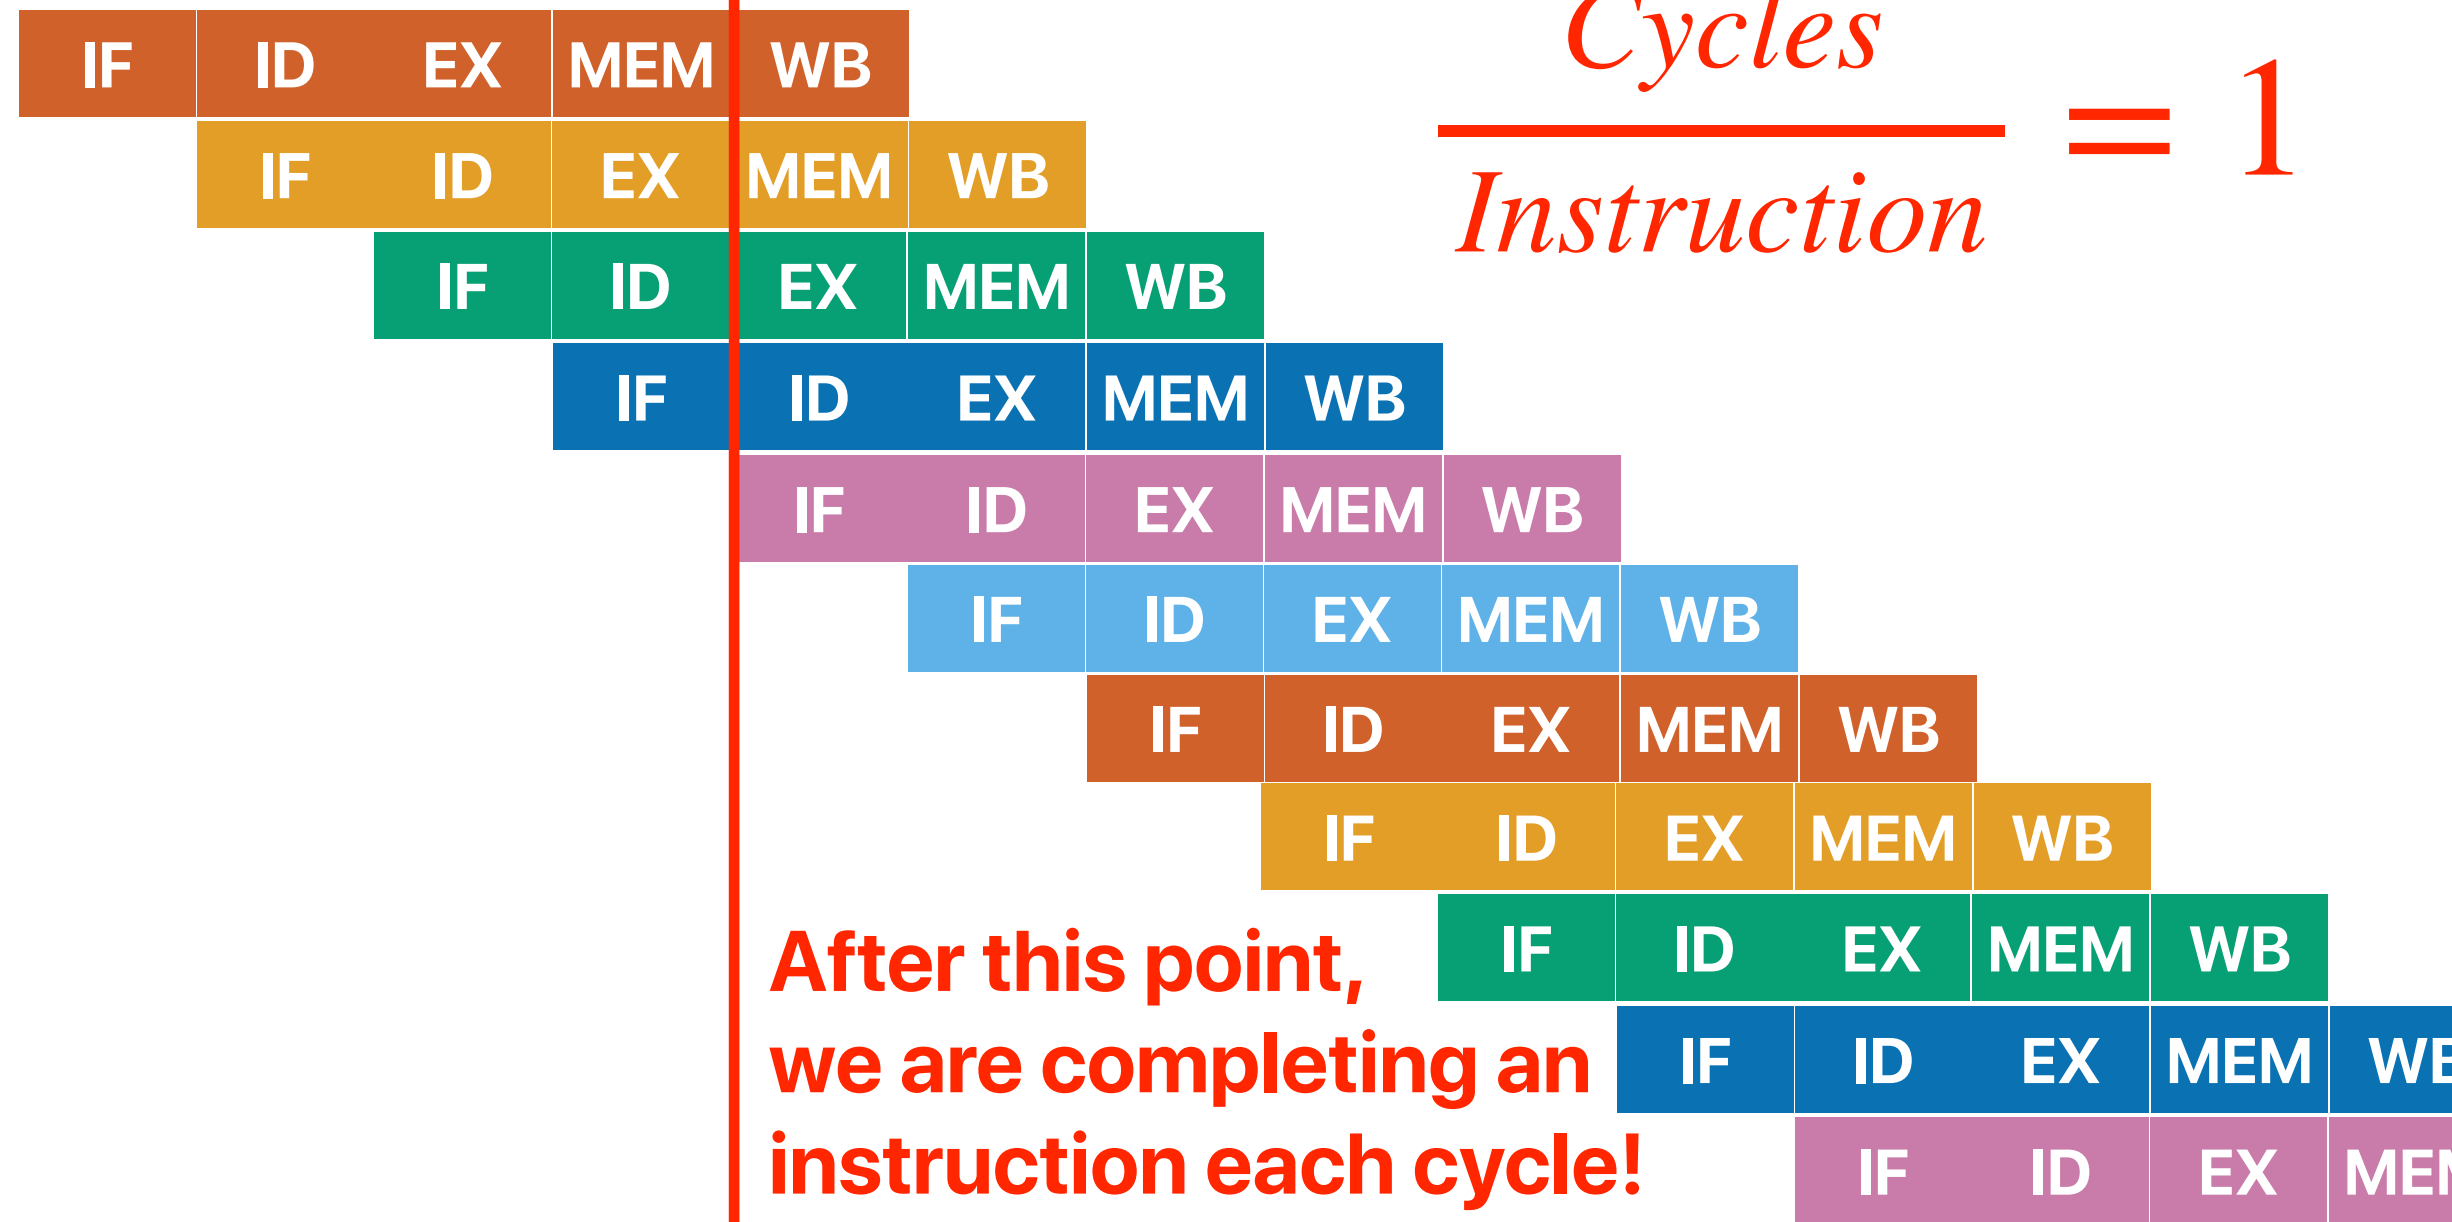
Hung-Wei Tseng

# Recap: Pipelining



# Recap: Pipelining

```
add x1, x2, x3
ld x4, 0(x5)
sub x6, x7, x8
sub x9, x10, x11
sd x1, 0(x12)
xor x13, x14, x15
and x16, x17, x18
add x19, x20, x21
sub x22, x23, x24
ld x25, 4(x26)
sd x27, 0(x28)
```



# Recap: Three pipeline hazards

- Structural hazards — resource conflicts cannot support simultaneous execution of instructions in the pipeline
- Control hazards — the PC can be changed by an instruction in the pipeline
- Data hazards — an instruction depending on a the result that's not yet generated or propagated when the instruction needs that

# Recap: addressing hazards

- Structural hazards
  - Stall
  - Modify hardware design
- Control hazards
  - Stall
  - Static prediction
  - Dynamic prediction

# Outline

- Data hazards
  - Data forwarding
- SuperScalar
- Out-of-order, Dynamic instruction scheduling

# Data hazards

# Data hazards

- An instruction currently in the pipeline cannot receive the “logically” correct value for execution
- Data dependencies
  - The output of an instruction is the input of a later instruction
  - May result in data hazard if the later instruction that consumes the result is still in the pipeline

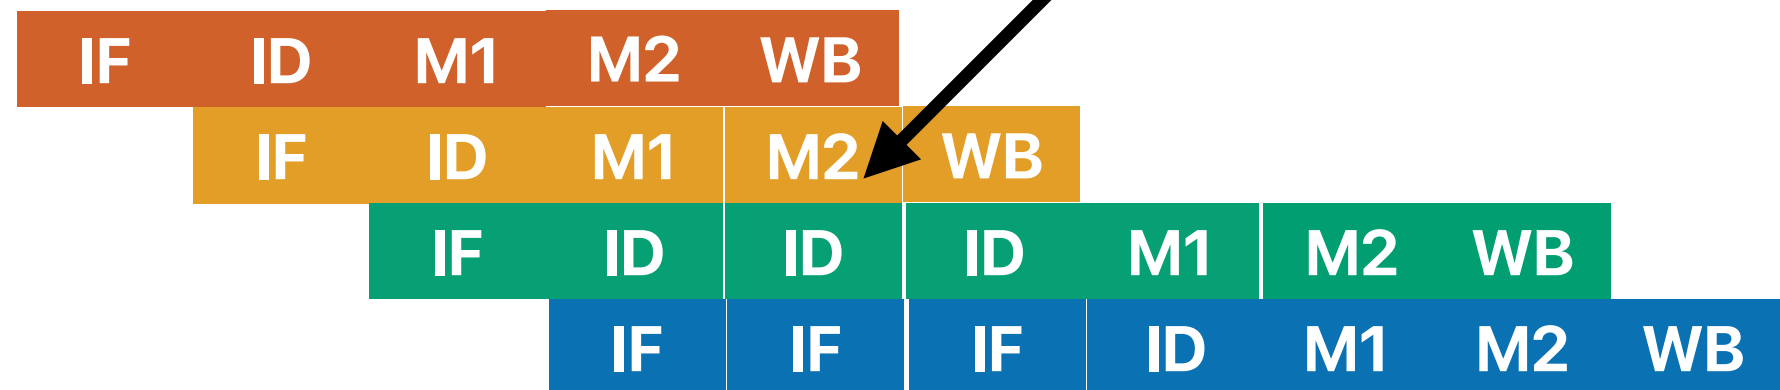


# Solution 1: Let's try "stall" again

- Whenever the input is not ready when the consumer is decoding, just stall — the consumer stays at ID.

**we have the value for %edx already! Why another cycle?**

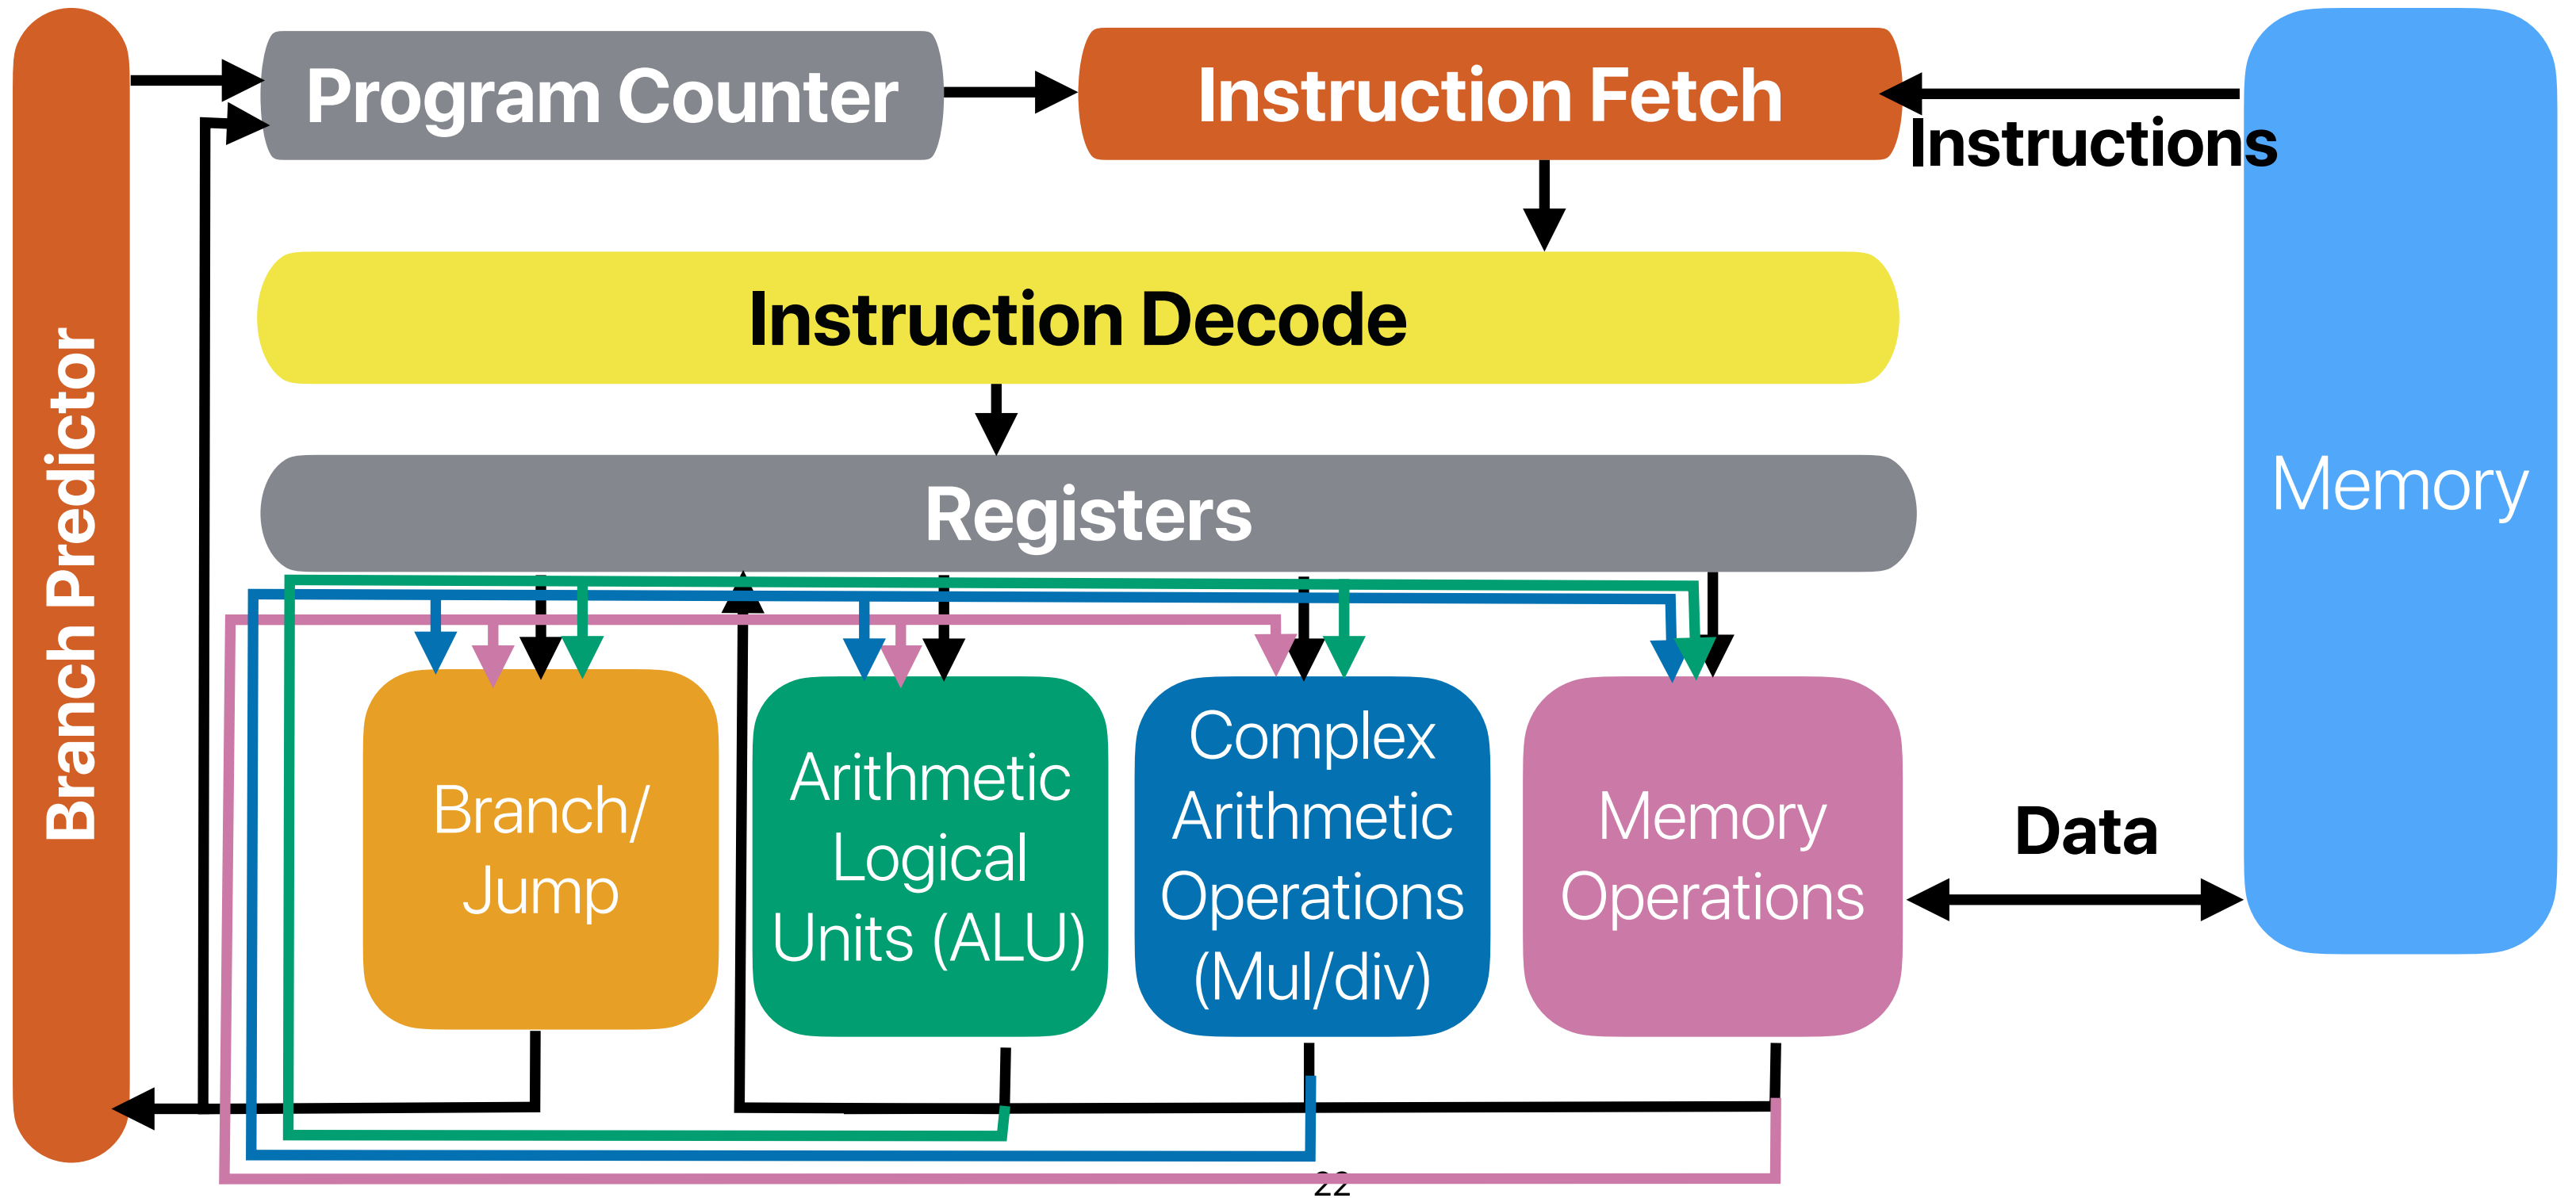
```
movl    (%rdi), %eax
movl    (%rsi), %edx
movl    %edx, (%rdi)
movl    %eax, (%rsi)
```



## Solution 2: Data forwarding

- Add logics/wires to forward the desired values to the demanding instructions

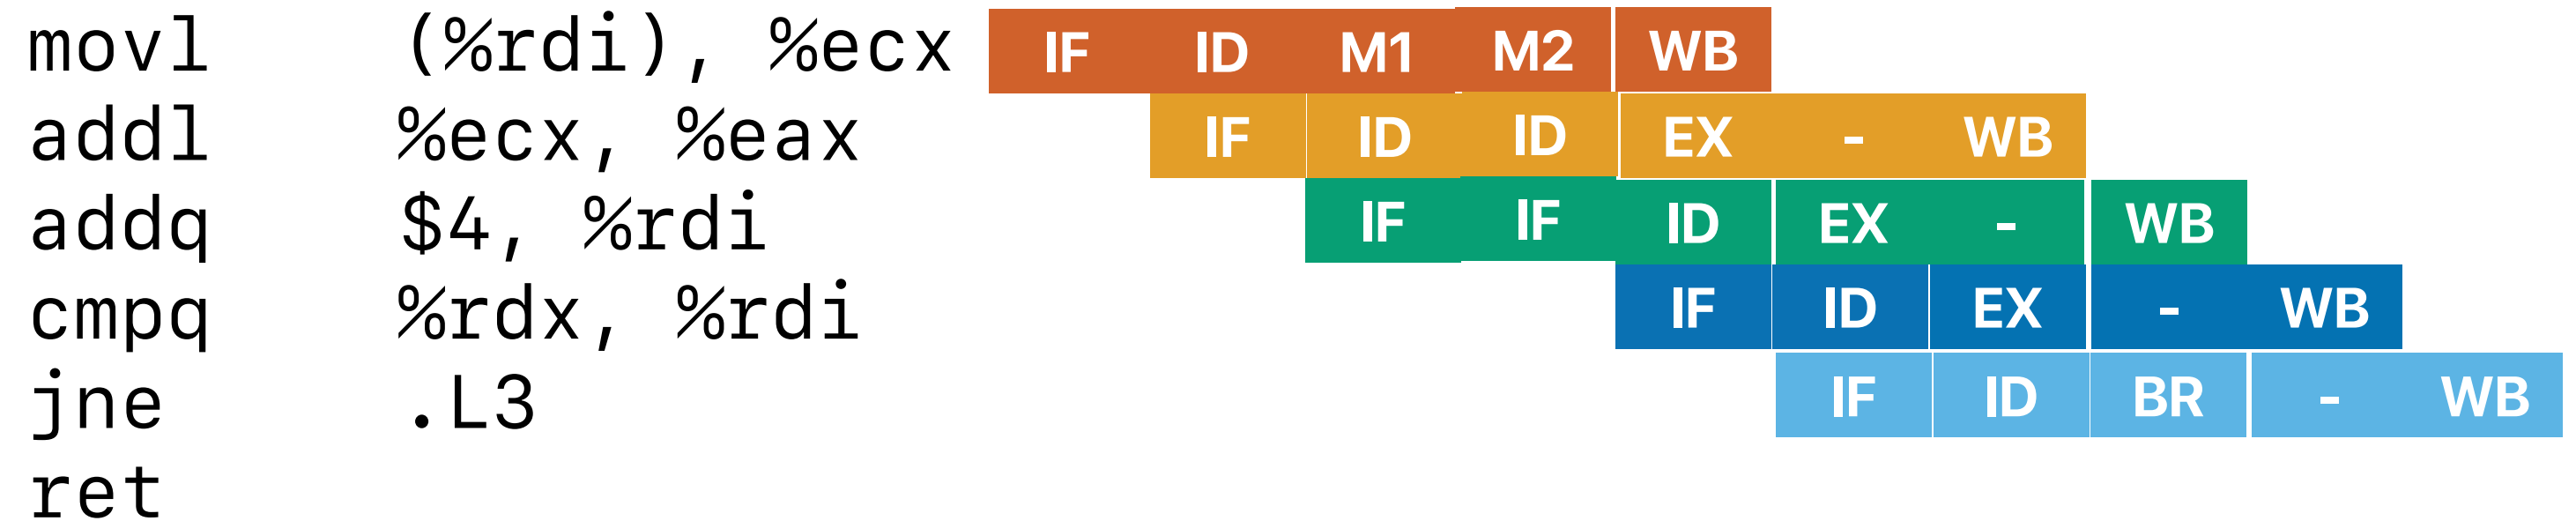
# Data "forwarding"



# Single pipeline

```
for(i = 0; i < count; i++) {
    s += a[i];
}
```

.L3:



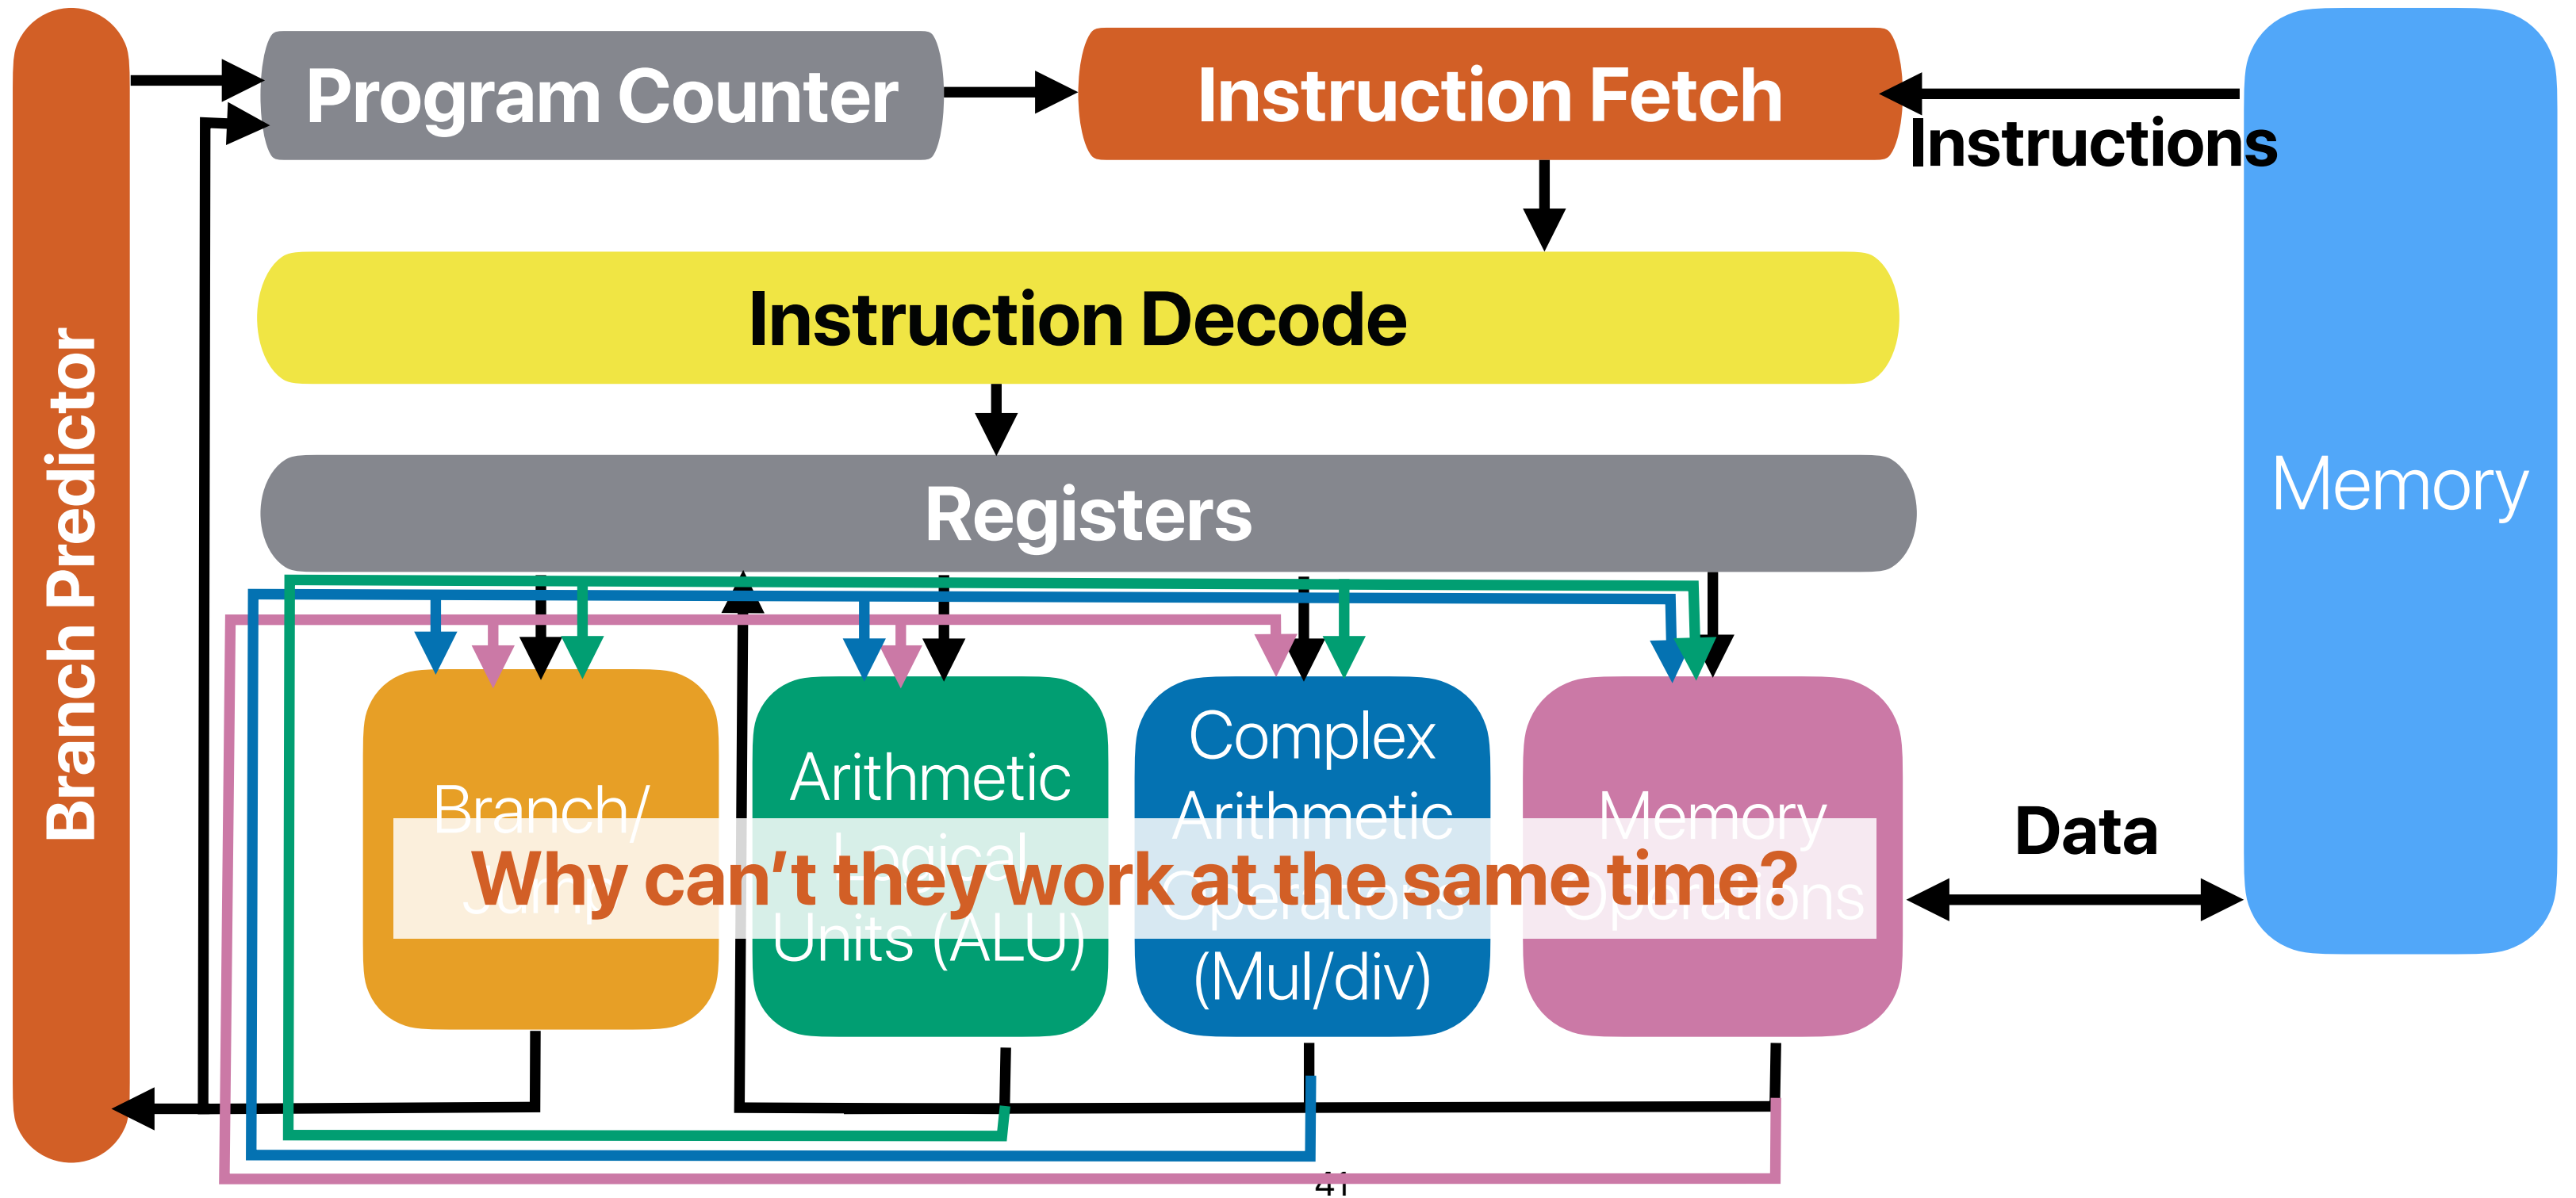
# Outline

- Make CPI go below 1 — Super Scalar
- Dynamic Out of Order execution
- Modern processor design

# Outline

- Make CPI go below 1 — Super Scalar
- Dynamic Out of Order execution
- Modern processor design

# Data "forwarding"




# Compiler optimization

```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

.L3:

```
movl    (%rdi), %ecx  
addq    $4, %rdi  
addl    %ecx, %eax  
cmpq    %rdx, %rdi  
jne     .L3  
ret
```

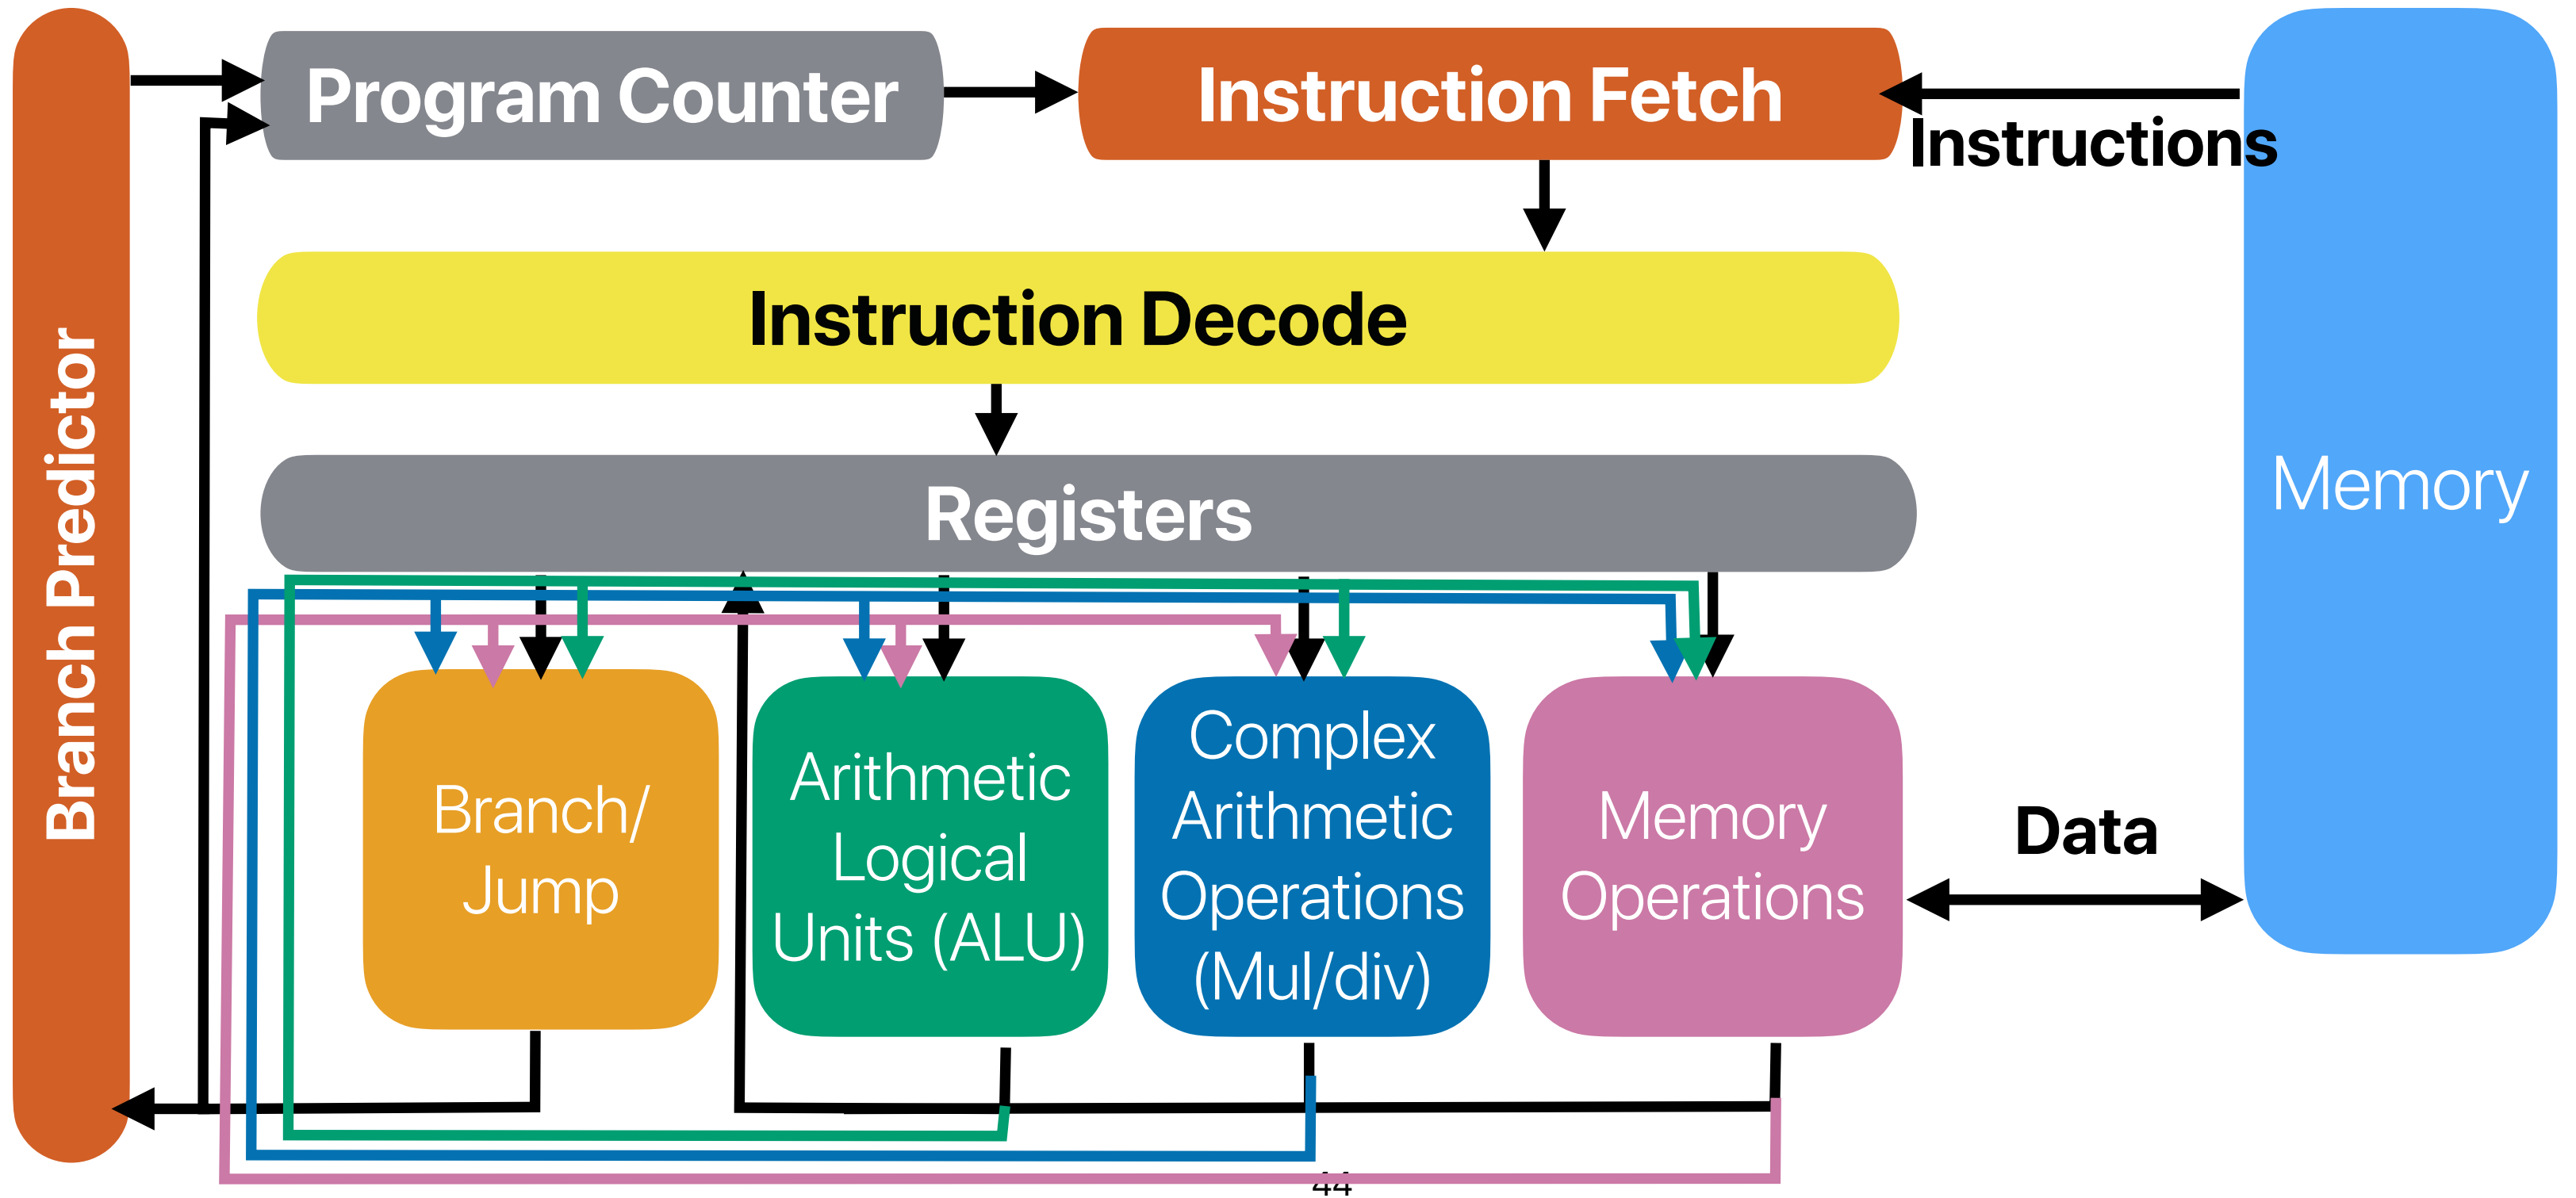


addq is not depending on movl and  
ALU is free! can we execute them  
together?

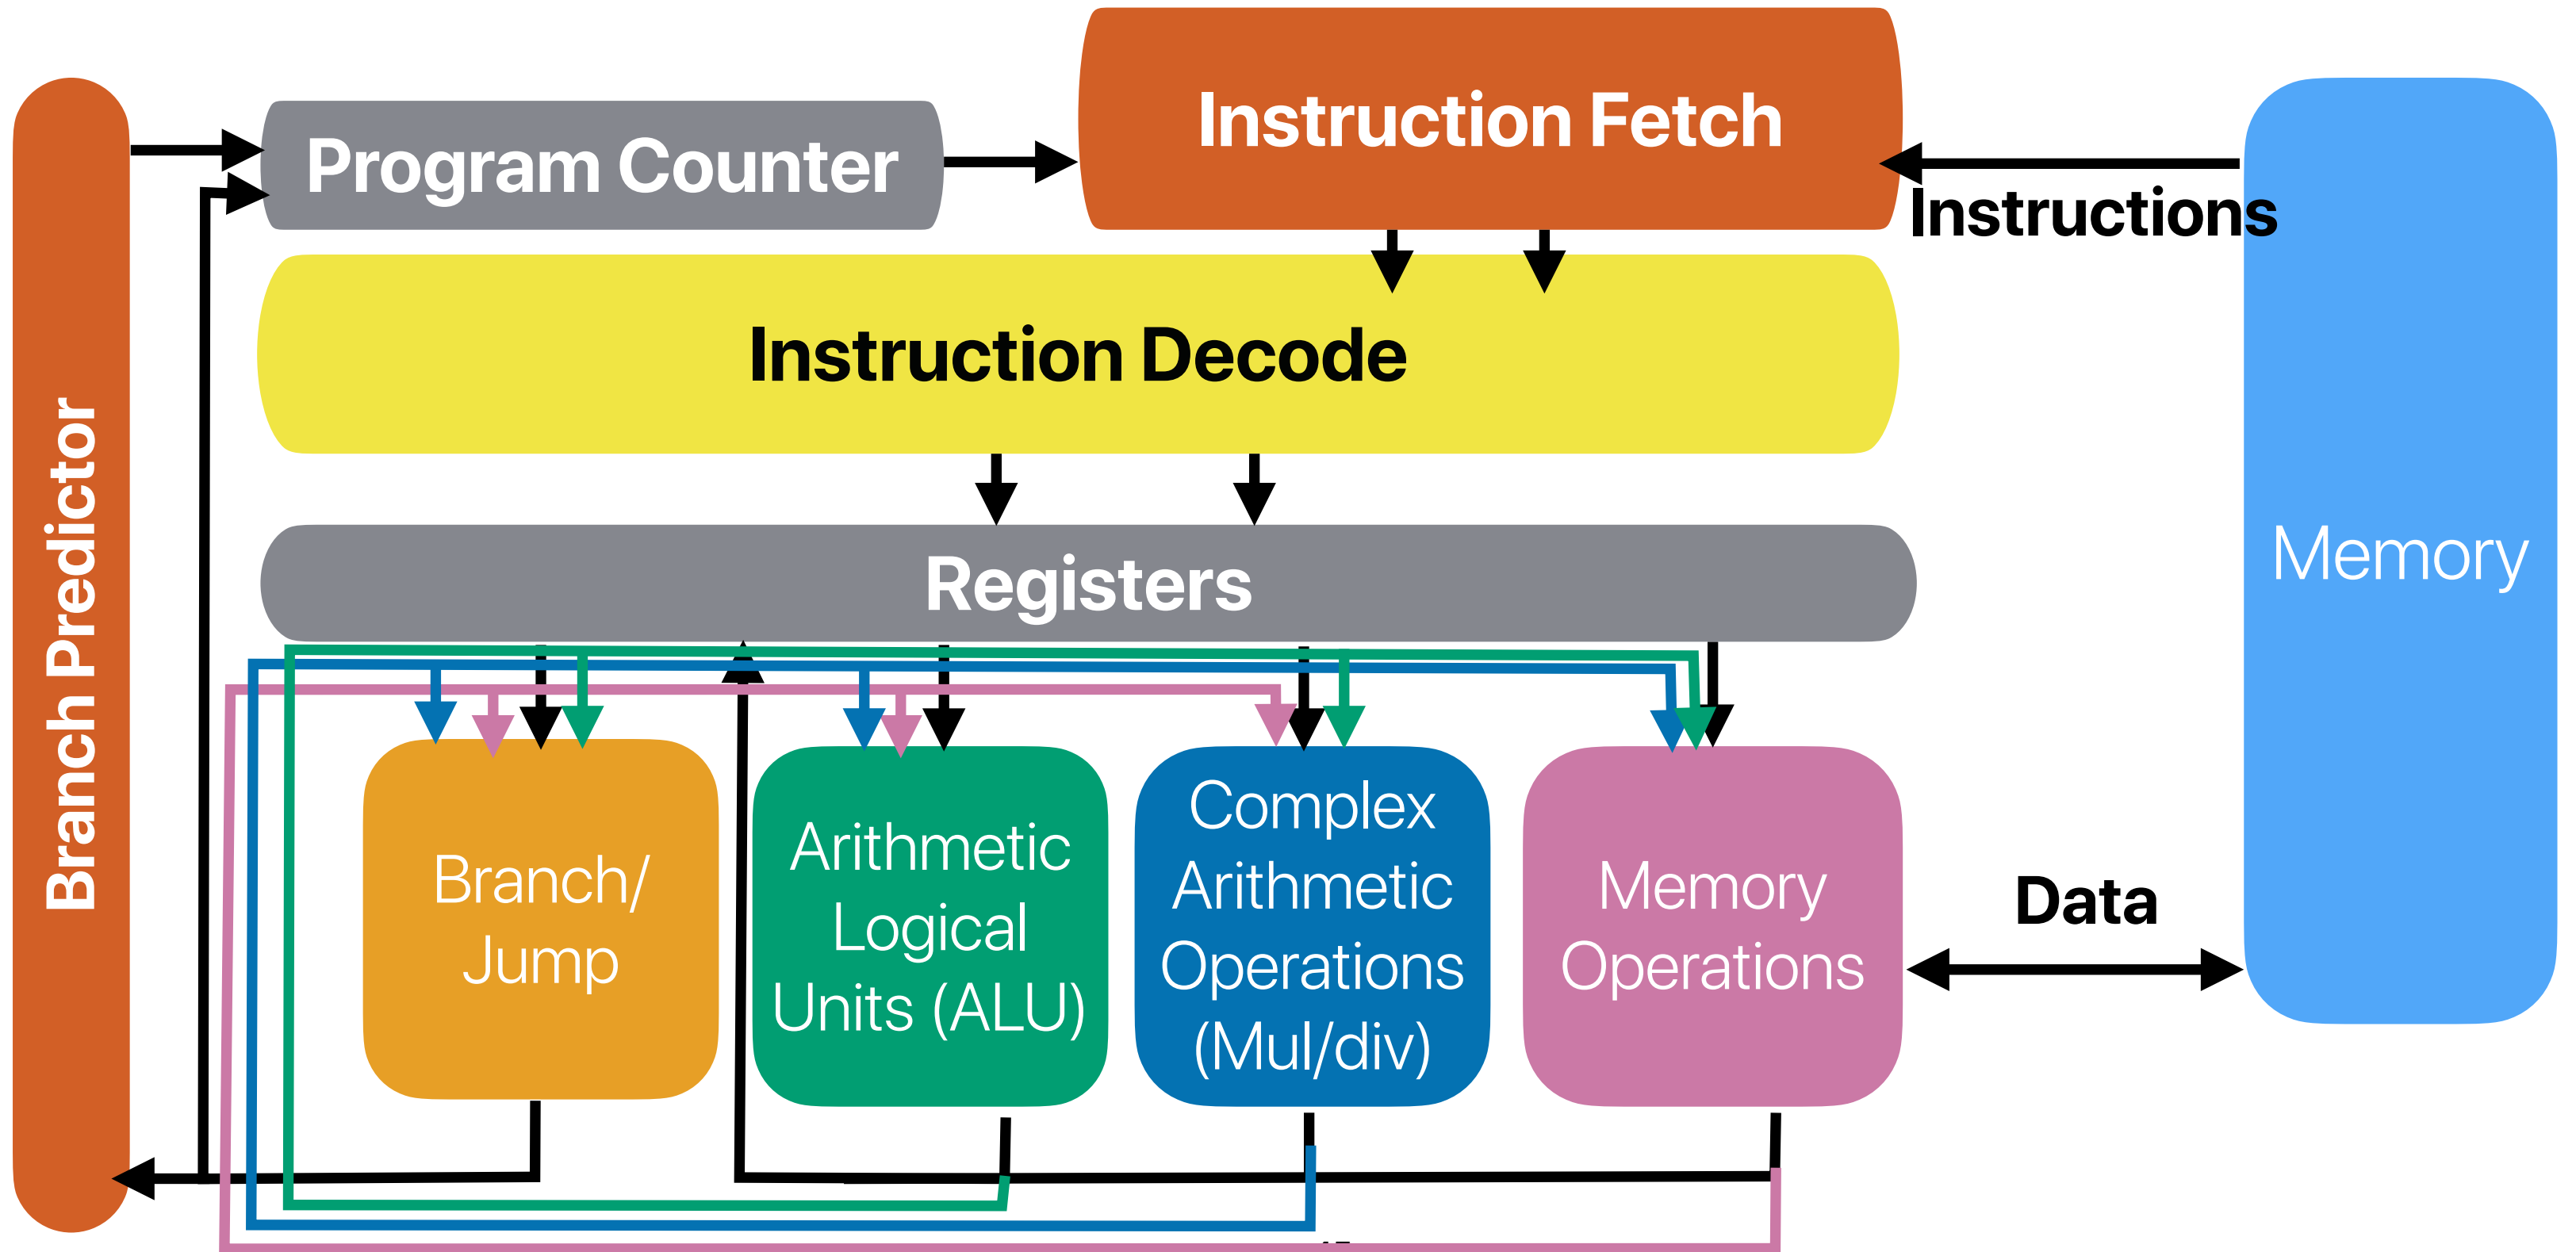


**If  $CPI == 1$  the limitation?**

# Data "forwarding"



# Super Scalar



# Super Scalar

# Superscalar


- Since we have many functional units now, we should fetch/decode more instructions each cycle so that we can have more instructions to issue!
- Super-scalar: fetch/decode/issue more than one instruction each cycle
  - **Fetch width:** how many instructions can the processor fetch/decode each cycle
  - **Issue width:** how many instructions can the processor issue each cycle

# Superscalar: fetch/issue width == 2

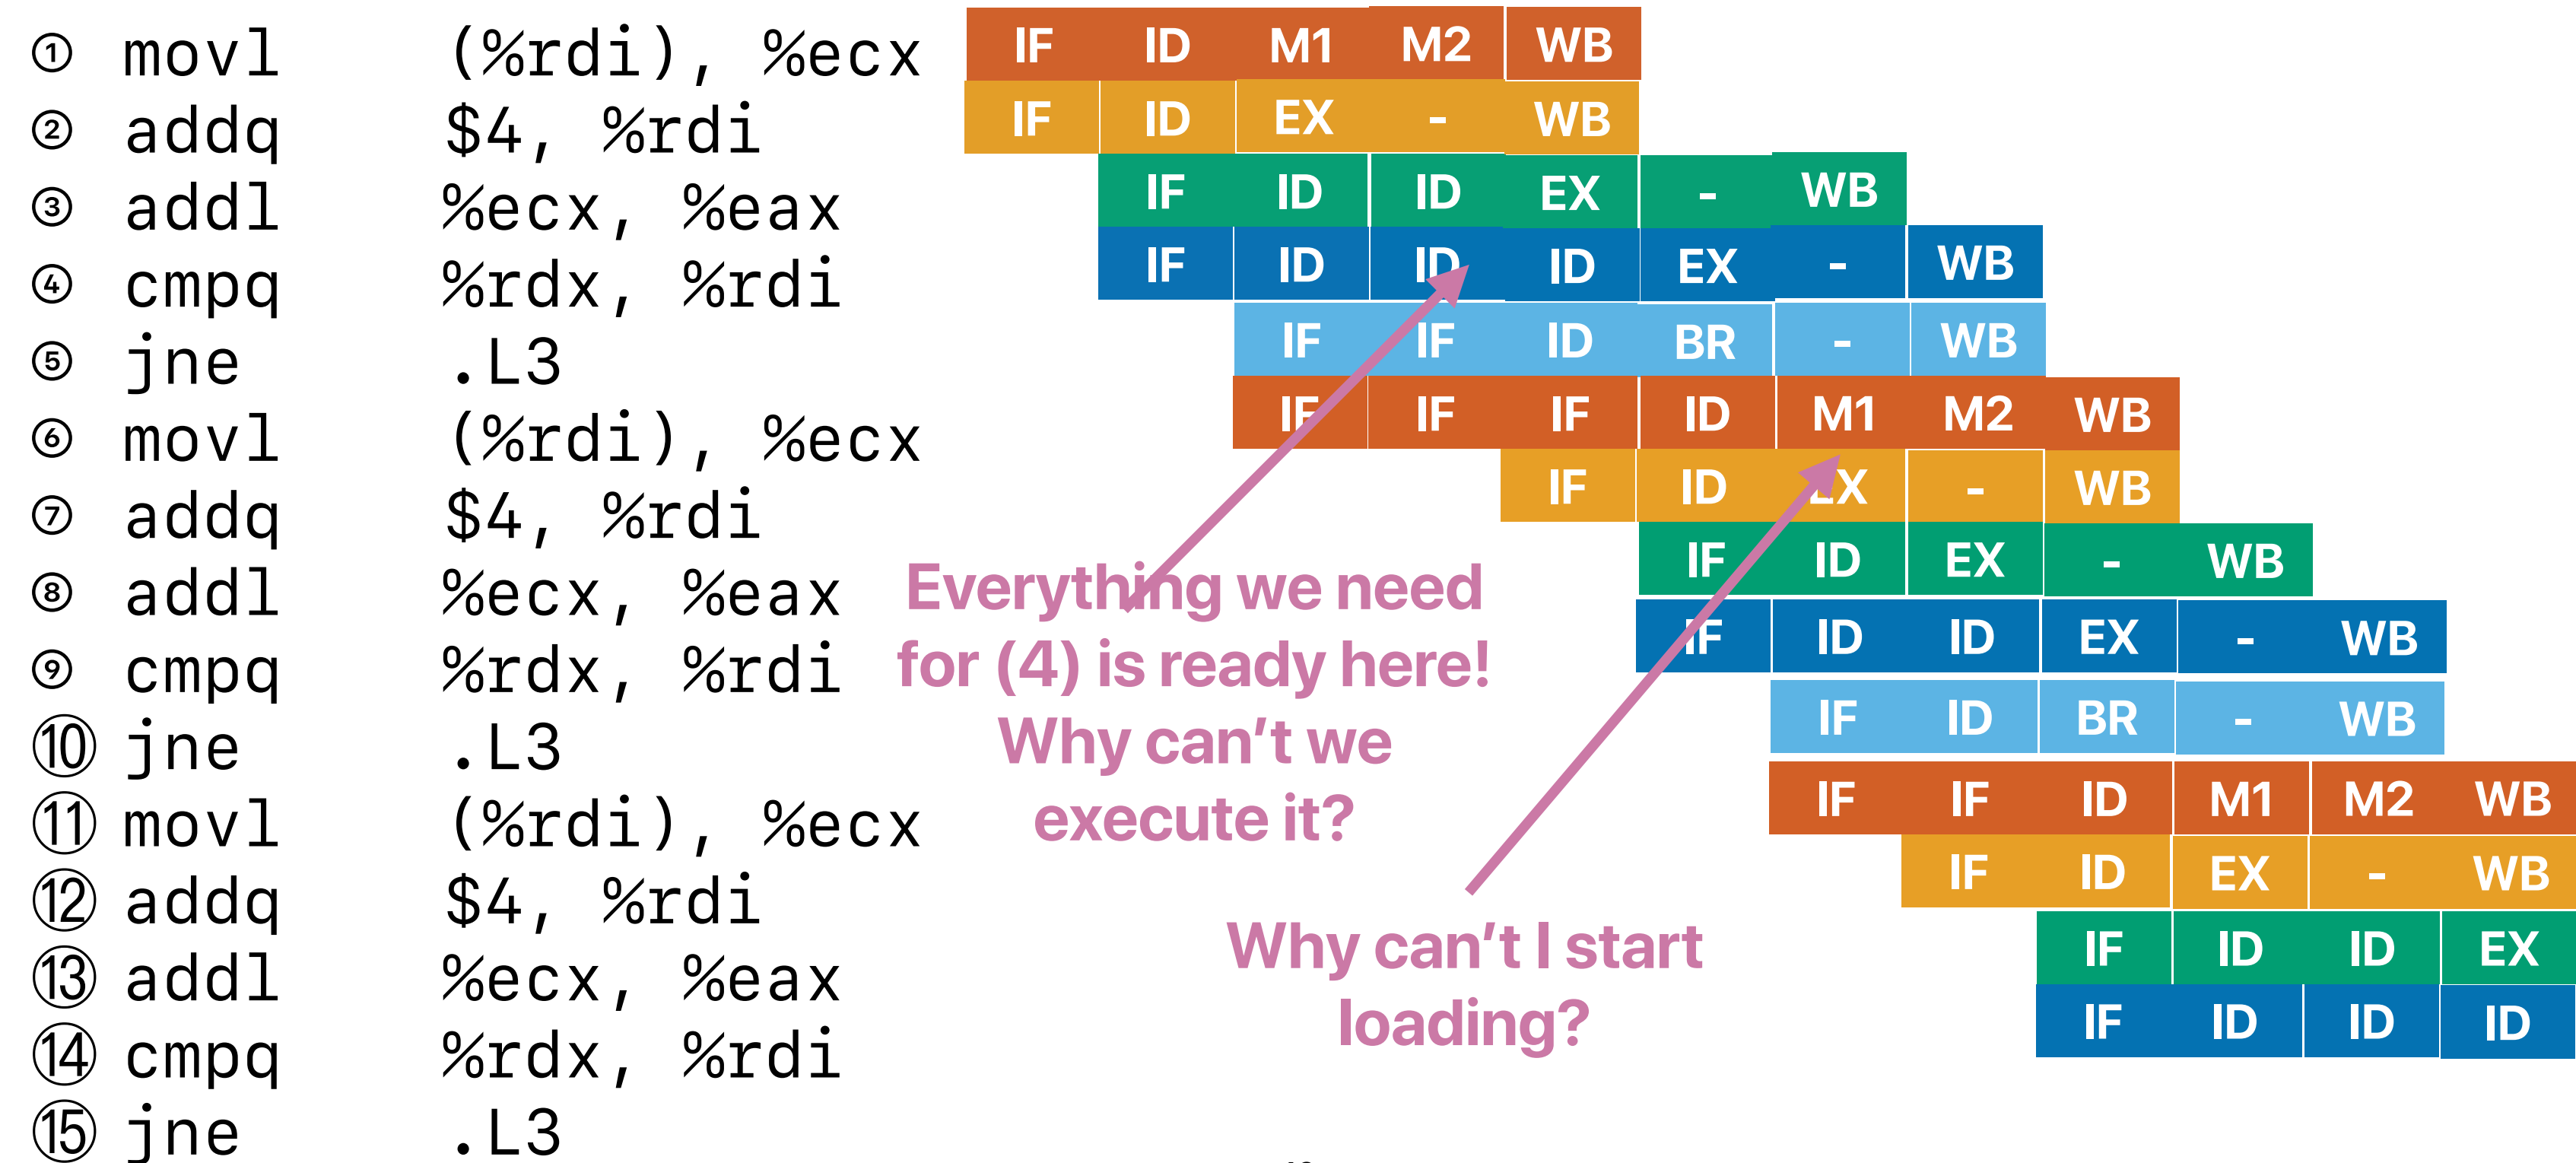
```
for(i = 0; i < count; i++) {  
    s += a[i];  
}
```

.L3:

```
movl    (%rdi), %ecx  
addq    $4, %rdi  
addl    %ecx, %eax  
cmpq    %rdx, %rdi  
jne     .L3  
ret
```



# If we loop many times (assume perfect predictor)



# Limitations of Compiler Optimizations

- If the hardware (e.g., pipeline changes), the same compiler optimization may not be that helpful
- The compiler can only optimize on static instructions, but cannot optimize dynamic instructions



# What do you need to execution an instruction?

- Whenever the instruction is decoded — put decoded instruction somewhere
- Whenever the inputs are ready — **all data dependencies are resolved**
- Whenever the target functional unit is available

# **Dynamic instruction scheduling/ Out-of-order (OoO) execution**

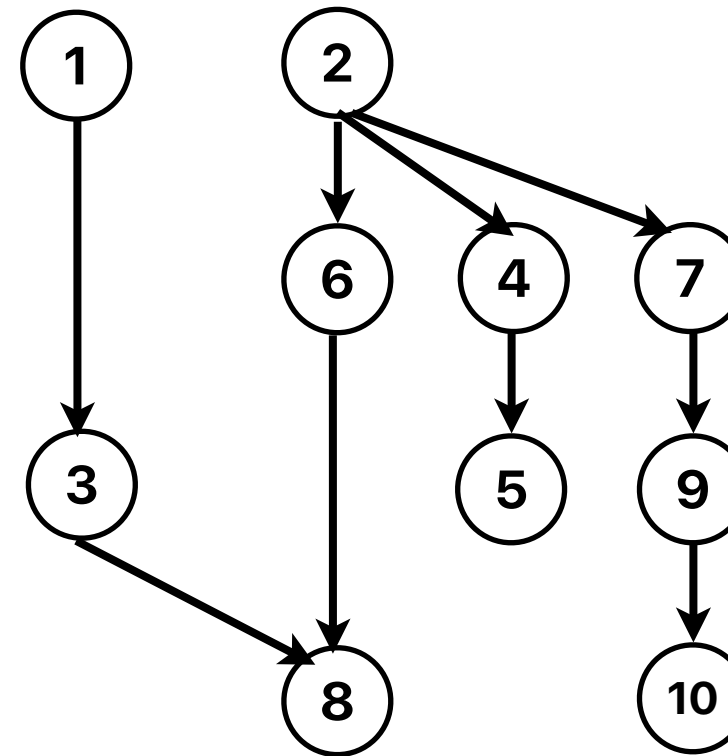
# What do you need to execution an instruction?

- Whenever the instruction is decoded — put decoded instruction somewhere
- Whenever the inputs are ready — **all data dependencies are resolved**
- Whenever the target functional unit is available

# Scheduling instructions: based on data dependencies

- Draw the data dependency graph, put an arrow if an instruction depends on the other.

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```

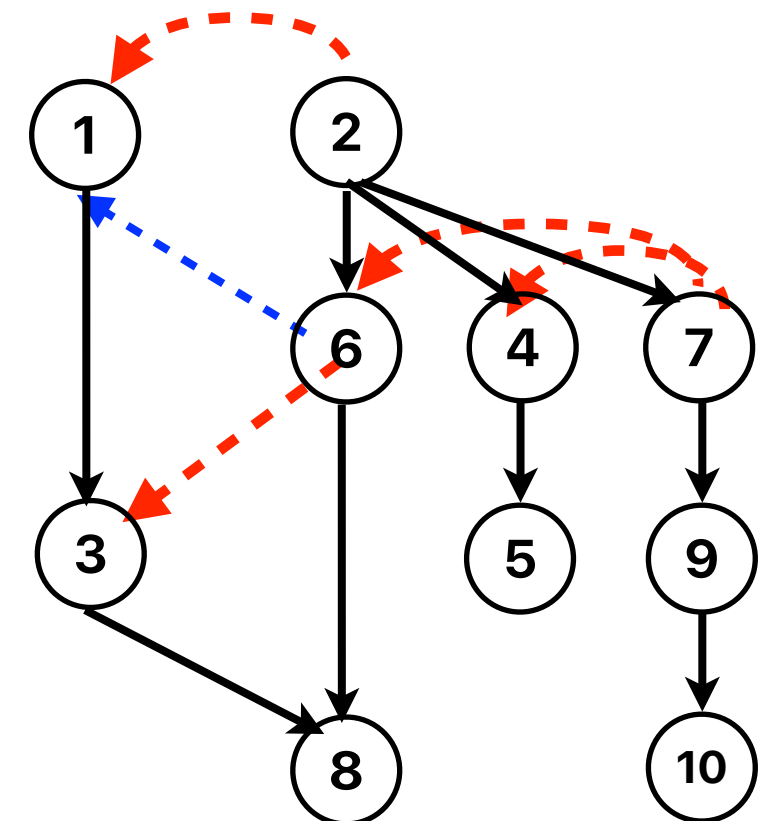


- In theory**, instructions without dependencies can be executed in parallel or out-of-order
- Instructions with dependencies can never be reordered

# False dependencies

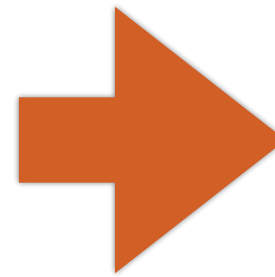
- We are still limited by **false dependencies**
- They are not “true” dependencies because they don’t have an arrow in data dependency graph
  - **WAR (Write After Read):** a later instruction overwrites the source of an earlier one
    - 2 and 1, 6 and 3, 7 and 4, 7 and 6
  - **WAW (Write After Write):** a later instruction overwrites the output of an earlier one
    - 6 and 1

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```



# What if we can use more registers...

```
① movl    (%rdi), %ecx
② addq    $4, %rdi
③ addl    %ecx, %eax
④ cmpq    %rdx, %rdi
⑤ jne     .L3
⑥ movl    (%rdi), %ecx
⑦ addq    $4, %rdi
⑧ addl    %ecx, %eax
⑨ cmpq    %rdx, %rdi
⑩ jne     .L3
```



```
① movl    (%rdi), %ecx
② addq    $4, %rdi, %t0
③ addl    %ecx, %eax, %t1
④ cmpq    %rdx, %t0
⑤ jne     .L3
⑥ movl    (%t0), %t2
⑦ addq    $4, %t0, %t3
⑧ addl    %t1, %t2, %t4
⑨ cmpq    %rdx, %t3
⑩ jne     .L3
```

**All false dependencies are gone!!!**

# Limitations of Compiler Optimizations

- If the hardware (e.g., pipeline changes), the same compiler optimization may not be that helpful
- The compiler can only optimize on static instructions, but cannot optimize dynamic instructions
- Compilers are limited by the registers an ISA provides

# Register renaming + speculative execution

- K. C. Yeager, "The Mips R10000 superscalar microprocessor," in IEEE Micro, vol. 16, no. 2, pp. 28-41, April 1996.
- R. E. Kessler, "The Alpha 21264 microprocessor," in IEEE Micro, vol. 19, no. 2, pp. 24-36, March-April 1999.



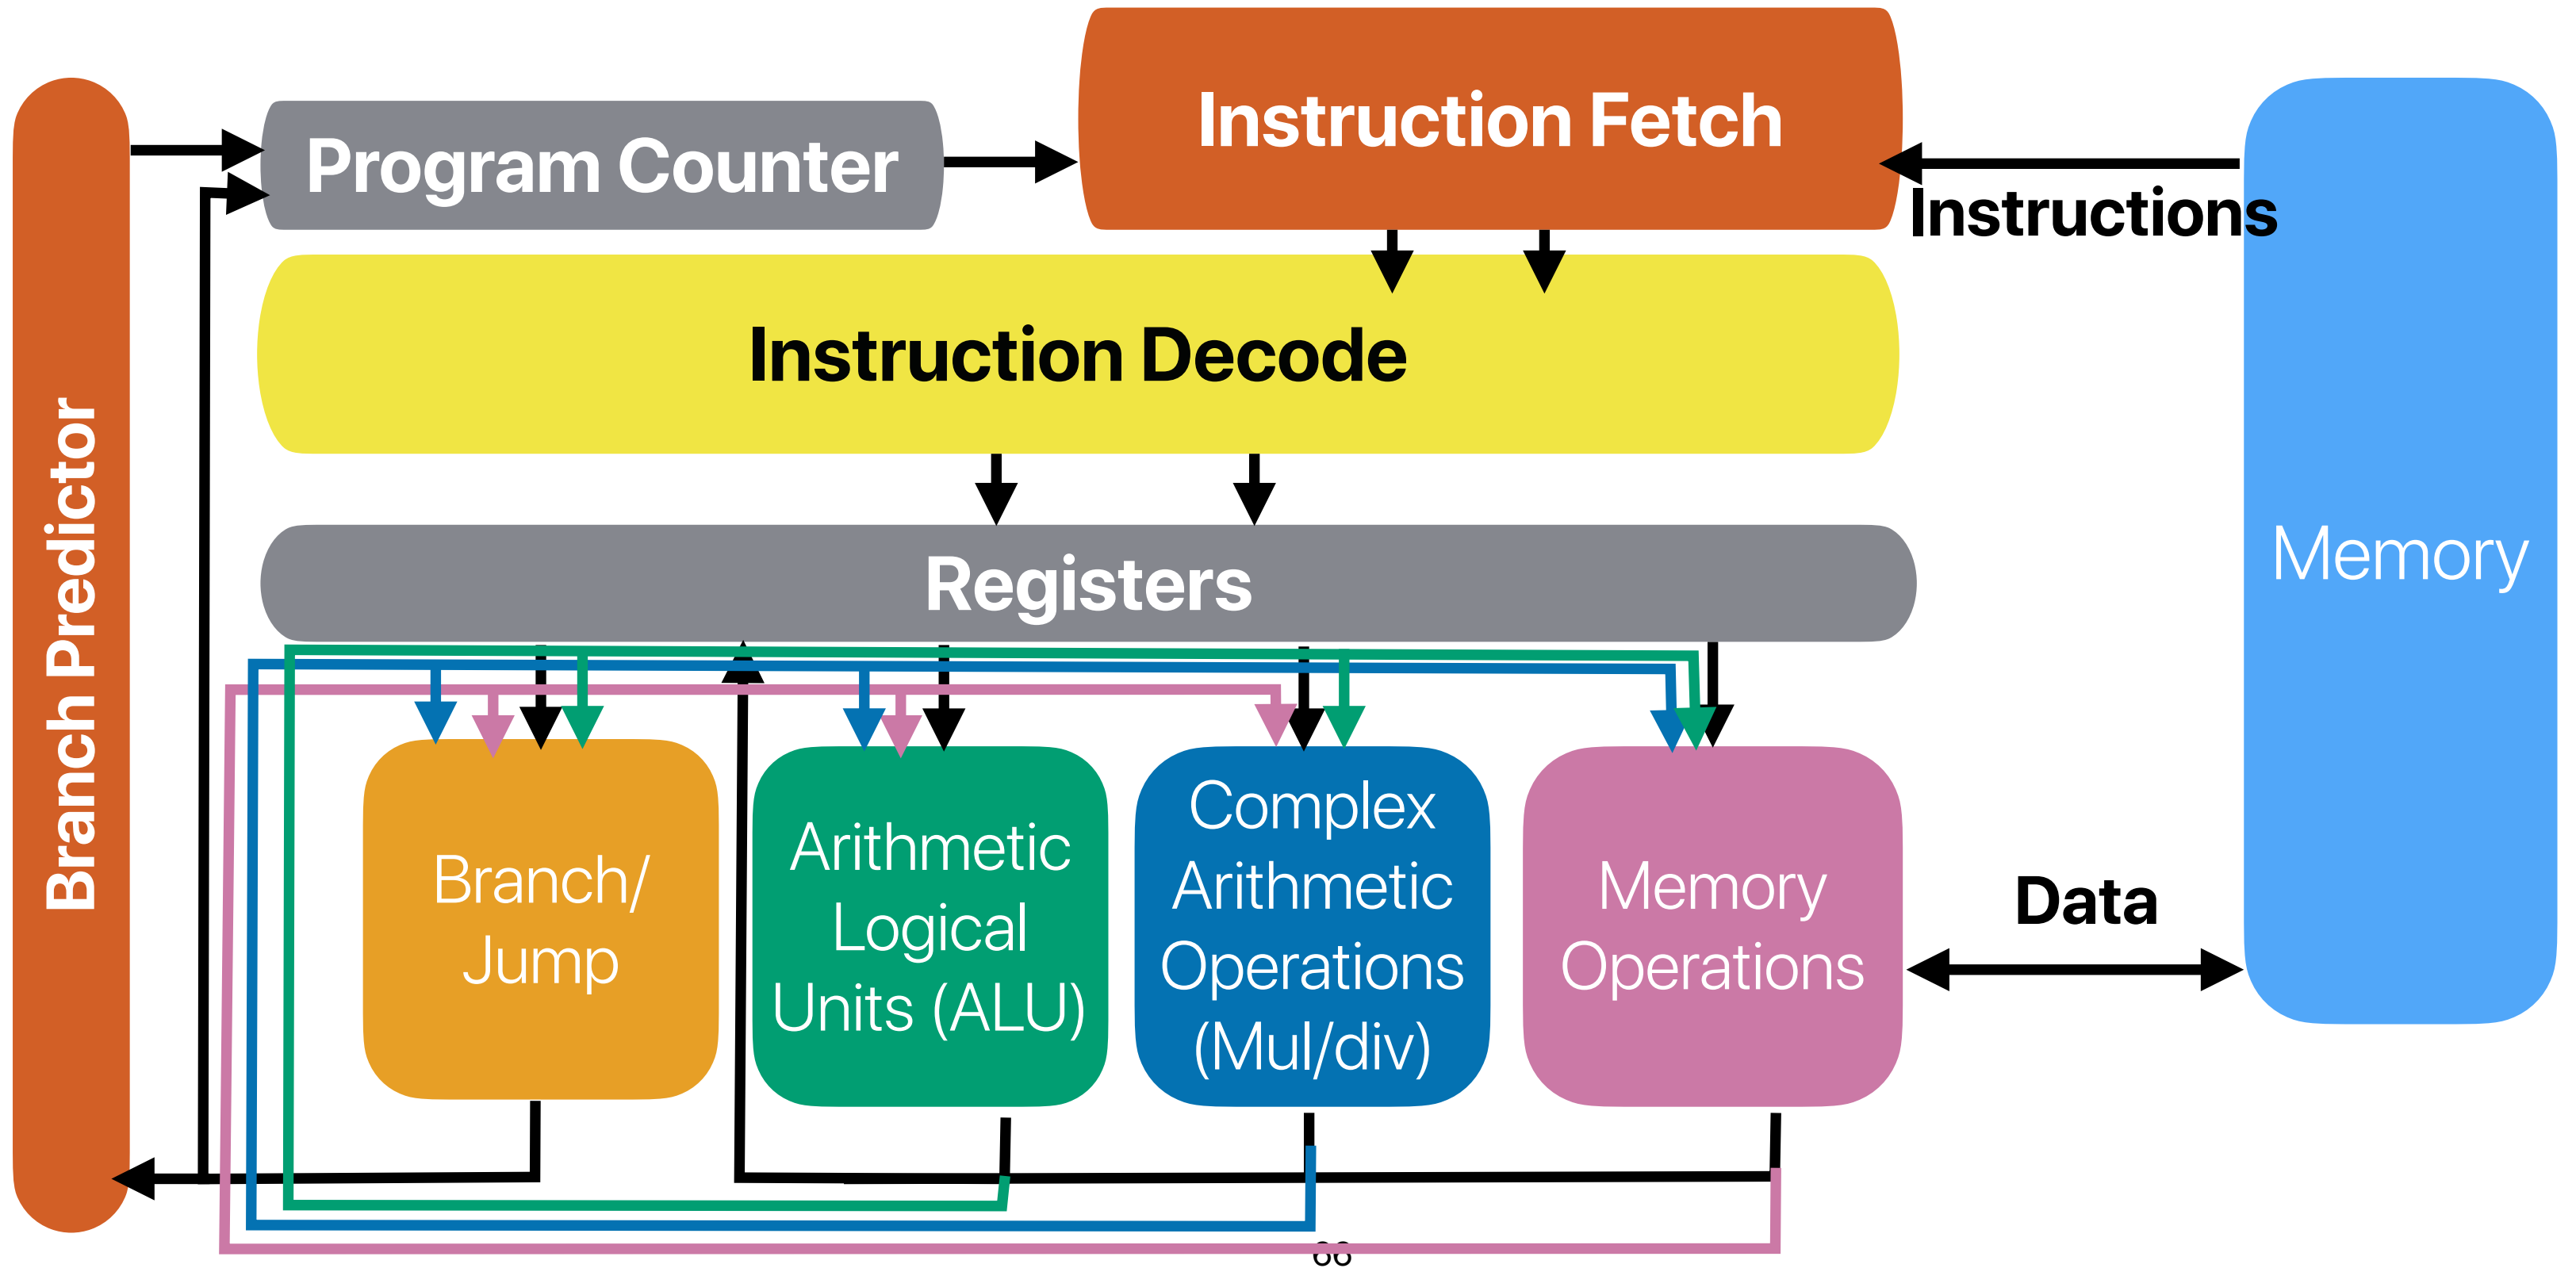
# Register renaming

- Provide a set of **physical registers** and a mapping table mapping **architectural registers** to physical registers
  - Architectural registers are virtual registers that software can see/use
- Allocate a physical register for a new output
- Stages
  - Dispatch/Rename (REN) — allocate a "physical register" for the output of a decoded instruction
  - Execute (EX, M1/M2, BR) — send the instruction to its corresponding pipeline if no structural hazards
  - Write Back (WB) — broadcast the result through CDB

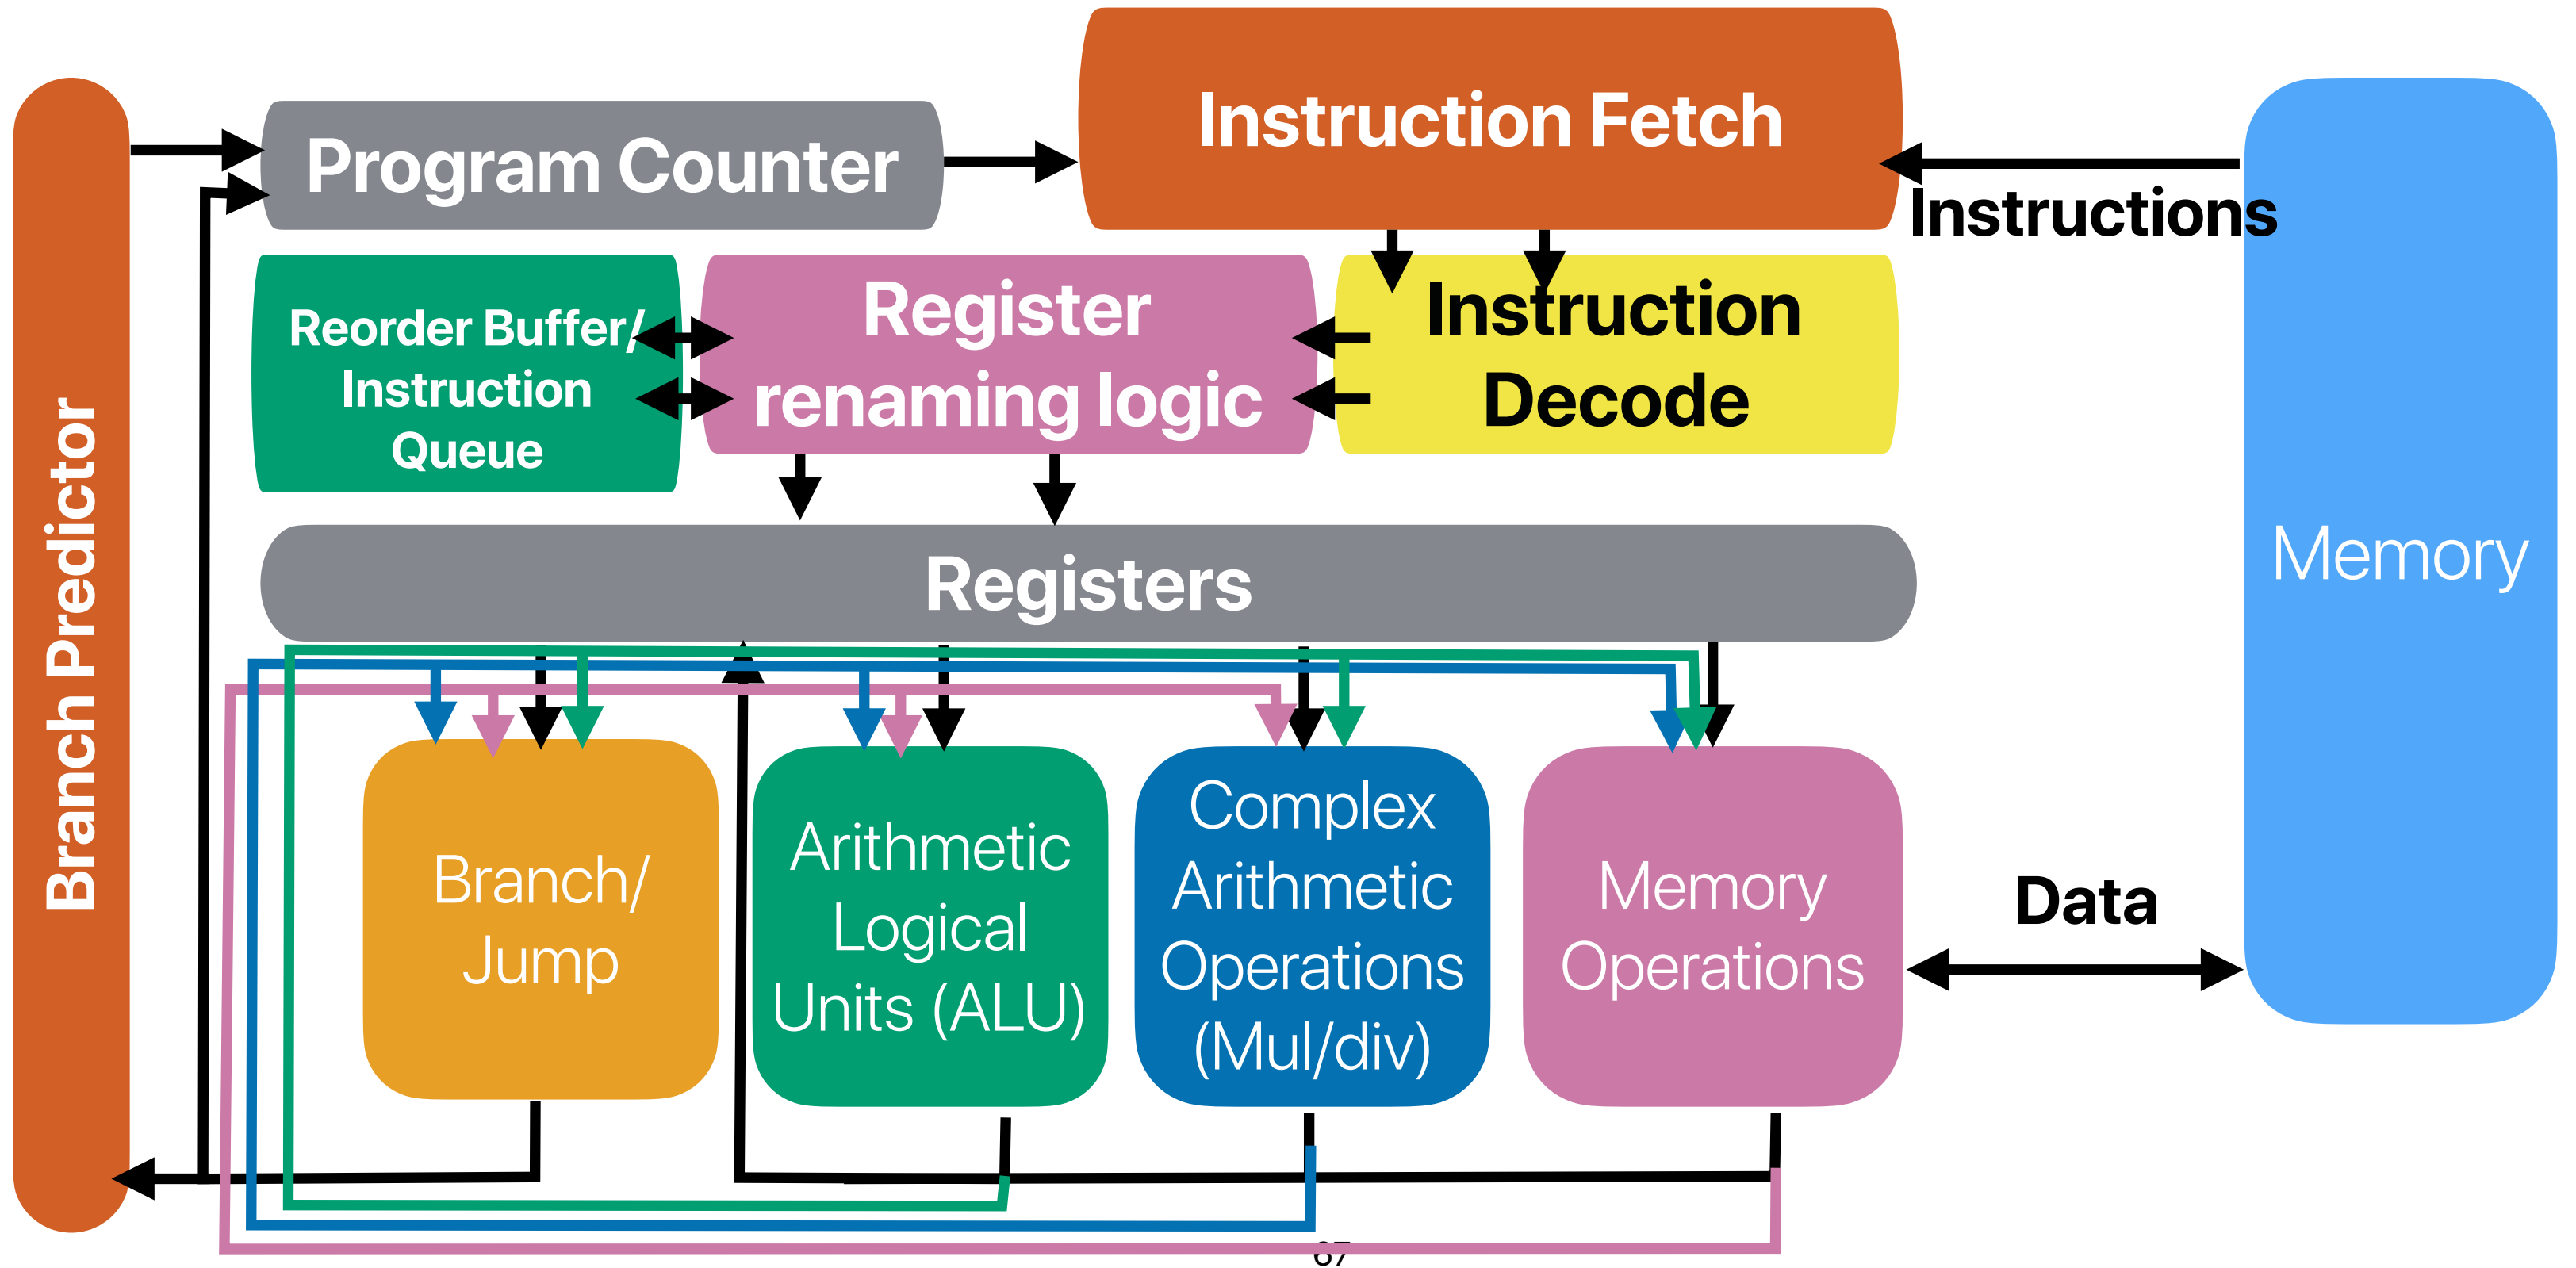
# Speculative Execution

- Exceptions (e.g. divided by 0, page fault) may occur anytime
  - A later instruction cannot write back its own result otherwise the architectural states won't be correct
- Hardware can schedule instruction across branch instructions with the help of branch prediction
  - Fetch instructions according to the branch prediction
  - However, branch predictor can never be perfect
- Execute instructions across branches
  - Speculative execution: execute an instruction before the processor know if we need to execute or not
  - Execute an instruction all operands are ready (the values of depending physical registers are generated)
  - Store results in **reorder buffer** before the processor knows if the instruction is going to be executed or not.

# Super Scalar



# Register renaming



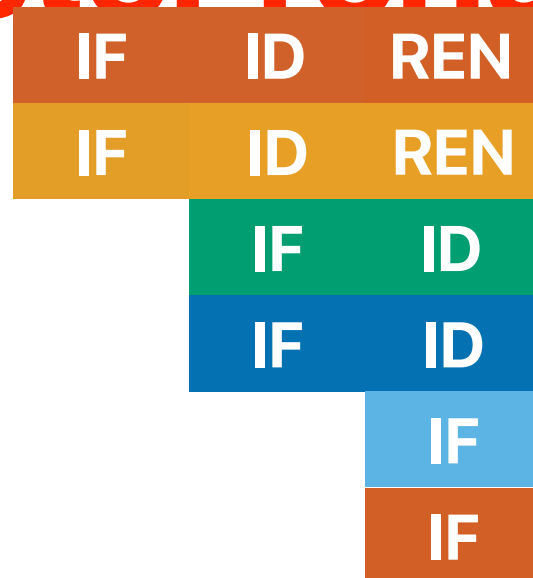
# Register renaming

```
①  movl    (%rdi), %ecx
②  addq    $4, %rdi
③  addl    %ecx, %eax
④  cmpq    %rdx, %rdi
⑤  jne     .L3
⑥  movl    (%rdi), %ecx
⑦  addq    $4, %rdi
⑧  addl    %ecx, %eax
⑨  cmpq    %rdx, %rdi
⑩  jne     .L3
```

	IF	ID	REN	M1	M2	EX	-	BR	-	WB
1	(1)(2)									
2	(3)(4)	(1)(2)								
3	(5)(6)	(3)(4)	(1)(2)							
4	(7)(8)	(5)(6)	(3)(4)	(1)		(2)				
5	(9)(10)	(7)(8)	(5)(6)		(1)	(4)				
6		(9)(10)	(7)(8)			(3)	(4)	(5)		(1)(2)
7			(9)(10)	(6)		(7)	(3)		(5)	
8					(6)	(9)				(3)(4)
9						(8)		(10)		(5)(6)
10										(7)(8)
11										(9)(10)

# Register renaming in motion

① `movl (%rdi), %ecx`  
② `addq $4, %rdi`  
③ `addl %ecx, %eax`  
④ `cmpq %rdx, %rdi`  
⑤ `jne .L3`  
⑥ `movl (%rdi), %ecx`  
⑦ `addq $4, %rdi`  
⑧ `addl %ecx, %eax`  
⑨ `cmpq %rdx, %rdi`  
⑩ `jne .L3`




Renamed instruction		
1	<code>movl</code>	<code>(%rdi), P1</code>
2	<code>addq</code>	<code>4, P2</code>
3		
4		
5		
6		
7		
8		
9		
10		

Physical Register	
eax	
ecx	P1
rdi	P2
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6			
P2	0		1	P7			
P3				P8			
P4				P9			
P5				P10			

# Register renaming in motion

① `movl (%rdi), %ecx`  
② `addq $4, %rdi`  
③ `addl %ecx, %eax`  
④ `cmpq %rdx, %rdi`  
⑤ `jne .L3`  
⑥ `movl (%rdi), %ecx`  
⑦ `addq $4, %rdi`  
⑧ `addl %ecx, %eax`  
⑨ `cmpq %rdx, %rdi`  
⑩ `jne .L3`



Renamed instruction		
1	<b>movl</b>	<b>(%rdi), P1</b>
2	<b>addq</b>	<b>\$4,%rdi, P2</b>
3	<b>addl</b>	<b>P1, %eax, P3</b>
4	<b>cmpq</b>	<b>%rdx, P2</b>
5		
6		
7		
8		
9		
10		

Physical Register	
eax	<b>P3</b>
ecx	<b>P1</b>
rdi	<b>P2</b>
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	<b>0</b>		<b>1</b>	P6			
P2	<b>0</b>		<b>1</b>	P7			
P3	<b>0</b>		<b>1</b>	P8			
P4				P9			
P5				P10			

# Register renaming in motion

① `movl (%rdi), %ecx`  
② `addq $4, %rdi`  
③ `addl %ecx, %eax`  
④ `cmpq %rdx, %rdi`  
⑤ `jne .L3`  
⑥ `movl (%rdi), %ecx`  
⑦ `addq $4, %rdi`  
⑧ `addl %ecx, %eax`  
⑨ `cmpq %rdx, %rdi`  
⑩ `jne .L3`

IF	ID	REN	M1	M2
IF	ID	REN	EX	-
	IF	ID	REN	REN
	IF	ID	REN	EX
		IF	ID	REN
		IF	ID	REN
			IF	ID
			IF	ID
				IF
				IF

(4) is now executing before (3)!

Renamed instruction		
1	<code>movl</code>	<code>(%rdi), P1</code>
2	<code>addq</code>	<code>\$4,%rdi, P2</code>
3	<code>addl</code>	<code>P1, %eax, P3</code>
4	<code>cmpq</code>	<code>%rdx, P2</code>
5	<code>jne</code>	<code>.L3</code>
6	<code>movl</code>	<code>(P2), P4</code>
7		
8		
9		
10		

Physical Register	
eax	P3
ecx	P4
rdi	P2
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	0		1	P6			
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5				P10			



# Register renaming in motion

① movl (%rdi), %ecx  
② addq \$4, %rdi  
③ addl %ecx, %eax  
④ cmpq %rdx, %rdi  
⑤ jne .L3  
⑥ movl (%rdi), %ecx  
⑦ addq \$4, %rdi  
⑧ addl %ecx, %eax  
⑨ cmpq %rdx, %rdi  
⑩ jne .L3

IF	ID	REN	M1	M2	WB
IF	ID	REN	EX	-	WB
	IF	ID	REN	REN	EX
	IF	ID	REN	EX	-
		IF	ID	REN	BR
		IF	ID	REN	REN
			IF	ID	REN
			IF	ID	REN
				IF	ID
				IF	ID

Assume issue width == 2, can only put 2 instructions into execution

Renamed instruction		
1	movl	(%rdi), P1
2	addq	\$4,%rdi, P2
3	addl	P1, %eax, P3
4	cmpq	%rdx, P2
5	jne	.L3
6	movl	(P2), P4
7	addq	\$4, P2, P5
8	addl	P4, P3, P6
9		
10		

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	0		1	P10			

# Register renaming in motion

① movl (%rdi), %ecx  
② addq \$4, %rdi  
③ addl %ecx, %eax  
④ cmpq %rdx, %rdi  
⑤ jne .L3  
⑥ movl (%rdi), %ecx  
⑦ addq \$4, %rdi  
⑧ addl %ecx, %eax  
⑨ cmpq %rdx, %rdi  
⑩ jne .L3

IF	ID	REN	M1	M2	WB
IF	ID	REN	EX	-	WB
IF	ID	REN	REN	EX	-
IF	ID	REN	EX	-	-
	IF	ID	REN	BR	-
	IF	ID	REN	REN	M1
		IF	ID	REN	EX
		IF	ID	REN	REN
			IF	ID	REN
			IF	ID	REN

Renamed instruction		
1	<del>movl (%rdi), P1</del>	
2	<del>addq \$4,%rdi, P2</del>	
3	addl P1, %eax, P3	
4	cmpq %rdx, P2	
5	jne .L3	
6	movl (P2), P4	
7	addq \$4, P2, P5	
8	addl P4, P3, P6	
9	cmpq %rdx, P5	
10	jne .L3	

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	0		1	P10			

# Register renaming in motion

① movl (%rdi), %ecx  
② addq \$4, %rdi  
③ addl %ecx, %eax  
④ cmpq %rdx, %rdi  
⑤ jne .L3  
⑥ movl (%rdi), %ecx  
⑦ addq \$4, %rdi  
⑧ addl %ecx, %eax  
⑨ cmpq %rdx, %rdi  
⑩ jne .L3

IF	ID	REN	M1	M2	WB		
IF	ID	REN	EX	-	WB		
	IF	ID	REN	REN	EX	-	WB
	IF	ID	REN	EX	-	WB	WB
		IF	ID	REN	BR	-	WB
		IF	ID	REN	REN	M1	M2
			IF	ID	REN	EX	-
			IF	ID	REN	REN	REN
				IF	ID	REN	EX
				IF	ID	REN	REN

Renamed instruction	
1	<del>movl (%rdi), P1</del>
2	<del>addq \$4,%rdi, P2</del>
3	<del>addl P1, %eax, P3</del>
4	<del>cmpq %rdx, P2</del>
5	jne .L3
6	movl (P2), P4
7	addq \$4, P2, P5
8	addl P4, P3, P6
9	cmpq %rdx, P5
10	jne .L3

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	0		1	P9			
P5	1		1	P10			

# Register renaming in motion

① movl (%rdi), %ecx  
② addq \$4, %rdi  
③ addl %ecx, %eax  
④ cmpq %rdx, %rdi  
⑤ jne .L3  
⑥ movl (%rdi), %ecx  
⑦ addq \$4, %rdi  
⑧ addl %ecx, %eax  
⑨ cmpq %rdx, %rdi  
⑩ jne .L3

IF	ID	REN	M1	M2	WB			
IF	ID	REN	EX	-	WB			
	IF	ID	REN	REN	EX	-	WB	
	IF	ID	REN	EX	-	WB	WB	
		IF	ID	REN	BR	-	WB	WB
		IF	ID	REN	REN	M1	M2	WB
			IF	ID	REN	EX	-	WB
			IF	ID	REN	REN	REN	EX
				IF	ID	REN	EX	-
				IF	ID	REN	REN	BR

Renamed instruction	
1	<del>movl (%rdi), P1</del>
2	<del>addq \$4,%rdi, P2</del>
3	<del>addl P1, %eax, P3</del>
4	<del>cmpq %rdx, P2</del>
5	jne .L3
6	movl (P2), P4
7	addq \$4, P2, P5
8	addl P4, P3, P6
9	cmpq %rdx, P5
10	jne .L3

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	1		1
P2	1		1	P7			
P3	0		1	P8			
P4	1		1	P9			
P5	1		1	P10			

# Register renaming in motion

① movl (%rdi), %ecx  
② addq \$4, %rdi  
③ addl %ecx, %eax  
④ cmpq %rdx, %rdi  
⑤ jne .L3  
⑥ movl (%rdi), %ecx  
⑦ addq \$4, %rdi  
⑧ addl %ecx, %eax  
⑨ cmpq %rdx, %rdi  
⑩ jne .L3

IF	ID	REN	M1	M2	WB			
IF	ID	REN	EX	-	WB			
	IF	ID	REN	REN	EX	-	WB	
	IF	ID	REN	EX	-	WB	WB	
		IF	ID	REN	BR	-	WB	WB
		IF	ID	REN	REN	M1	M2	WB
			IF	ID	REN	EX	-	WB
			IF	ID	REN	REN	REN	EX
				IF	ID	REN	EX	WB
					IF	ID	REN	EX
						IF	ID	REN
							IF	ID


Renamed instruction	
1	<del>movl (%rdi), P1</del>
2	<del>addq \$4,%rdi, P2</del>
3	<del>addl P1, %eax, P3</del>
4	<del>cmpq %rdx, P2</del>
5	<del>jne .L3</del>
6	<del>movl (P2), P4</del>
7	<del>addq \$4, P2, P5</del>
8	<del>addl P4, P3, P6</del>
9	cmpq %rdx, P5
10	jne .L3

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	1		1	P9			
P5	1		1	P10			

# Register renaming in motion

① movl (%rdi), %ecx  
② addq \$4, %rdi  
③ addl %ecx, %eax  
④ cmpq %rdx, %rdi  
⑤ jne .L3  
⑥ movl (%rdi), %ecx  
⑦ addq \$4, %rdi  
⑧ addl %ecx, %eax  
⑨ cmpq %rdx, %rdi  
⑩ jne .L3



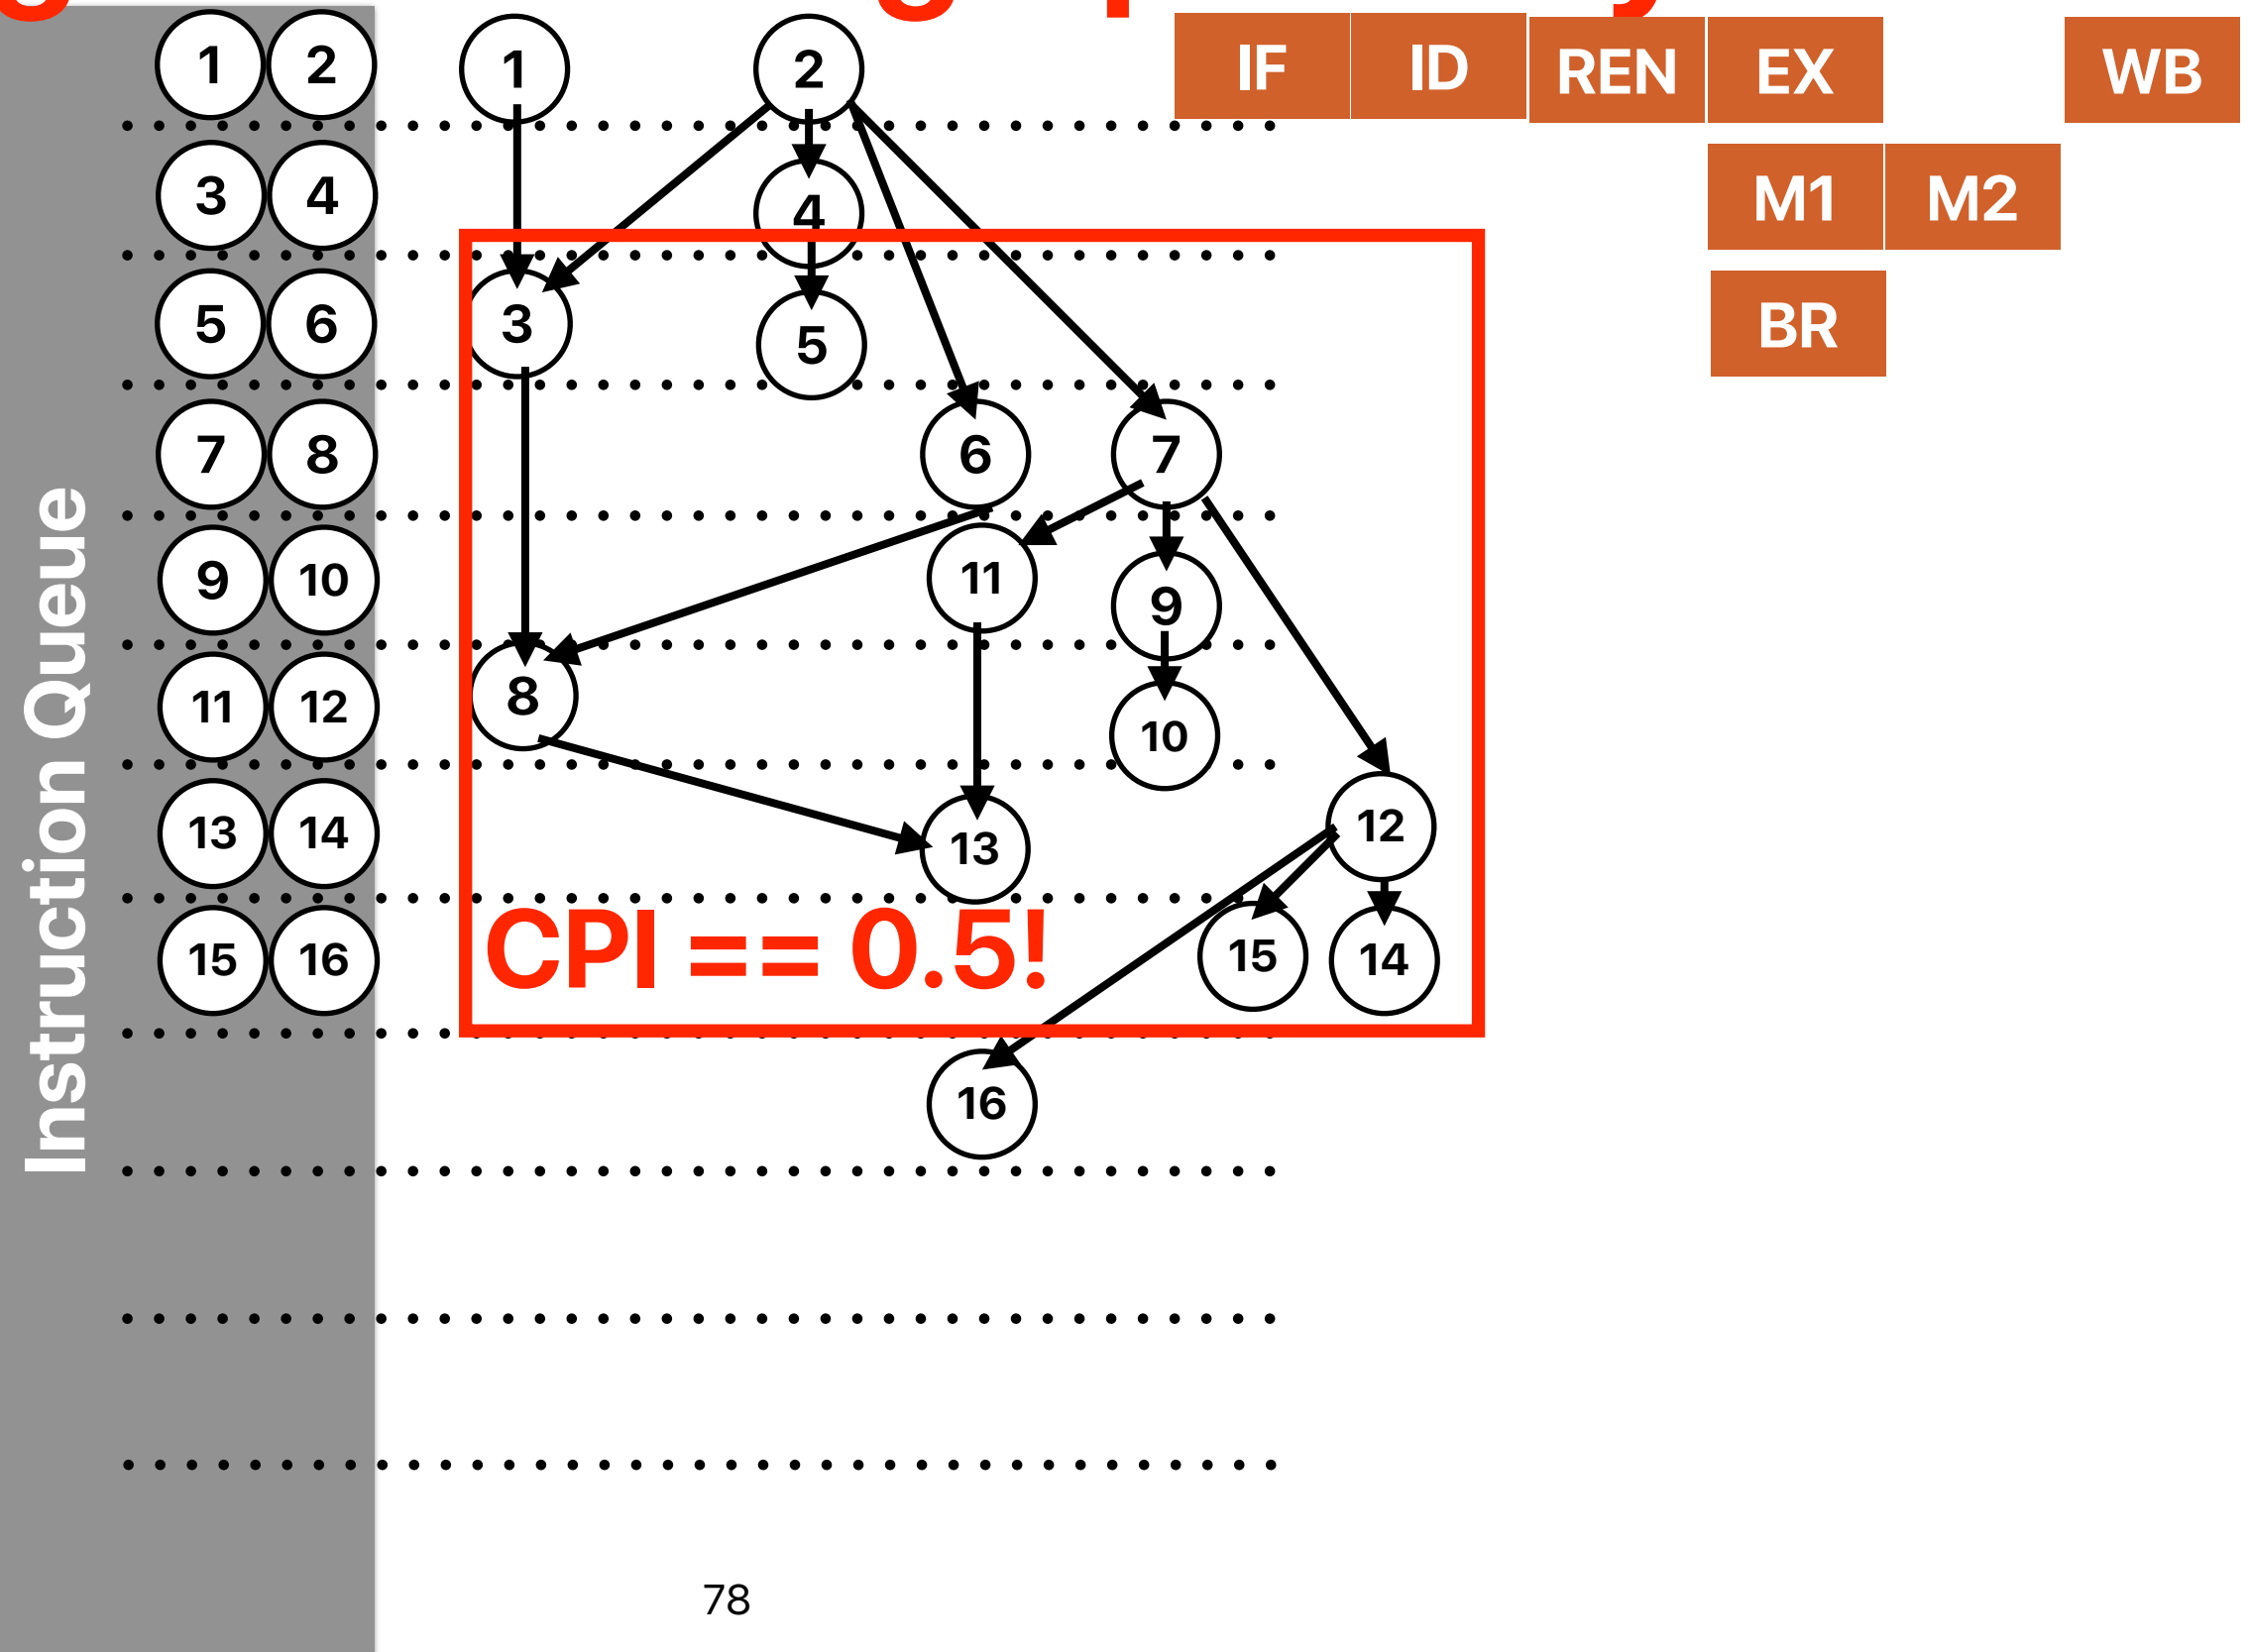
Renamed instruction	
1	<del>movl (%rdi), P1</del>
2	<del>addq \$4,%rdi, P2</del>
3	<del>addl P1, %eax, P3</del>
4	<del>cmpq %rdx, P2</del>
5	<del>jne .L3</del>
6	<del>movl (P2), P4</del>
7	<del>addq \$4, P2, P5</del>
8	<del>addl P4, P3, P6</del>
9	<del>cmpq %rdx, P5</del>
10	<del>jne .L3</del>

Physical Register	
eax	P3
ecx	P4
rdi	P5
rdx	

	Valid	Value	In use		Valid	Value	In use
P1	1		1	P6	0		1
P2	1		1	P7			
P3	0		1	P8			
P4	1		1	P9			
P5	1		1	P10			

# Through data flow graph analysis

```
① movl (%rdi), %ecx
② addq $4, %rdi
③ addl %ecx, %eax
④ cmpq %rdx, %rdi
⑤ jne .L3
⑥ movl (%rdi), %ecx
⑦ addq $4, %rdi
⑧ addl %ecx, %eax
⑨ cmpq %rdx, %rdi
⑩ jne .L3
⑪ movl (%rdi), %ecx
⑫ addq $4, %rdi
⑬ addl %ecx, %eax
⑭ cmpq %rdx, %rdi
⑮ jne .L3
⑯ movl (%rdi), %ecx
```



# Reorder Buffer (ROB)



# Speculative Execution

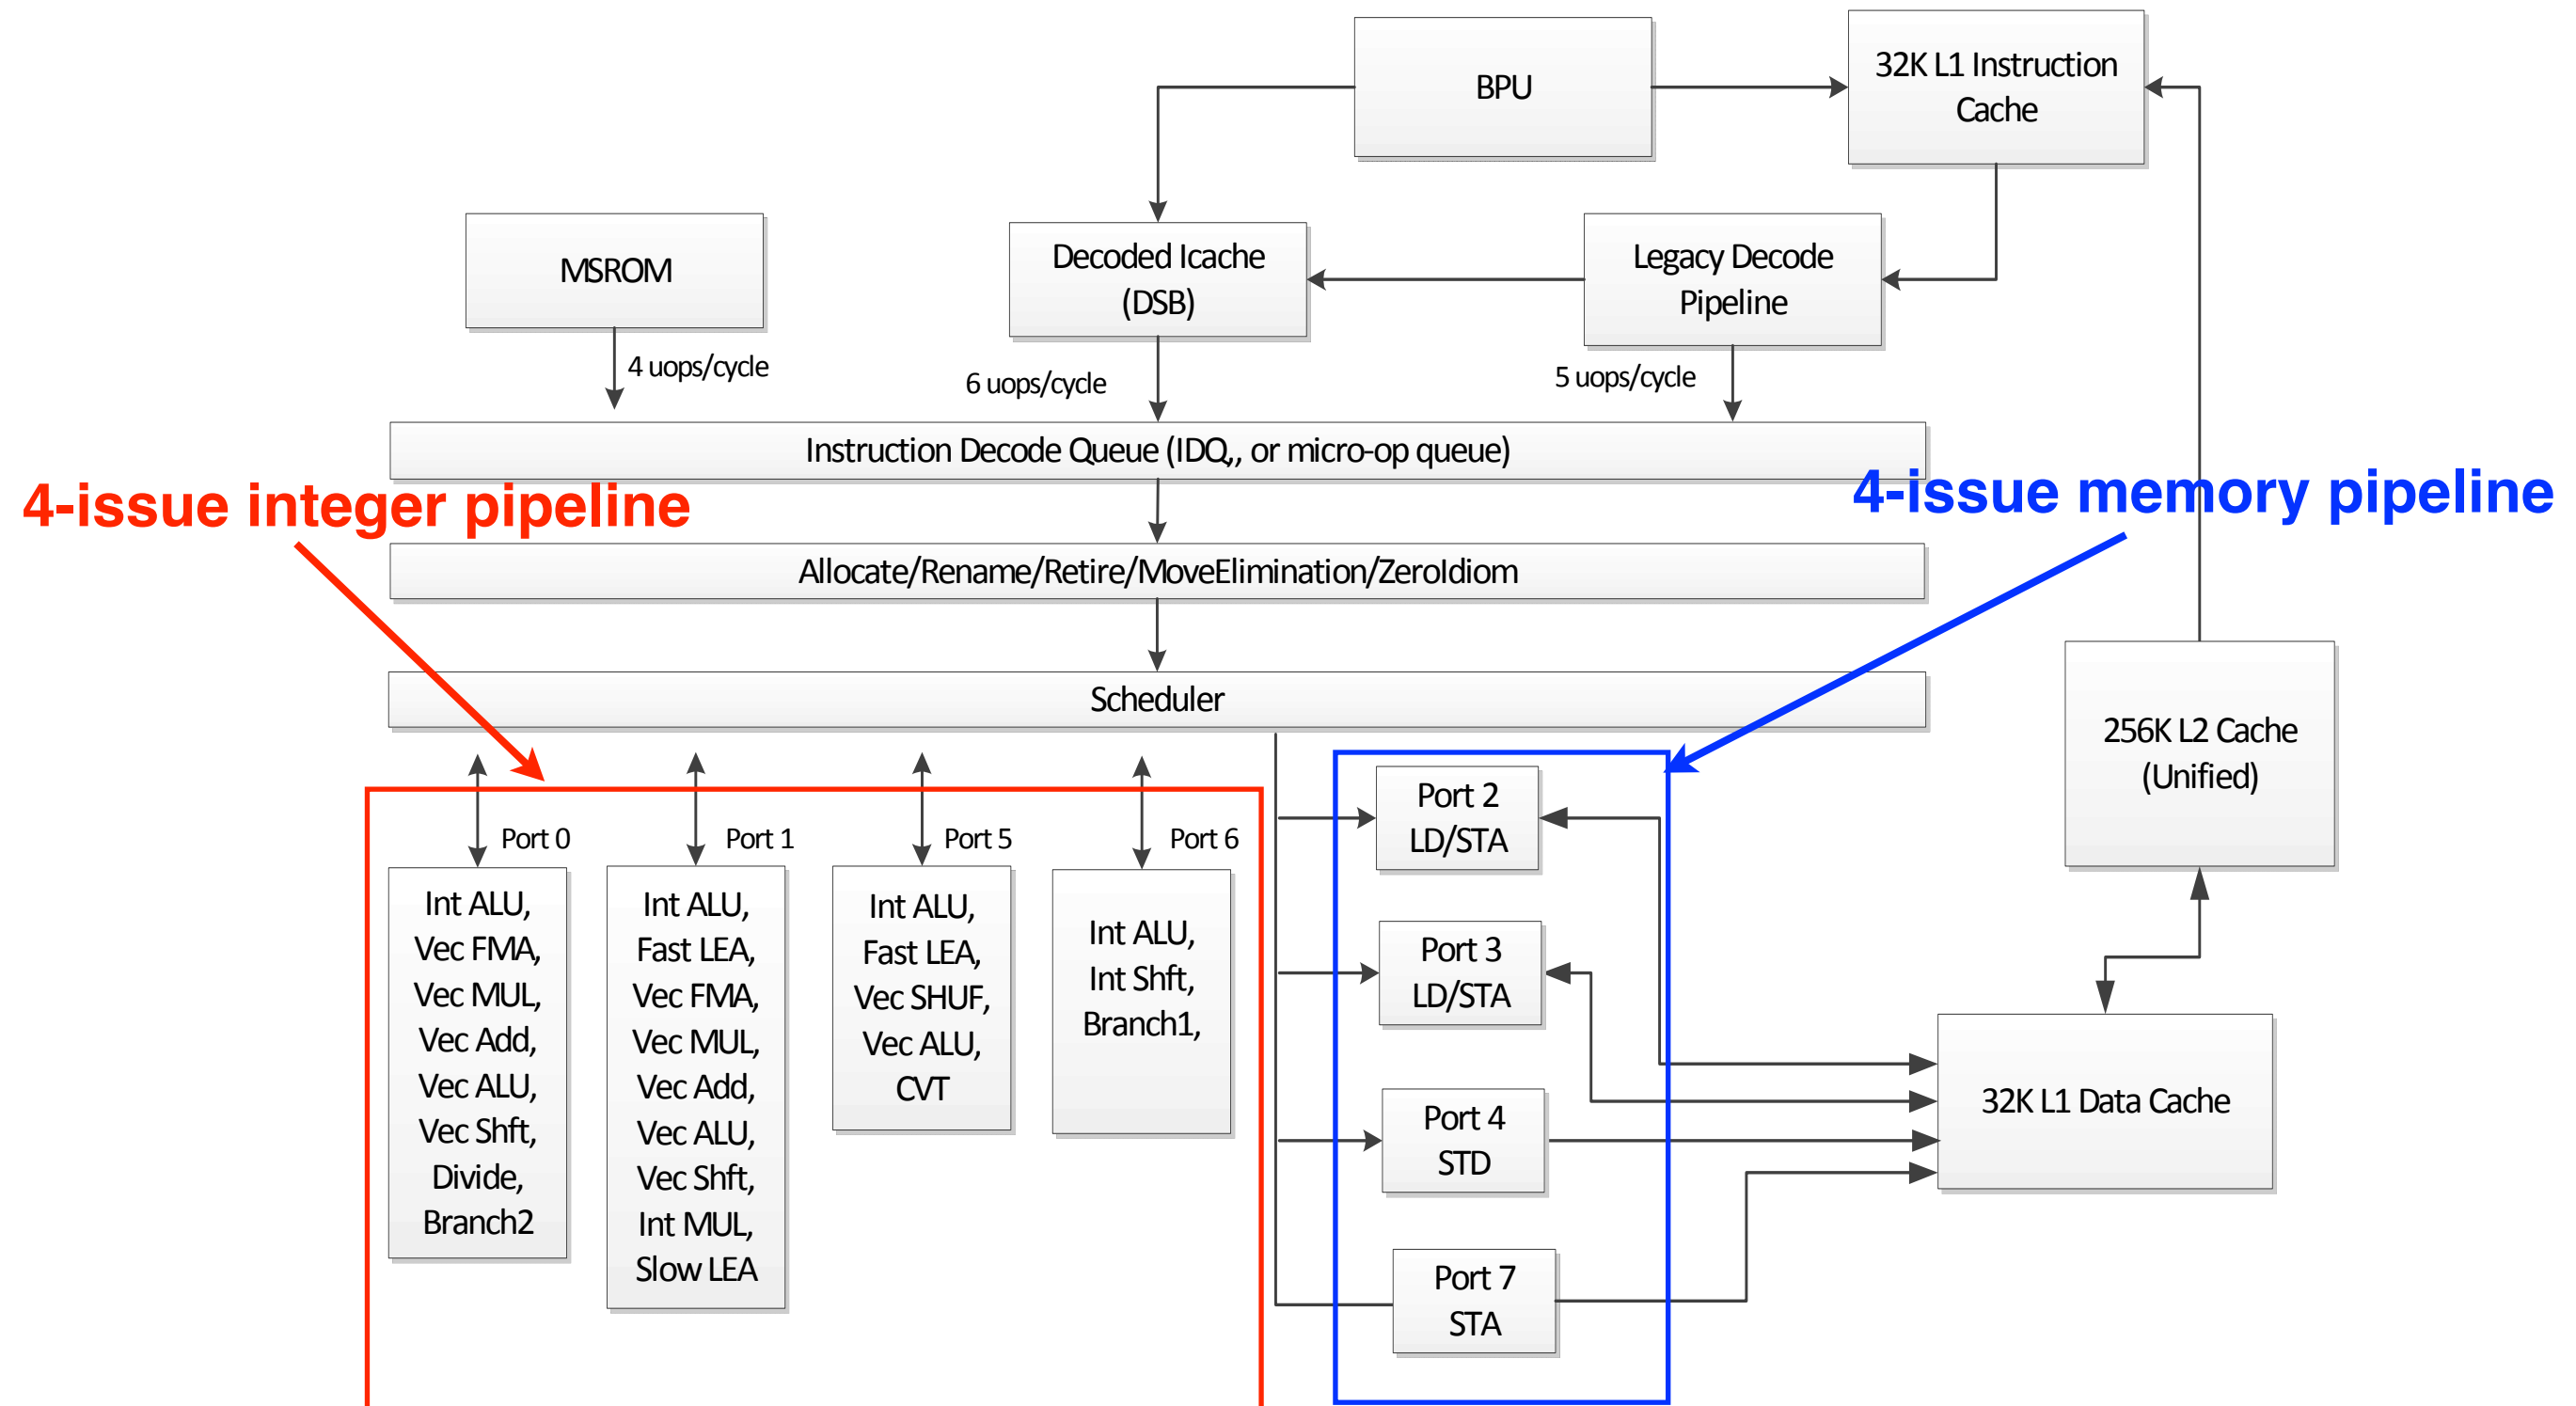
- Any execution of an instruction before a prior instruction finishes is considered as **speculative execution**
- Because it's speculative, we need to preserve the capability to restore to the states before it's executed
  - Branch mis-prediction
  - Exceptions

# Reorder buffer/Commit stage

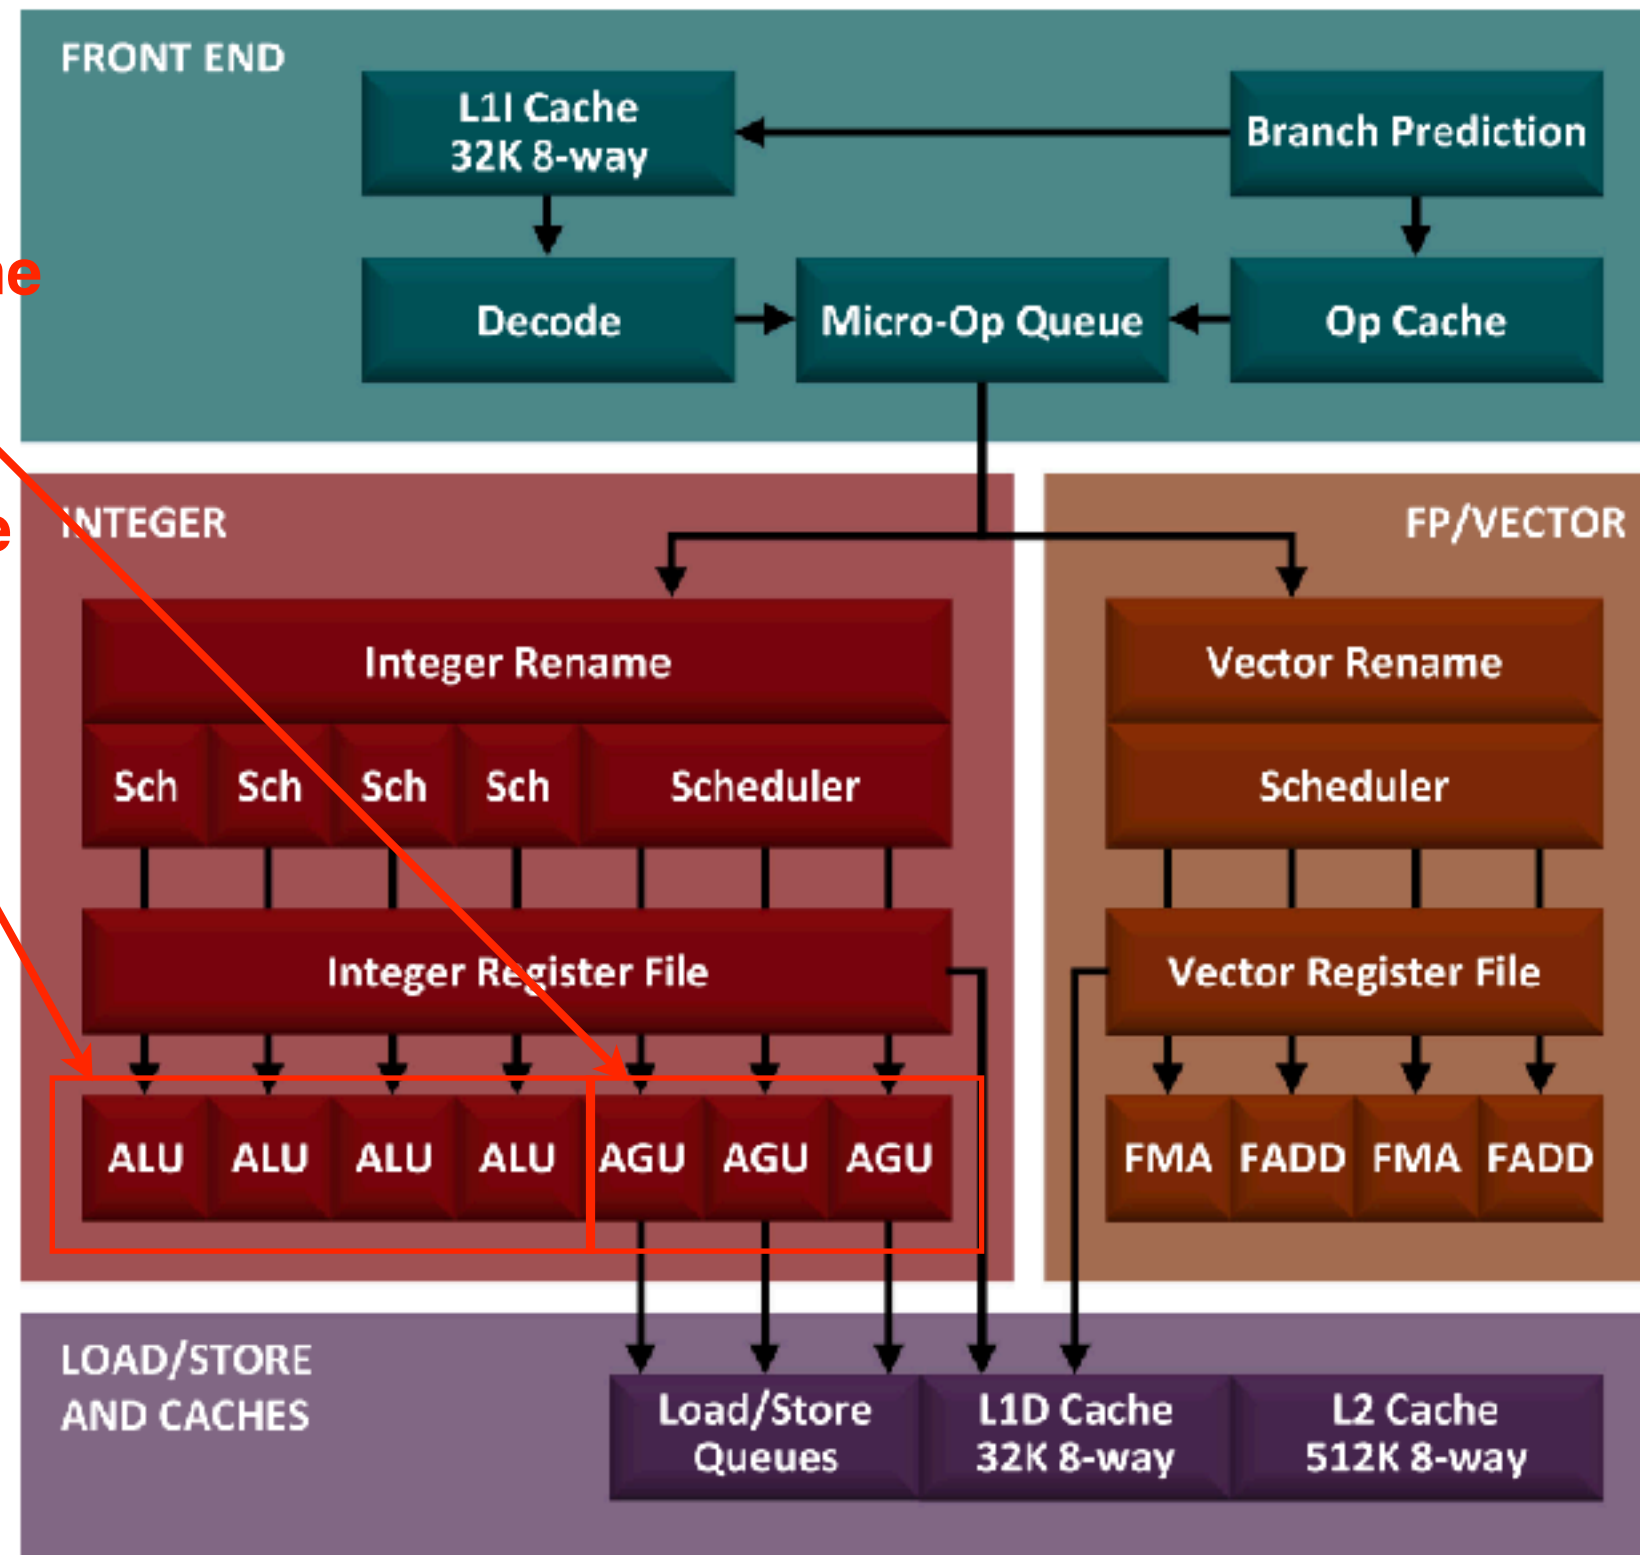
- Reorder buffer — a buffer keep track of the program order of instructions
  - Can be combined with IQ or physical registers — make either as a circular queue
- Commit stage — should the outcome of an instruction be realized
  - An instruction can only leave the pipeline if all it's previous are committed
  - If any prior instruction failed to commit, the instruction should yield it's ROB entry, restore all it's architectural changes

# **The pipelines of Modern Processors**

# Intel Skylake



# AMD Zen 2 (RyZen 3000 Series)



3-issue memory pipeline

4-issue integer pipeline

# Demo: ILP within a program

- perf is a tool that captures performance counters of your processors and can generate results like branch mis-prediction rate, cache miss rates and ILP.

# Announcements

- Assignment #3 due **Friday**
- Reading Quiz due next Monday